

Introduction to Alignment, Mapping and Indexing

The problem we want to solve

GGTACCCCCAAATTCGAAAGAGCCGAATGAAGTGGAAAGAATGGCCATTCAAATGGCTTGGGA

CAGGGCGCGCCGCTGGCCGAGAAAACCGCATTACAAATACTCTCAAAGCAGGATTTCTTTCTATAGGTAGTCAGTCCACAATTTTCACATTTGCAGCCA
GGTGGACATCTTTGGGCCAATGATGAGGAGAAGGTGGAGGAAGAAGGAAATGAGGAAGAGGAGAAAGAAGAGTATGGCATTTTTACAAACTGTGACCGTT
TCTGTGTGAAGATTTTGTAGCTGTCTGTTCGGACTTGGGGGTCTCAGGGAAACTCACTTTGCCCCCAGCTGAGGTTTTTCAGGAAATCTGGAAACCTACAG
TCTCCAAGCCTGCTCAGCCAAGAACGCGGGCGGGCGGGCGGGCGGTGGCGACGGCGGGCGGGCGGCAGGACCTCCGCGCCTCCATTATGCTATTCTGCCCGC
CGTGGGTGACAAACAGACGGATGCTACCAGCCCAGCCCAGTCCCGGGGAGCCAGCTGGCCTGGGGTTCGGTCCCGCGTCTTCCCTCATTCTGTGCCGCT
GCCGAGCCTGTCTCAGCTCCACACACGCTTGGGAGCTGCAGATGCCTCCGCCCCCTCCTCTCTCCCAGGCTCTTCCTGCCGTTGAACCCCGGCGGGCGGGC
TCTCGGCCAGCGGGCGCGCCCTGGTACCCCCAAATTCCAAAGAGCCGAATGAAGTGGAAAGAATGGCCACAAATGGCTTGGGCCCAGGTGACCATGGGAT
GGTTAGGTAGGATTTTAGAGGCGACTGCTCCTGGAATTAGAGAAAGAGTTTCATTACAACCGCTACCTGACCCGAAGGCACTCGGGCCCATGCCTTCCT
CTCCTTCGCTGTTTGATTTCTATTCTGTTTAGCCTGAGTTGGAGATGCCGGAGGTGCCGTGGTGGAAAGAAGTCTTGAACGAATTTGGAGGGCGTCTCCGT
GGCAGCTAAGCGAGCACGGGTTCTGCTGGTGCAGGATGACACTGGCAGCCACTGCCGCGGACTTGTATCTCTTGTCTTCCTGCTTTTATAGAGAATAGA
ATGACACTCACAACCTCTAACTACCTGTCAGAAGCAGGCAGGAGCTAGTAAGGATGAATTTGTAGCAAAATTAGCAAGTGGACTTCTTTCTCCTCTTCCT
CATTTCTTCTTCTCCTCCACCTTCTCCTCATCTTAAATCTTTAACATACTACCTAAAGGGAACCTGCAATAATCTTGAAAAAGGACTTCAATCCGACGTT
TTCGTGTCAAATAAGGATTAAAGAGAAACTCCTCCGCGAGCCGTGCGCCGAGGGTGGCGGGCGGGGGCCTGAAGCGTGAGGAGCCTTCAATATGTATTTA
ACCAGGGACCGTCGGTATGAGGTGGCCCCGGGTTCTTATTTGTTTGGGGGCTGGAGGGGGGAGACGGAGAAACAGTGAAAAGTTCTTGAGCCCCATAAAG
GGACTGTCTGGGGAGCGCCTCGTAGCCATAGAATTCCACCGCCGCCGCCGCCGCGTAGTCGTACTTGAAGCCGAGCGCAGGCGGGTGGTTCATTAAC
CTGACTTTGCCTTTGATTTTGTCTGACCTCTGCTTCGTCAAAATCTGGTTTCAGAATCGAAGGATGAAGATGAAAAAGATGAATAAAGGAGGAAAAGGA
AGAAAACAAGGACTAAGCAAAAAAGAAAGACCCCCCTATAGCAGGATTTTAAAATTTTTCTCTTTTTCTTTTTCAAGATTATTGCAAGGCGAGCGTGG
TGCAATATCCCGACTGTAAATCCTCCGCCAACACTAACTTTTAAAAAAAACACCCAGCAGGTACCATGCTAAGACAACATCACATGCATTATTATGAC
TCACGTATACAATACAAAGTACTTGGACCAGGAACAGGGTCTTTAATCCTTATTTGACACGAAAACGTCGGATTGAAGTCTCTCATGCCCAACTAGTGG
GGTTTCCTGGCACTGGACCCCAGCAAGTGGTCCTAGAGGCGAAAAGGAAGAAAACAAGGACTAAGCAAAAAAGAAAGACCCCCCCCCCTGAAAGGATA
TCATGGATTCTTAAAGAATGAGAACTTCGACATGGTGATATCACATGGACTTCTGGGGCCGAGTGAAGAATTTTCTGATGTTCTTTAGTTAGCAGTCC
TAAACCCTACCCAGCCTGCTGCCTCAGCACAGCCAAGGGAAAATCAGCAAGGGCTATGCTGTGATTCTGTCTCTGAGTGACTTGGACCACTGTTGATTT
TTATTTTTTTCTCTTTTTTTCTCCTATAGCAGGATTTTAAAATCGGGCCACCTTAACTCGGGAGGGCCGCGCTGAGGCTGGGAGCCGGAGATTCGGGC
GAGGGCAGTGTCTGCGGGGCGCGGTGCGCGAGCTCCCCGGGCGAGCCAGGTGCAGCCTTGGCGGGGTCTGTTTGGTGGGCGATGTCACCATTTCCCGCC
GCCGCCGTCGCCACCGCCGCCGCCGCCGCCGCGTAGTCGTACTTGAATAGCTGGACATAAAGACAAATGACAAAAAATTATTATTATAGATATATTTGGTC
TGTGTGTTATGTCCTAAGGTGTTTTGTCTGCAGTTTGAGAGCATGTTGCTGGTAGCCTGAGTTGGAGAT

The problem we want to solve

GGTACCCCCAAATTCCGAAAGAGCCGAATGAAGTGGAAAGAATGGCCAATCAAATGGCTTGGGA

CAGGGCGCGCCGCTGGCCGAGAAAACCGCATTACAAATACTCTCAAAGCAGGATTTCTTTCTATAGGTAGTCAGTCCACAATTTTCACATTTGCAGCCA
GGTGGACATCTTTGGGCCAATGATGAGGAGAAGGTGGAGGAAGAAGGAAATGAGGAAGAGGAGAAAGAAGAGTATGGCATTTTTACAAACTGTGACCGTT
TCTGTGTGAAGATTTTGTAGCTGTCTGTCTGCGACTTTGGGGGTCTCAGGGAAACTCACTTTGCCCCCAGCTGAGGTTTTTCAGGAAATCTGGAAACCTACAG
TCTCCAAGCCTGCTCAGCCAAGAACGCGGGCGGGCGGGCGGGCGGTGGCGACGGCGGGCGGGCGGCAGGACCTCCGCGCCTCCATTATGCTATTCTGCCCGC
CGTGGGTGACAAACAGACGGATGCTACCAGCCCAGCCCAGTCCCGGGGAGCCAGCTGGCCTGGGGTTCGGTCCCGCGTCTTCCCTCATTCTGTGCCGCT
GCCGAGCCTGTCTCAGCTCCACACACGCTTGGGAGCTGCAGATGCCTCCGCCCCCTCCTCTCTCCAGGCTCTTCCTGCCGTTGAACCCCGGCGGGCGGGC
TCTCGGCCAGCGGCGCGCCCTGGTACCCCCAAATTCCAAAGAGCCGAATGAAGTGGAAAGAATGGCCACAATGGCTTGGGC

CCGAGTGACCATGGGAT
GGTTAGGTAGGATTTTAGAGGCGACTGCTCCTGGAATTAGAGAAAGAGTTTCATTACAACCGCTACCTGACCCGAAGGCACTCGGGCCCATGCCTTCCT
CTCCTTCGCTGTTTGTATTTCTATTCTGTTTAGCCTGAGTTGGAGATGCCGGAGGTGCCGTGGTGGAAAGAAGTCTTGAACGAATTTGGAGGCGTCTCCGT
GGCAGCTAAGCGAGCACGGGTTCTGCTGGTGCAGGATGACACTGGCAGCCACTGCCGCGGACTTGTATCTCTTGTCTTCCTGCTTTTATAGAGAATAGA
ATGACACTCACAACCTCTAACTACCTGTCAGAAGCAGGCAGGAGCTAGTAAGGATGAATTTGTAGCAAAATTAGCAAGTGGACTTCTTTCTCCTCTTCCT
CATTTCTTCTTCTCCTCCACCTTCTCCTCATCTTAAATCTTTAACATACTACCTAAAGGGAACCTGCAATAATCTTGAAAAAGGACTTCAATCCGACGTT
TTCGTGTCAAATAAGGATTAAGAGAAACTCCTCCGCGAGCCGTGCGCCGAGGGTGGCGGGCGGGGGCCTGAAGCGTGAGGAGCCTTCAATATGTATTTA
ACCAGGGACCGTCGGTATGAGGTGGCCCCGGGTTCTTATTTGTTTGGGGGCTGGAGGGGGGAGACGGAGAAACAGTGAAAAGTTCTTGAGCCCCATAAAG
GGACTGTCTGGGGAGCGCCTCGTAGCCATAGAATTCCACCGCCGCCGCCGCCGCGTAGTCGTACTTGAAGCCGAGCGCAGGCGGGTGGTTCATTAAC
CTGACTTTGCCTTTGATTTTGTCTGACCTCTGCTTCGTCAAAATCTGGTTTCAGAATCGAAGGATGAAGATGAAAAAGATGAATAAAGGAGGAAAAGGA
AGAAAACAAGGACTAAGCAAAAAAGAAAGACCCCCCTATAGCAGGATTTTAAAATTTTCTCTTTTTCTTTTTCAAGATTATTGCAAGGCGAGCGTGG
TGCAATATCCCGACTGTAAATCCTCCGCCAACACTAACTTTTAAAAAAAACACCCAGCAGGTACCATGCTAAGACAACATCACATGCATTATTATGAC
TCACGTATACAATACAAAGTACTTGGACCAGGAACAGGGTCTTTAATCCTTATTTGACACGAAAACGTCGGATTGAAGTCTCTCATGCCCAACTAGTGG
GGTTTCCTGGCACTGGACCCCAGCAAGTGGTCCTAGAGGCGAAAAGGAAGAAAACAAGGACTAAGCAAAAAAGAAAGACCCCCCCCCCTGAAAGGATA
TCATGGATTCTTAAAGAATGAGAACTTCGACATGGTGATATCACATGGACTTCTGGGGCCGAGTGAAGAATTTTCTGATGTTCTTTAGTTAGCAGTCC
TAAACCCTACCCAGCCTGCTGCCTCAGCACAGCCAAGGGAAAATCAGCAAGGGCTATGCTGTGATTCTGTCTCTGAGTGACTTGGACCACTGTTGATTT
TTATTTTTTTCTCTTTTTTTCTCCTATAGCAGGATTTTAAAATCGGGCCACCTTAACTCGGGAGGGCCGCGCTGAGGCTGGGAGCCGGAGATTCGGGC
GAGGGCAGTGTCTGCGGGGCGCGGTGCGCGAGCTCCCCGGGCGAGCCAGGTGCAGCCTTGGCGGGGTCTGTTTGGTGGGCGATGTCACCATTTCCCGCC
GCCGCCGTCGCCACCGCCGCCGCCGCCGCCGCGTAGTCGTACTTGAATAGCTGGACATAAAGACAAATGACAAAAAATTATTATTATAGATATATTTGGTC
TGTGTGTTATGTCCTAAGGTGTTTTGTCTGCAGTTTGAGAGCATGTTGCTGGTAGCCTGAGTTGGAGAT

The problem we want to solve

[illegible]

The problem we want to solve

[illegible]

Given: *A reference* text (e.g. genome) and a query (e.g. read)

Find: The location in the reference with the “most similarity” / “smallest distance” to the query.

The Language of Strings

A **string** **s** is a finite sequence of characters

|s| denotes the **length** of the string — the number of characters in the sequence.

A string is defined over an **alphabet**, Σ

$$\Sigma_{\text{DNA}} = \{A, T, C, G\}$$

$$\Sigma_{\text{RNA}} = \{A, U, C, G\}$$

$$\Sigma_{\text{AminoAcid}} = \{A, R, N, D, C, E, Q, G, H, I, L, K, M, F, P, S, T, W, Y, V\}$$

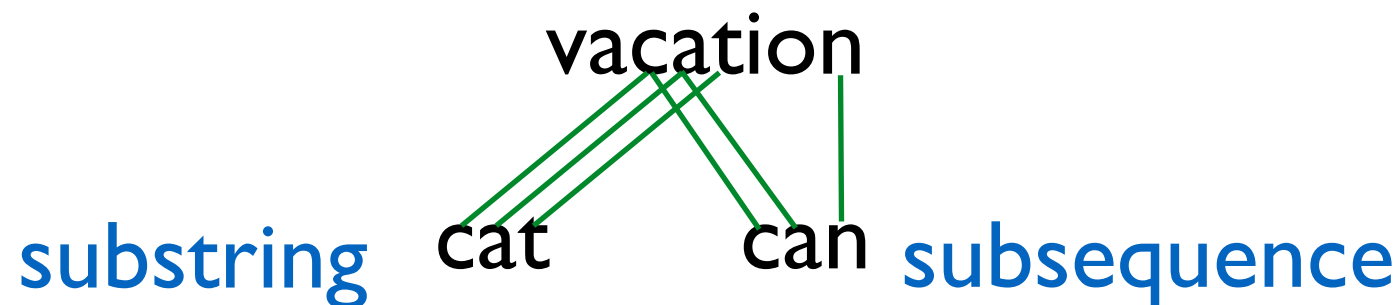
The **empty string** is denoted ϵ — $|\epsilon| = 0$

The Language of Strings

Given two strings **s**, **t** over the same alphabet Σ , we denote the **concatenation** as **st** — this is the sequence of **s** followed by the sequence of **t**

String **s** is a **substring** of **t** if there exist two (potentially empty) strings **u** and **v** such that **t = usv**

String **s** is a **subsequence** of **t** if the characters of **s** appear in order (but not necessarily consecutively) in **t**



String **s** is a **prefix/suffix** of **t** if **t = su/us** — if neither **s** nor **u** are ϵ , then **s** is a **proper prefix/suffix** of **t**

The Simplest String Comparison Problem

Given: Two strings

$$a = a_1a_2a_3a_4\dots a_m$$

$$b = b_1b_2b_3b_4\dots b_n$$

where a_i, b_i are letters from some alphabet, Σ , like $\{A,C,G,T\}$.

Compute how **similar** the two strings are.

What do we mean by “similar”?

Edit distance between strings a and b = the smallest number of the following operations that are needed to transform a into b :

- mutate (replace) a character
- delete a character
- insert a character

riddle $\xrightarrow{\text{delete}}$ ridle $\xrightarrow{\text{mutate}}$ riple $\xrightarrow{\text{insert}}$ triple

The String Alignment Problem

Parameters:

- “*gap*” is the cost of inserting a “-” character, representing an insertion or deletion (insertion/deletion are dual operations depending on the string)
- $cost(x,y)$ is the cost of aligning character x with character y .
In the simplest case, $cost(x,x) = 0$ and $cost(x,y) = \text{mismatch penalty}$.

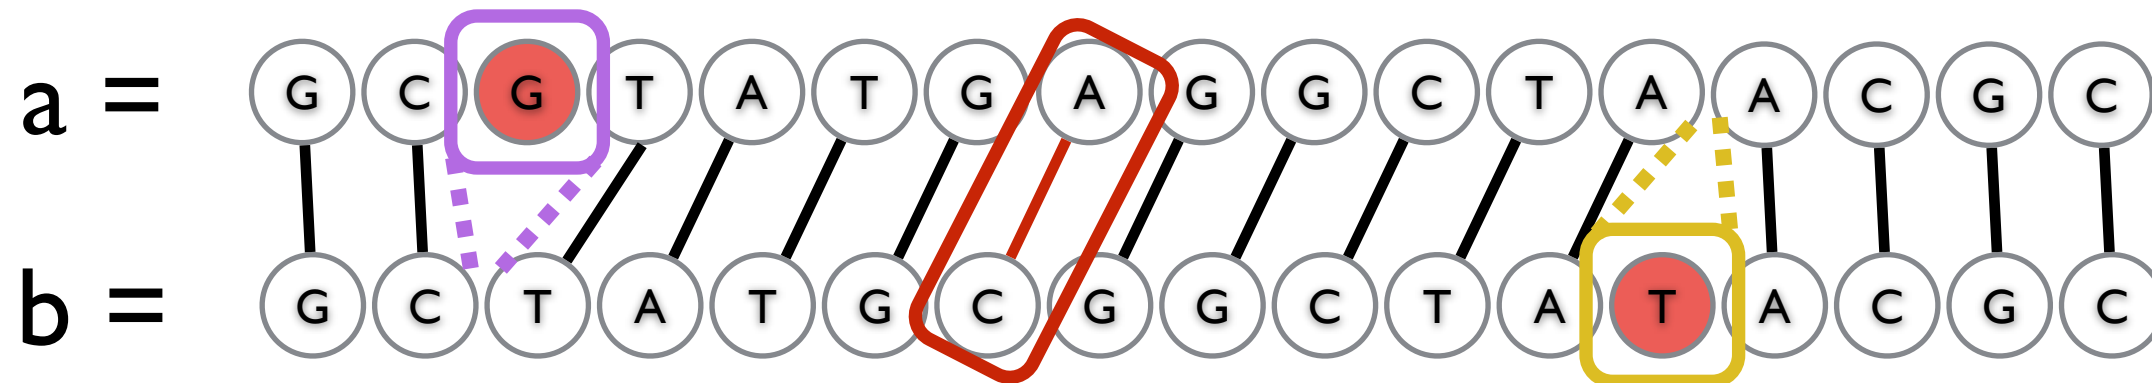
Goal:

- Can compute the edit distance by finding the **lowest cost alignment**. (often phrased as finding **highest scoring alignment**.)
- Cost of an alignment is: sum of the $cost(x,y)$ for the pairs of characters that are aligned + $gap \times \text{number of - characters inserted}$.

Another View: Alignment as a Matching

Each string is a set of nodes, one for each character.

Looking for a low-cost matching (pairing) between the sequences.



The operations at our disposal

Insertion (into **a** ~ deletion from **b**)

Mutation

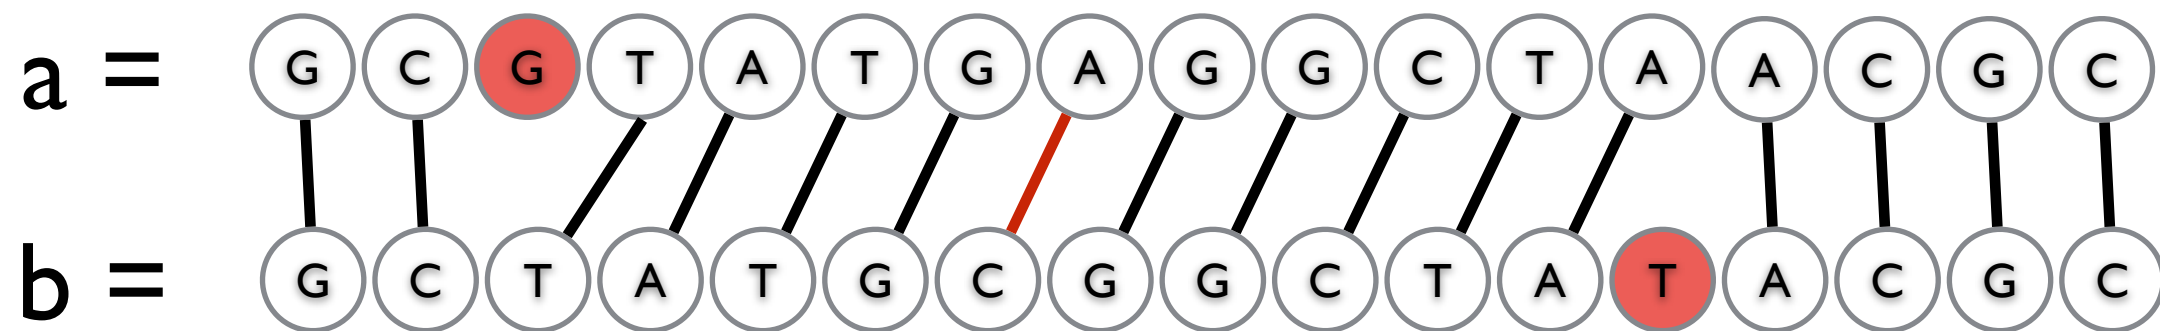
Deletion (from **a** ~ insertion into **b**)

When we “delete a” character in **a** this is the same as inserting the character “-” in **b**. Conceptually, you can think of this as aligning the deleted character with “-”. Under this model $\text{cost}(x, '-') = \text{cost}('-', x) = \text{gap}$ for any $x \in \Sigma$

Another View: Alignment as a Matching

Each string is a set of nodes, one for each character.

Looking for a low-cost matching (pairing) between the sequences.



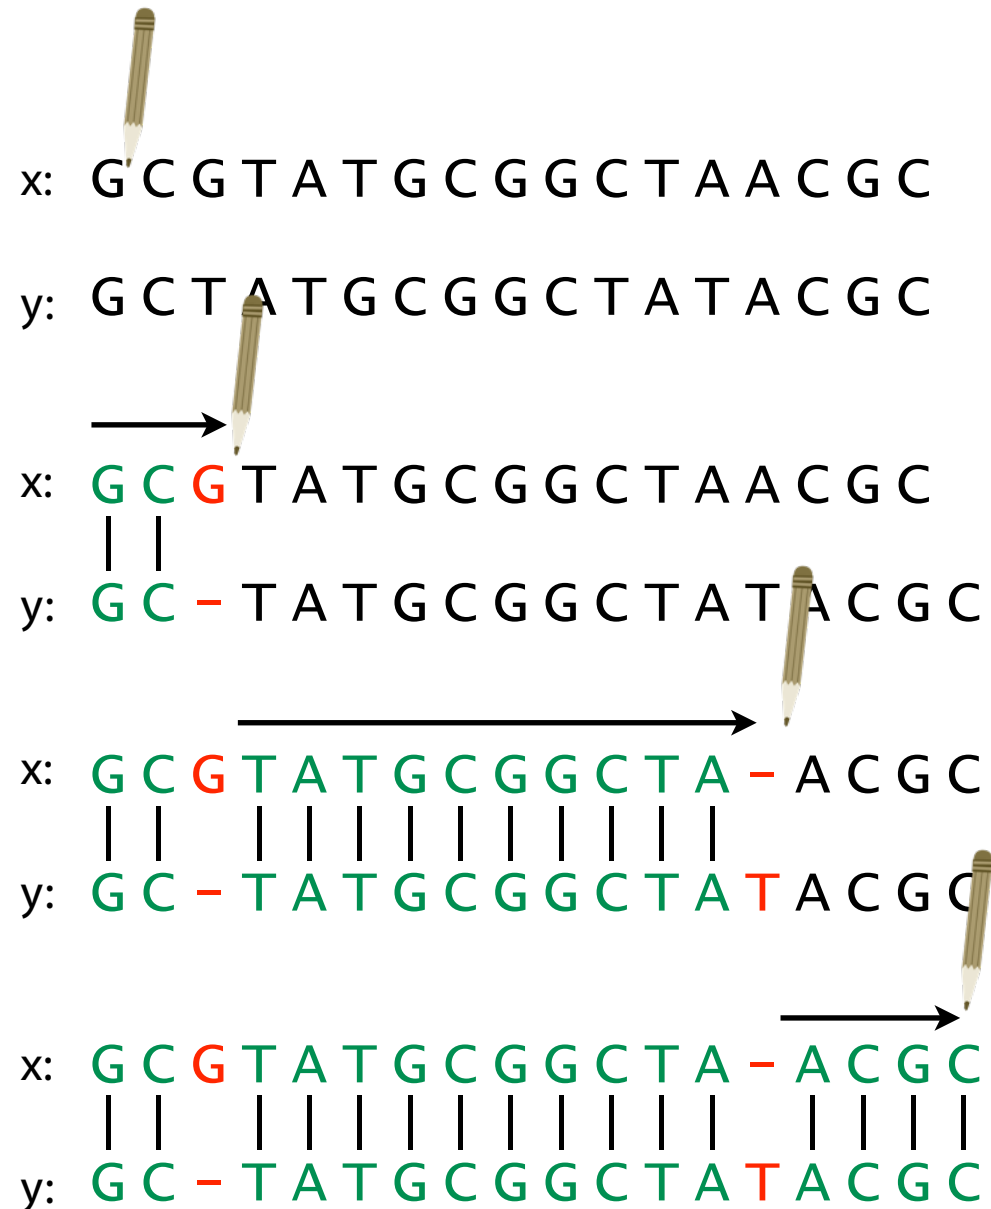
Cost of a matching is:

$$\text{gap} \times \#unmatched + \sum_{(a_i, b_j)} \text{cost}(a_i, b_j)$$

Edges are not allowed to cross!

Representing alignments as edit transcripts

Can think of edits as being introduced by an *optimal editor* working left-to-right. *Edit transcript* describes how editor turns x into y .



Operations:

M = match, **R** = replace,

I = insert into x , **D** = delete from x

MMD

MMDMMMMMMMMMMMI

MMDMMMMMMMMMMMI MMMM

Representing edits as alignments

prin-ciple
|||| |||xx
prinncipal
(1 gap, 2 mm)
MMMMIMMMRR

misspell
|||| ||||
mis-pell
(1 gap)
MMMIMMMM

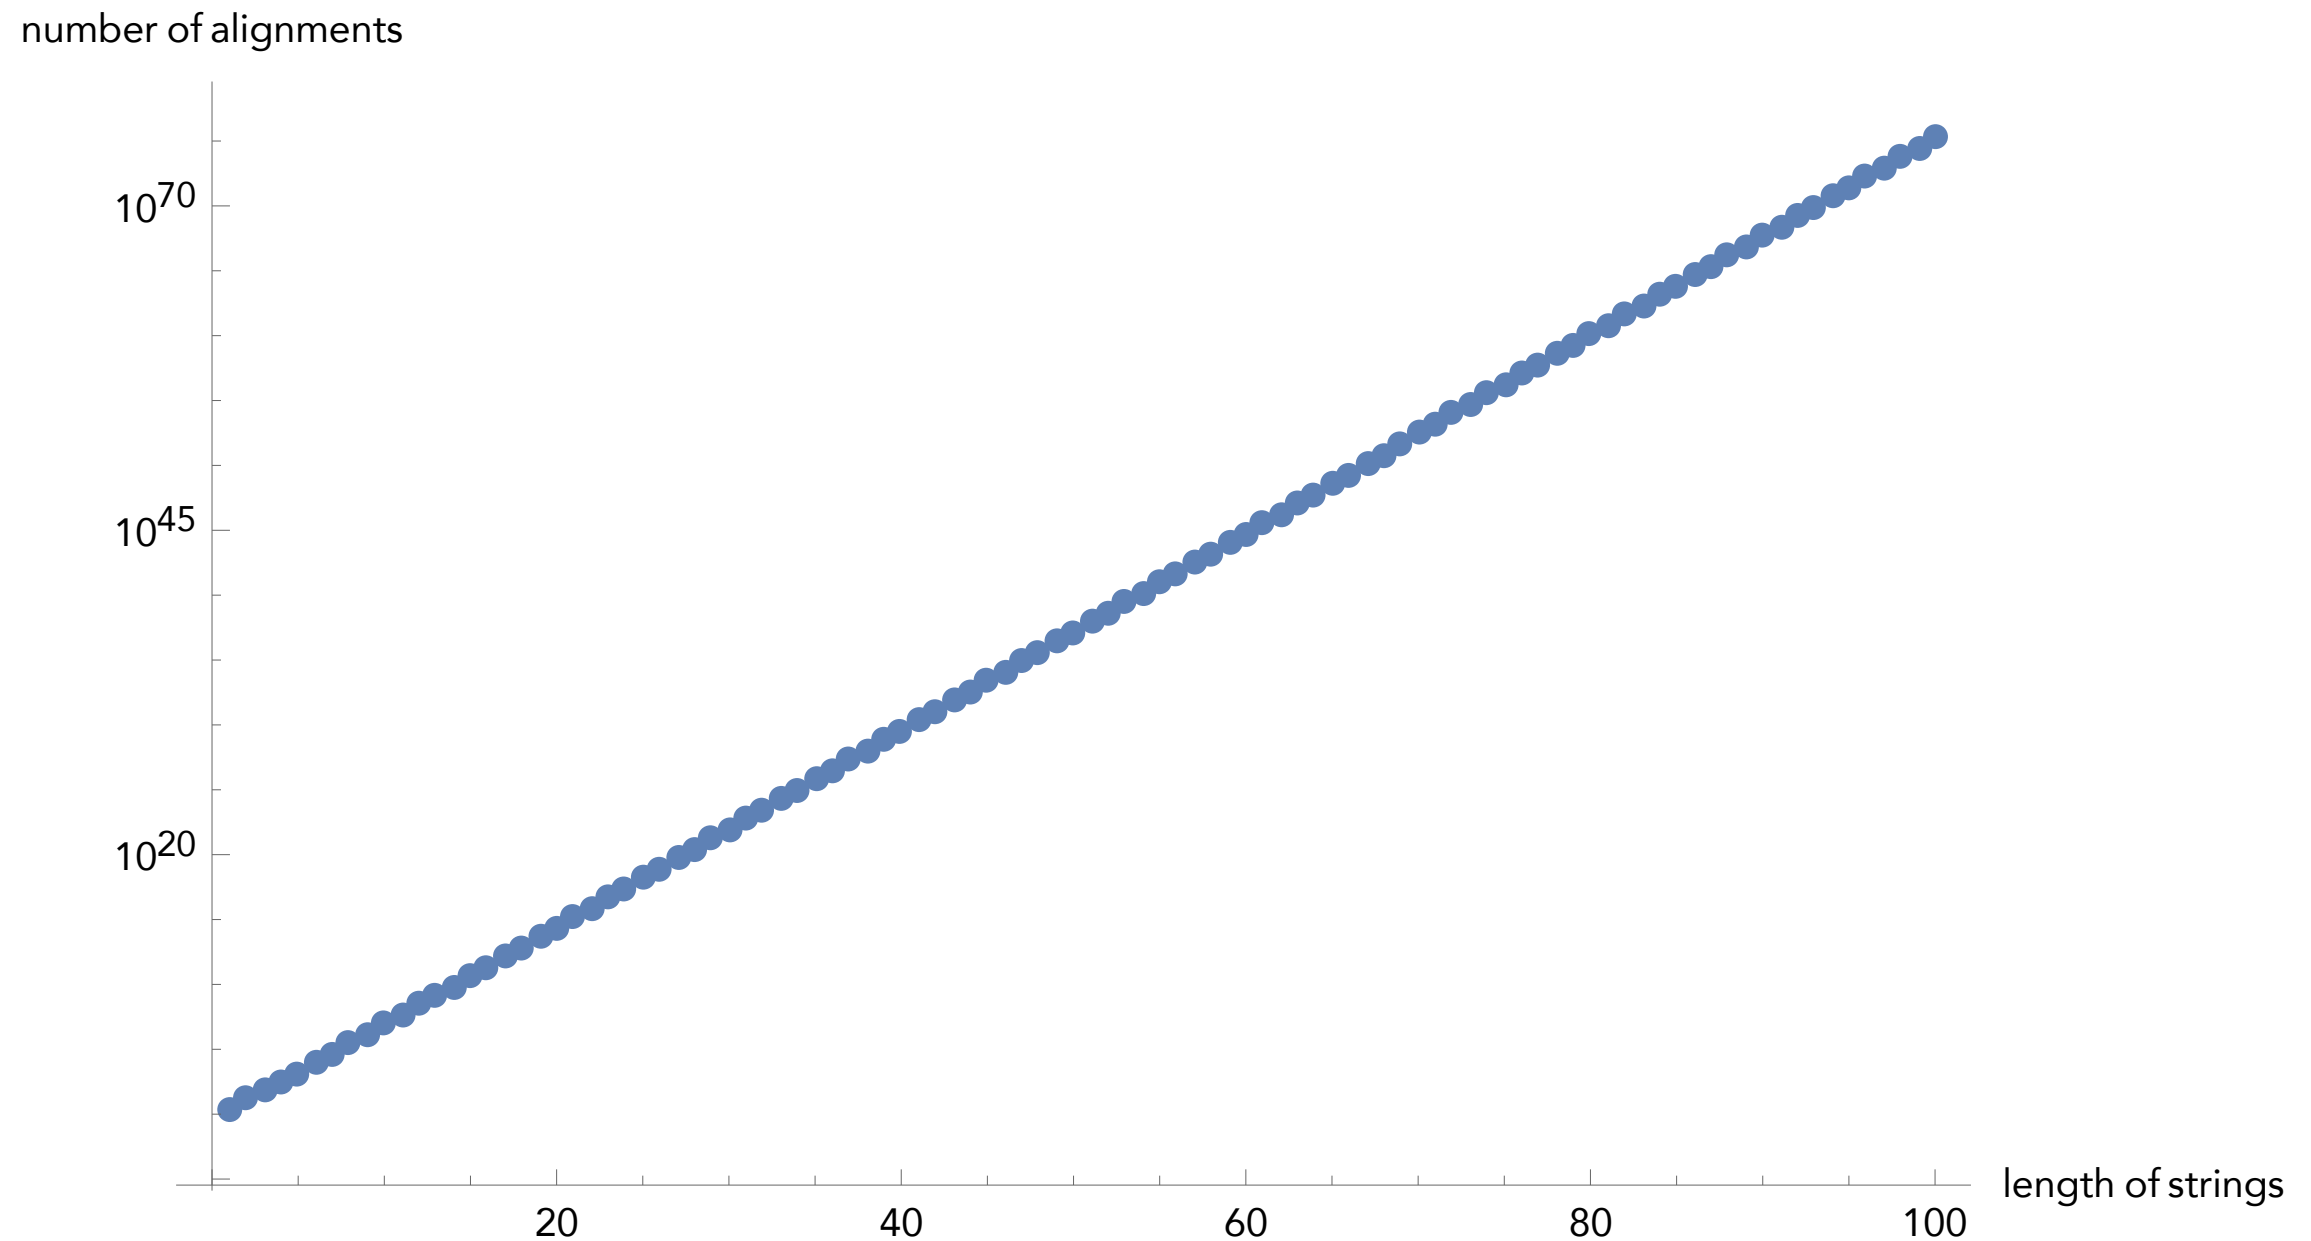
aa-bb-ccaabb
|x || ||
ababbbc-a-b-
(5 gaps, 1 mm)
MRIMMIMDMDMD

prin-cip-le
|||| ||||
prinncipal-
(3 gaps, 0 mm)
MMMMIMMMIMD

prehistoric
|||| ||||
---historic
(3 gaps)
DDDMMMMMMM

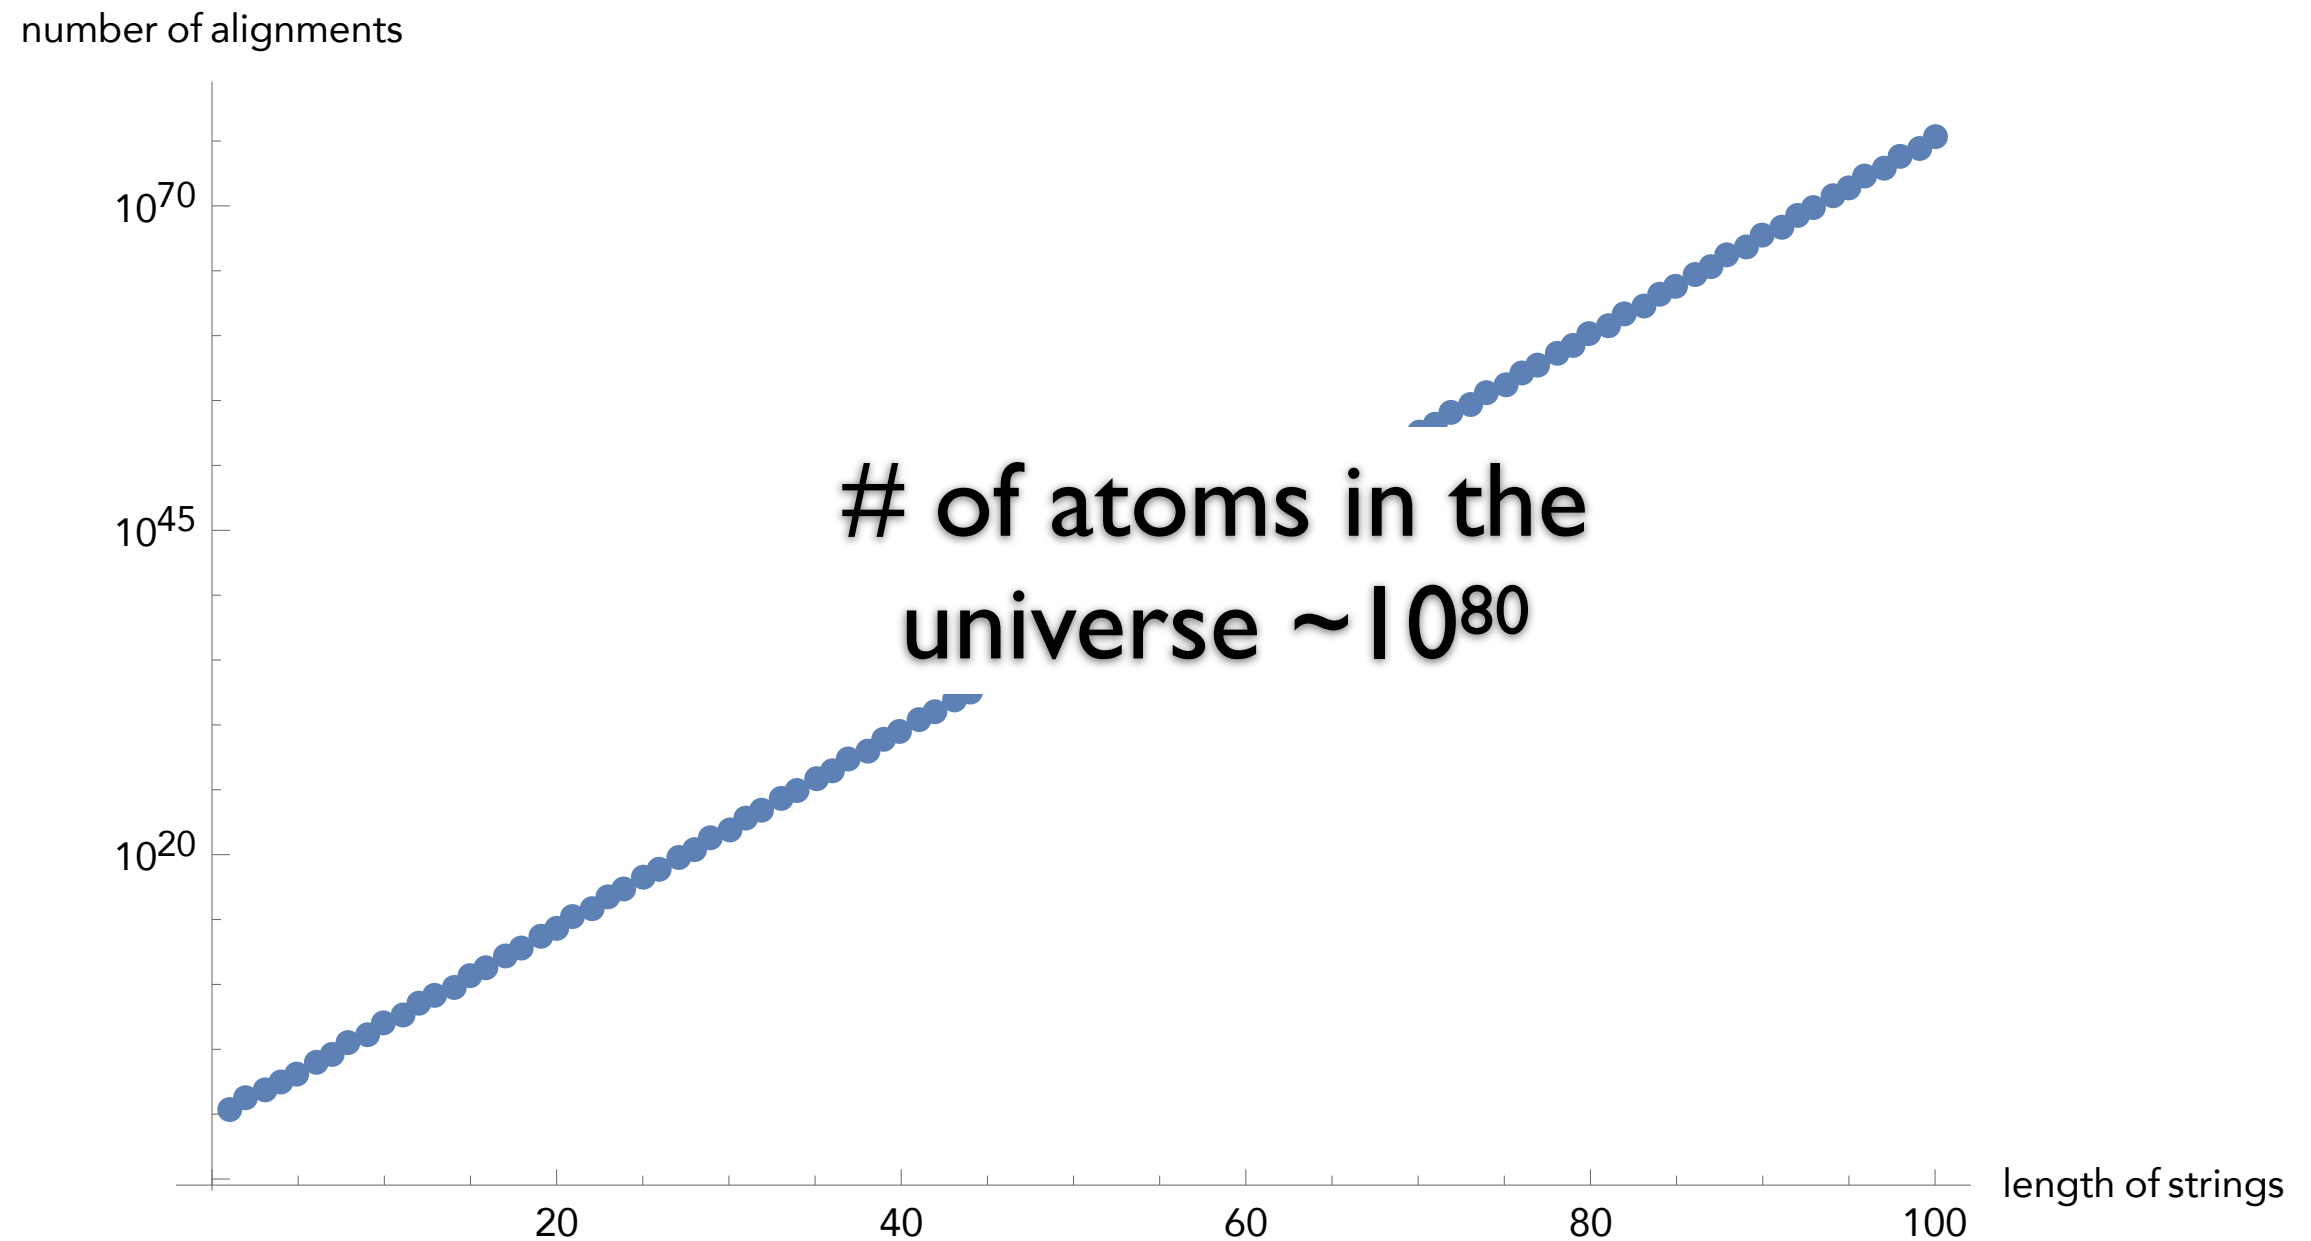
al-go-rithm-
|| xx ||x |
alKhwariz-mi
(4 gaps, 3 mm)
MMIRRIIMRDMI

How many alignments are there?



$$f(n, m) = \sum_{k=0}^{\min(m, n)} 2^k \binom{m}{k} \binom{n}{k}$$

How many alignments are there?



$$f(n, m) = \sum_{k=0}^{\min(m, n)} 2^k \binom{m}{k} \binom{n}{k}$$

Algorithm for Computing Edit Distance

Consider the last characters of each string:

$$\begin{aligned}a &= a_1a_2a_3a_4\dots a_m \\ b &= b_1b_2b_3b_4\dots b_n\end{aligned}$$

One of these possibilities must hold:

1. (a_m, b_n) are matched to each other
2. a_m is not matched at all
3. b_n is not matched at all
4. a_m is matched to some b_j ($j \neq n$) and b_n is matched to some a_k ($k \neq m$).

Algorithm for Computing Edit Distance

Consider the last characters of each string:

$$a = a_1a_2a_3a_4\dots a_m$$
$$b = b_1b_2b_3b_4\dots b_n$$

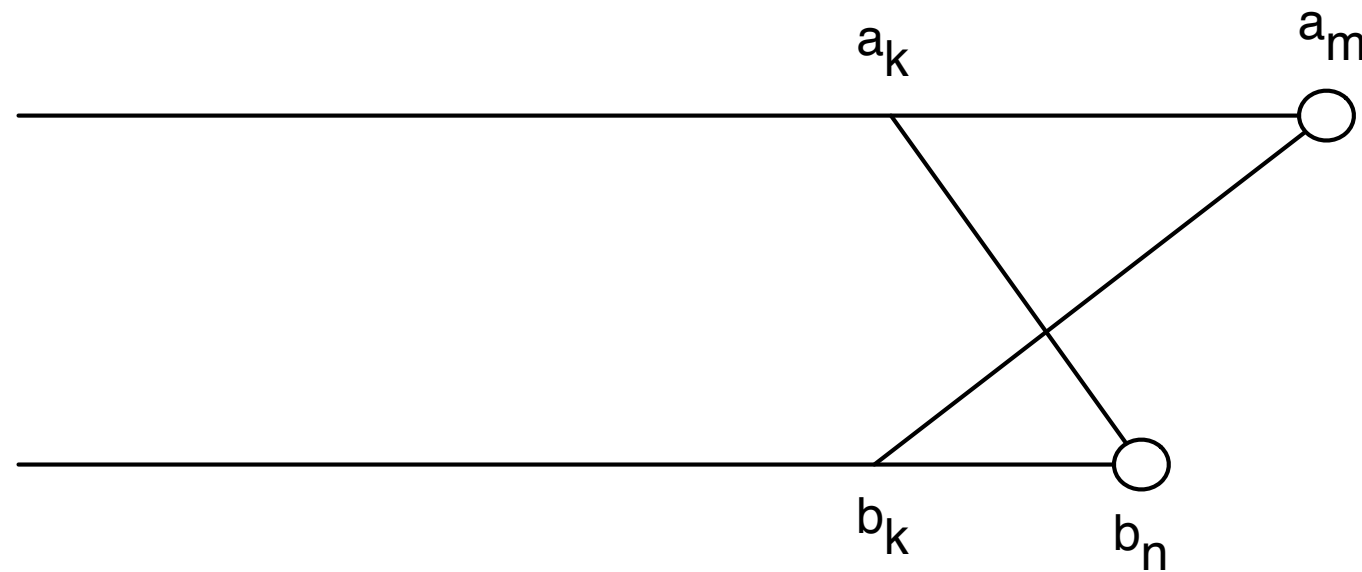
One of these possibilities must hold:

1. (a_m, b_n) are matched to each other
2. a_m is not matched at all
3. b_n is not matched at all
4. a_m is matched to some b_j ($j \neq n$) and b_n is matched to some a_k ($k \neq m$).

#4 can't happen! Why?

No Crossing Rule Forbids #4

4. a_m is matched to some b_j ($j \neq n$) and b_n is matched to some a_k ($k \neq m$).



So, the only possibilities for what happens to the last characters are:

1. (a_m, b_n) are matched to each other
2. a_m is not matched at all
3. b_n is not matched at all

Recursive Solution

Turn the 3 possibilities into 3 cases of a recurrence:

$$OPT(i, j) = \min \begin{cases} \text{cost}(a_i, b_j) + OPT(i-1, j-1) & \text{match } a_i, b_j \\ \text{gap} + OPT(i-1, j) & a_i \text{ is not matched} \\ \text{gap} + OPT(i, j-1) & b_j \text{ is not matched} \end{cases}$$

↑
Cost of the optimal alignment between $a_1 \dots a_i$ and $b_1 \dots b_j$

↑
Written in terms of the costs of smaller problems

Key: we don't know which of the 3 possibilities is the right one, so we try them all.

Base case: $OPT(i, 0) = i \times \text{gap}$ and $OPT(0, j) = j \times \text{gap}$.

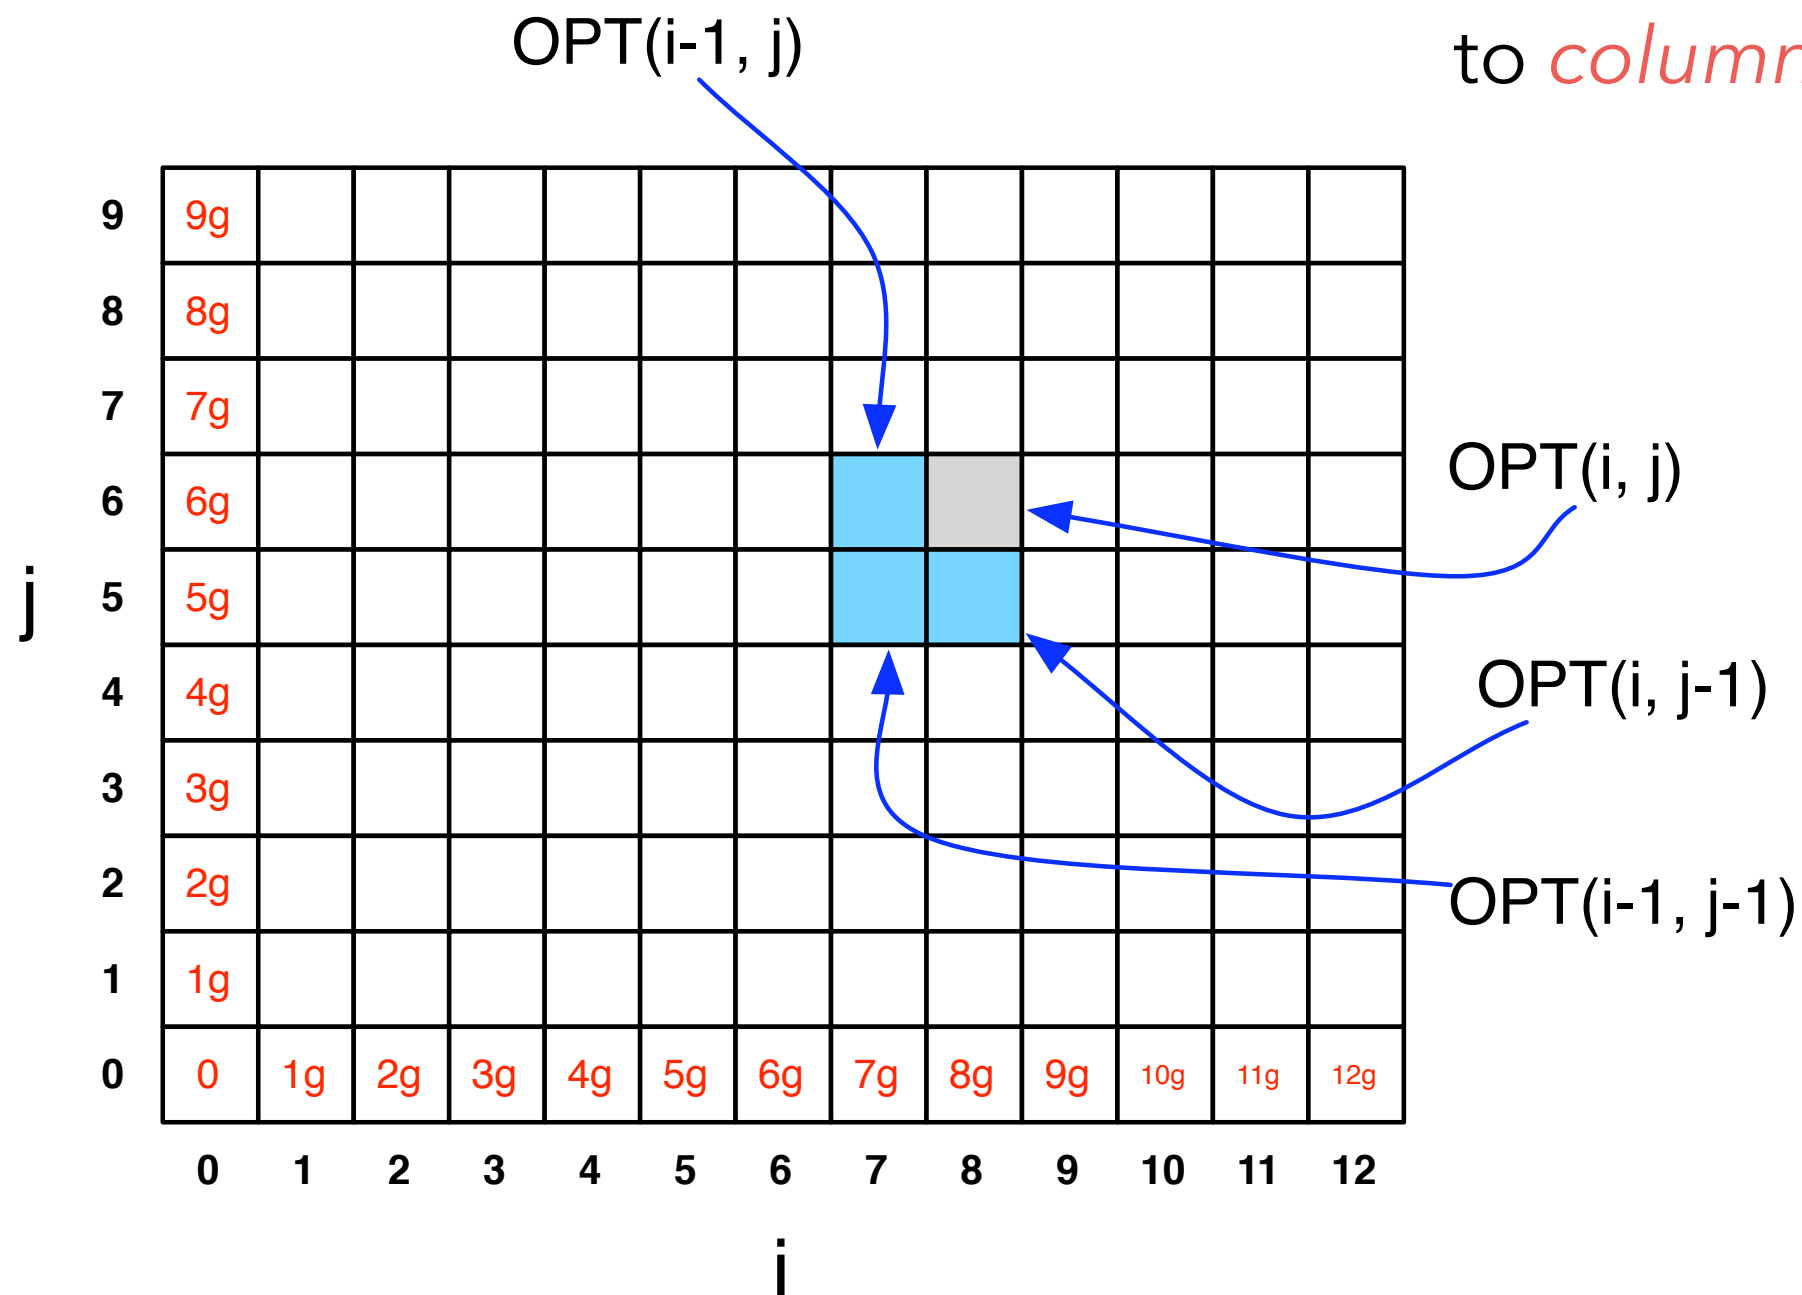
(Aligning i characters to 0 characters must use i gaps.)

Computing $OPT(i,j)$ Efficiently

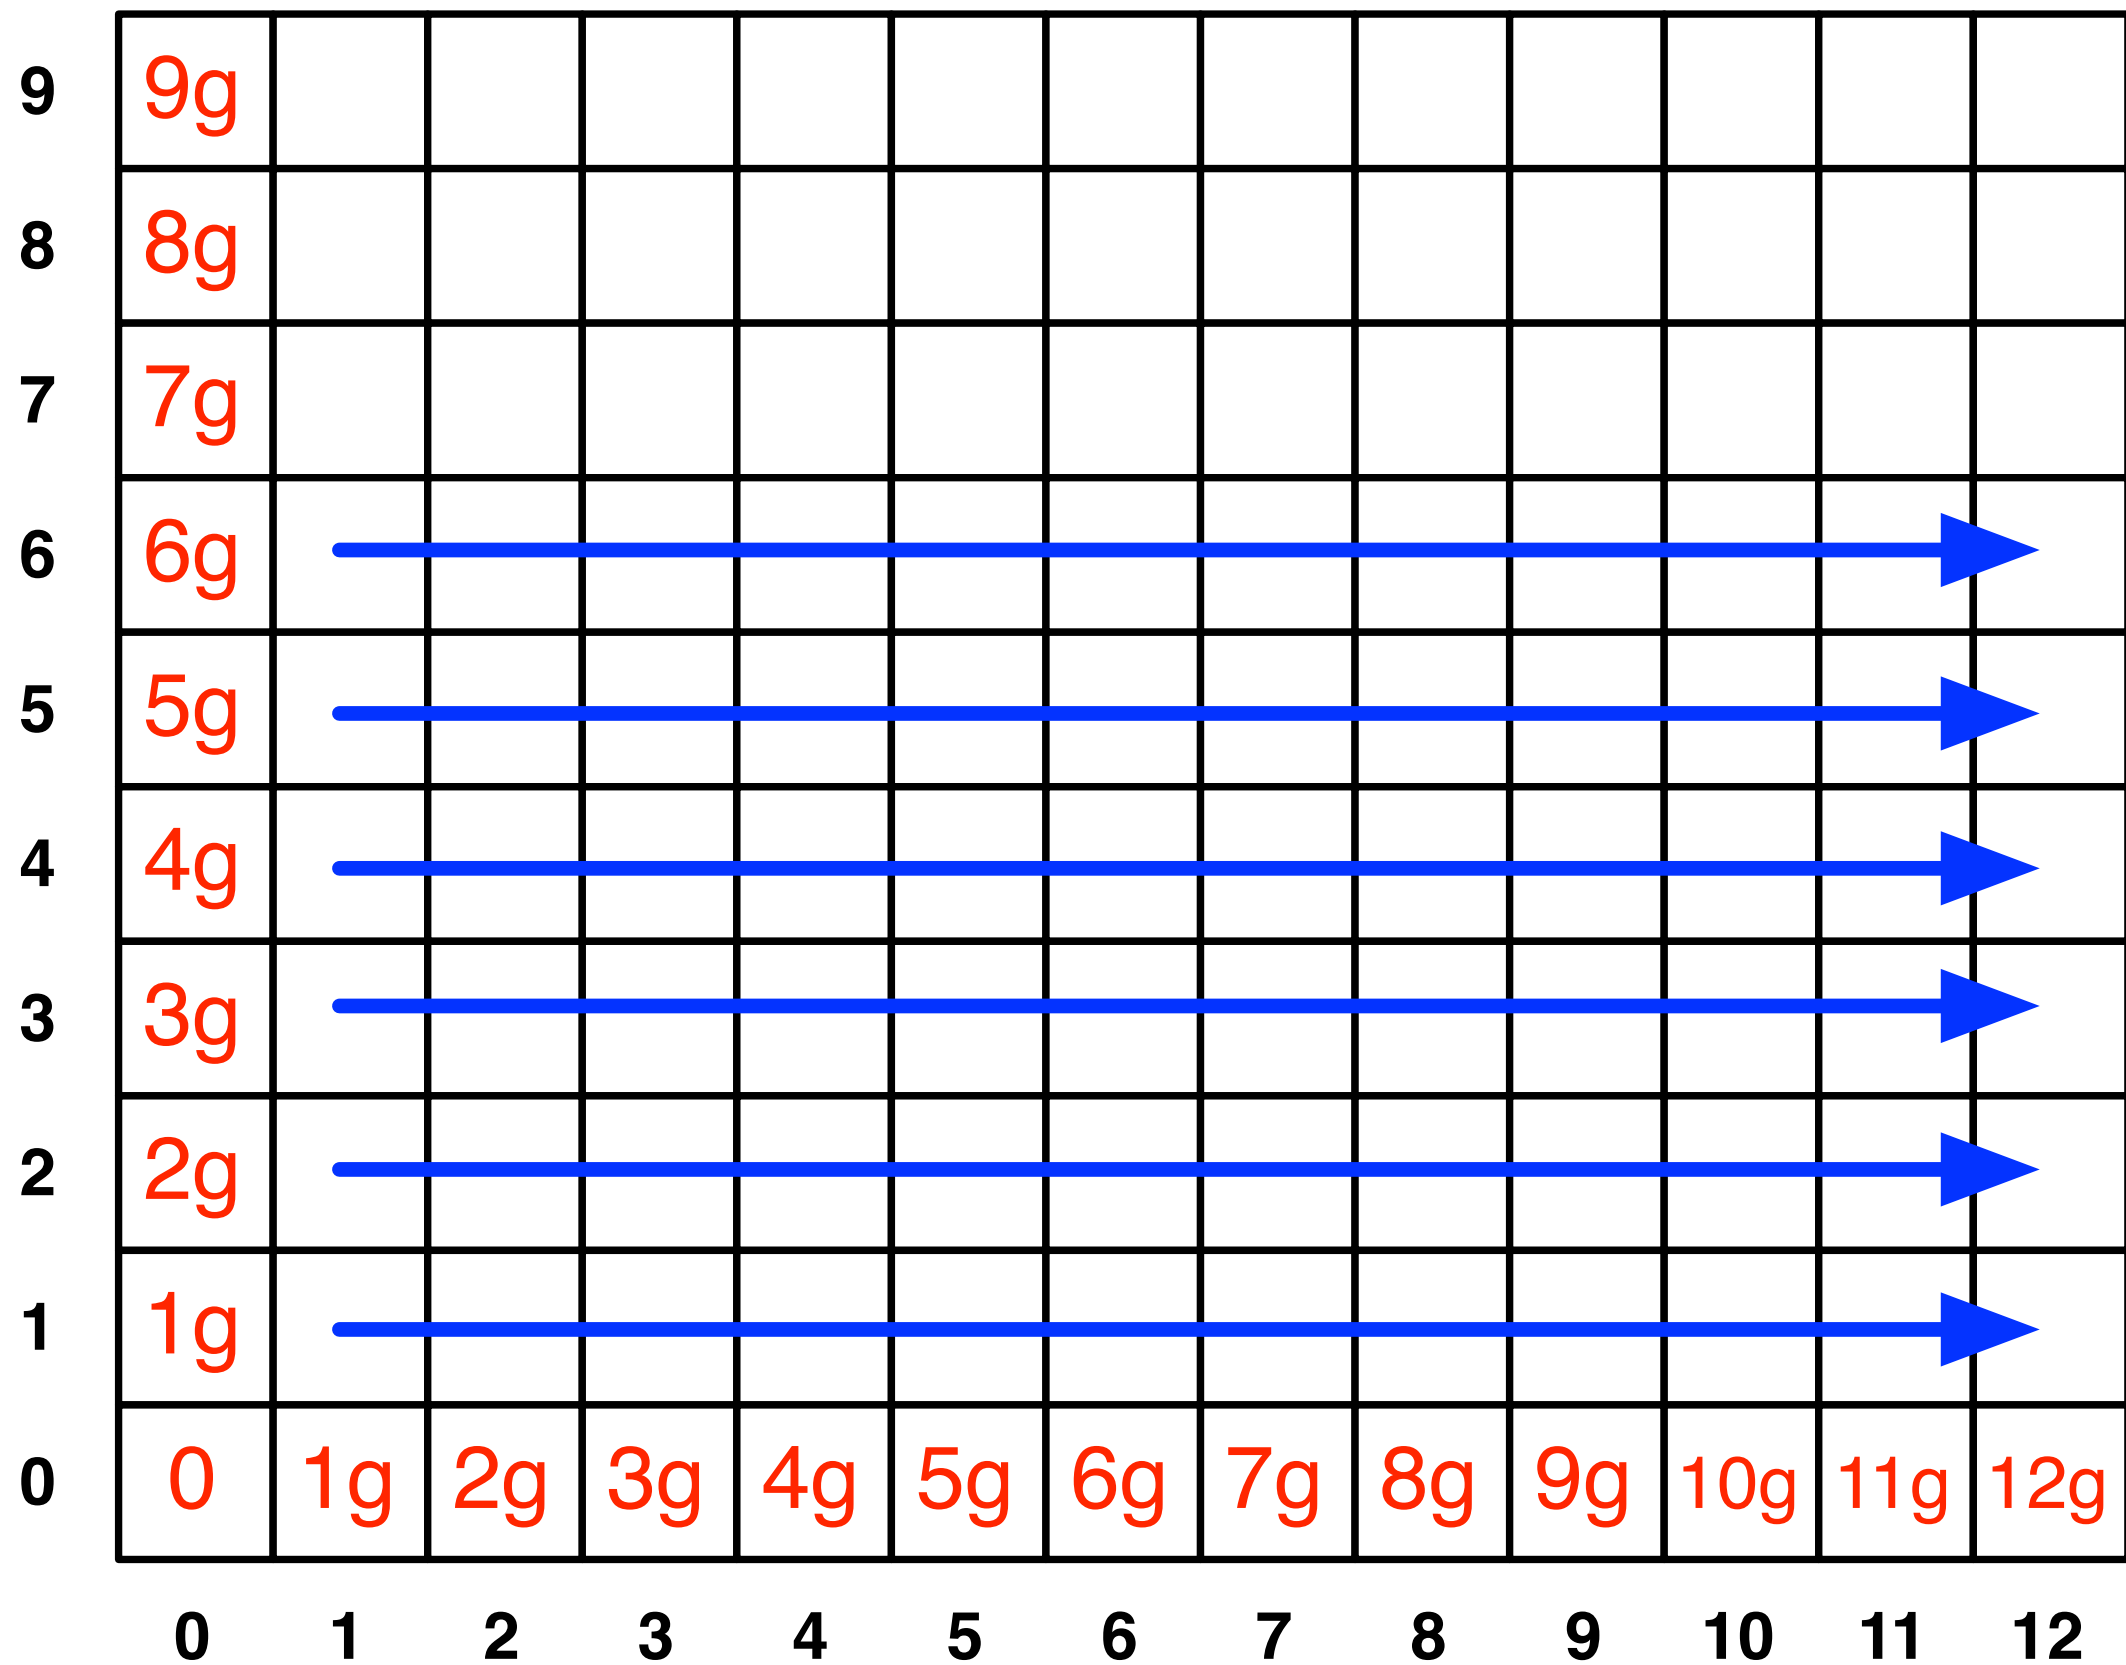
We're ultimately interested in $OPT(n,m)$, but we will compute all other $OPT(i,j)$ ($i \leq n, j \leq m$) on the way to computing $OPT(n,m)$.

Store those values in a 2D array:

NOTE: observe the non-standard notation here; $OPT(i,j)$ is referring to *column* i , *row* j of the matrix.



Filling in the 2D Array



Edit Distance Computation

```
EditDistance(X,Y):  
  For i = 1,...,m: A[i,0] = i*gap  
  For j = 1,...,n: A[0,j] = j*gap  
  
  For i = 1,...,m:  
    For j = 1,...,n:  
      A[i,j] = min(  
        cost(a[i],b[j]) + A[i-1,j-1],  
        gap + A[i-1,j],  
        gap + A[i,j-1]  
      )  
    EndFor  
  EndFor  
  Return A[m,n]
```

Where's the answer?

$\text{OPT}(n,m)$ contains the edit distance between the two strings.

Why? By induction: EVERY cell contains the optimal edit distance between some prefix of string 1 with some prefix of string 2.

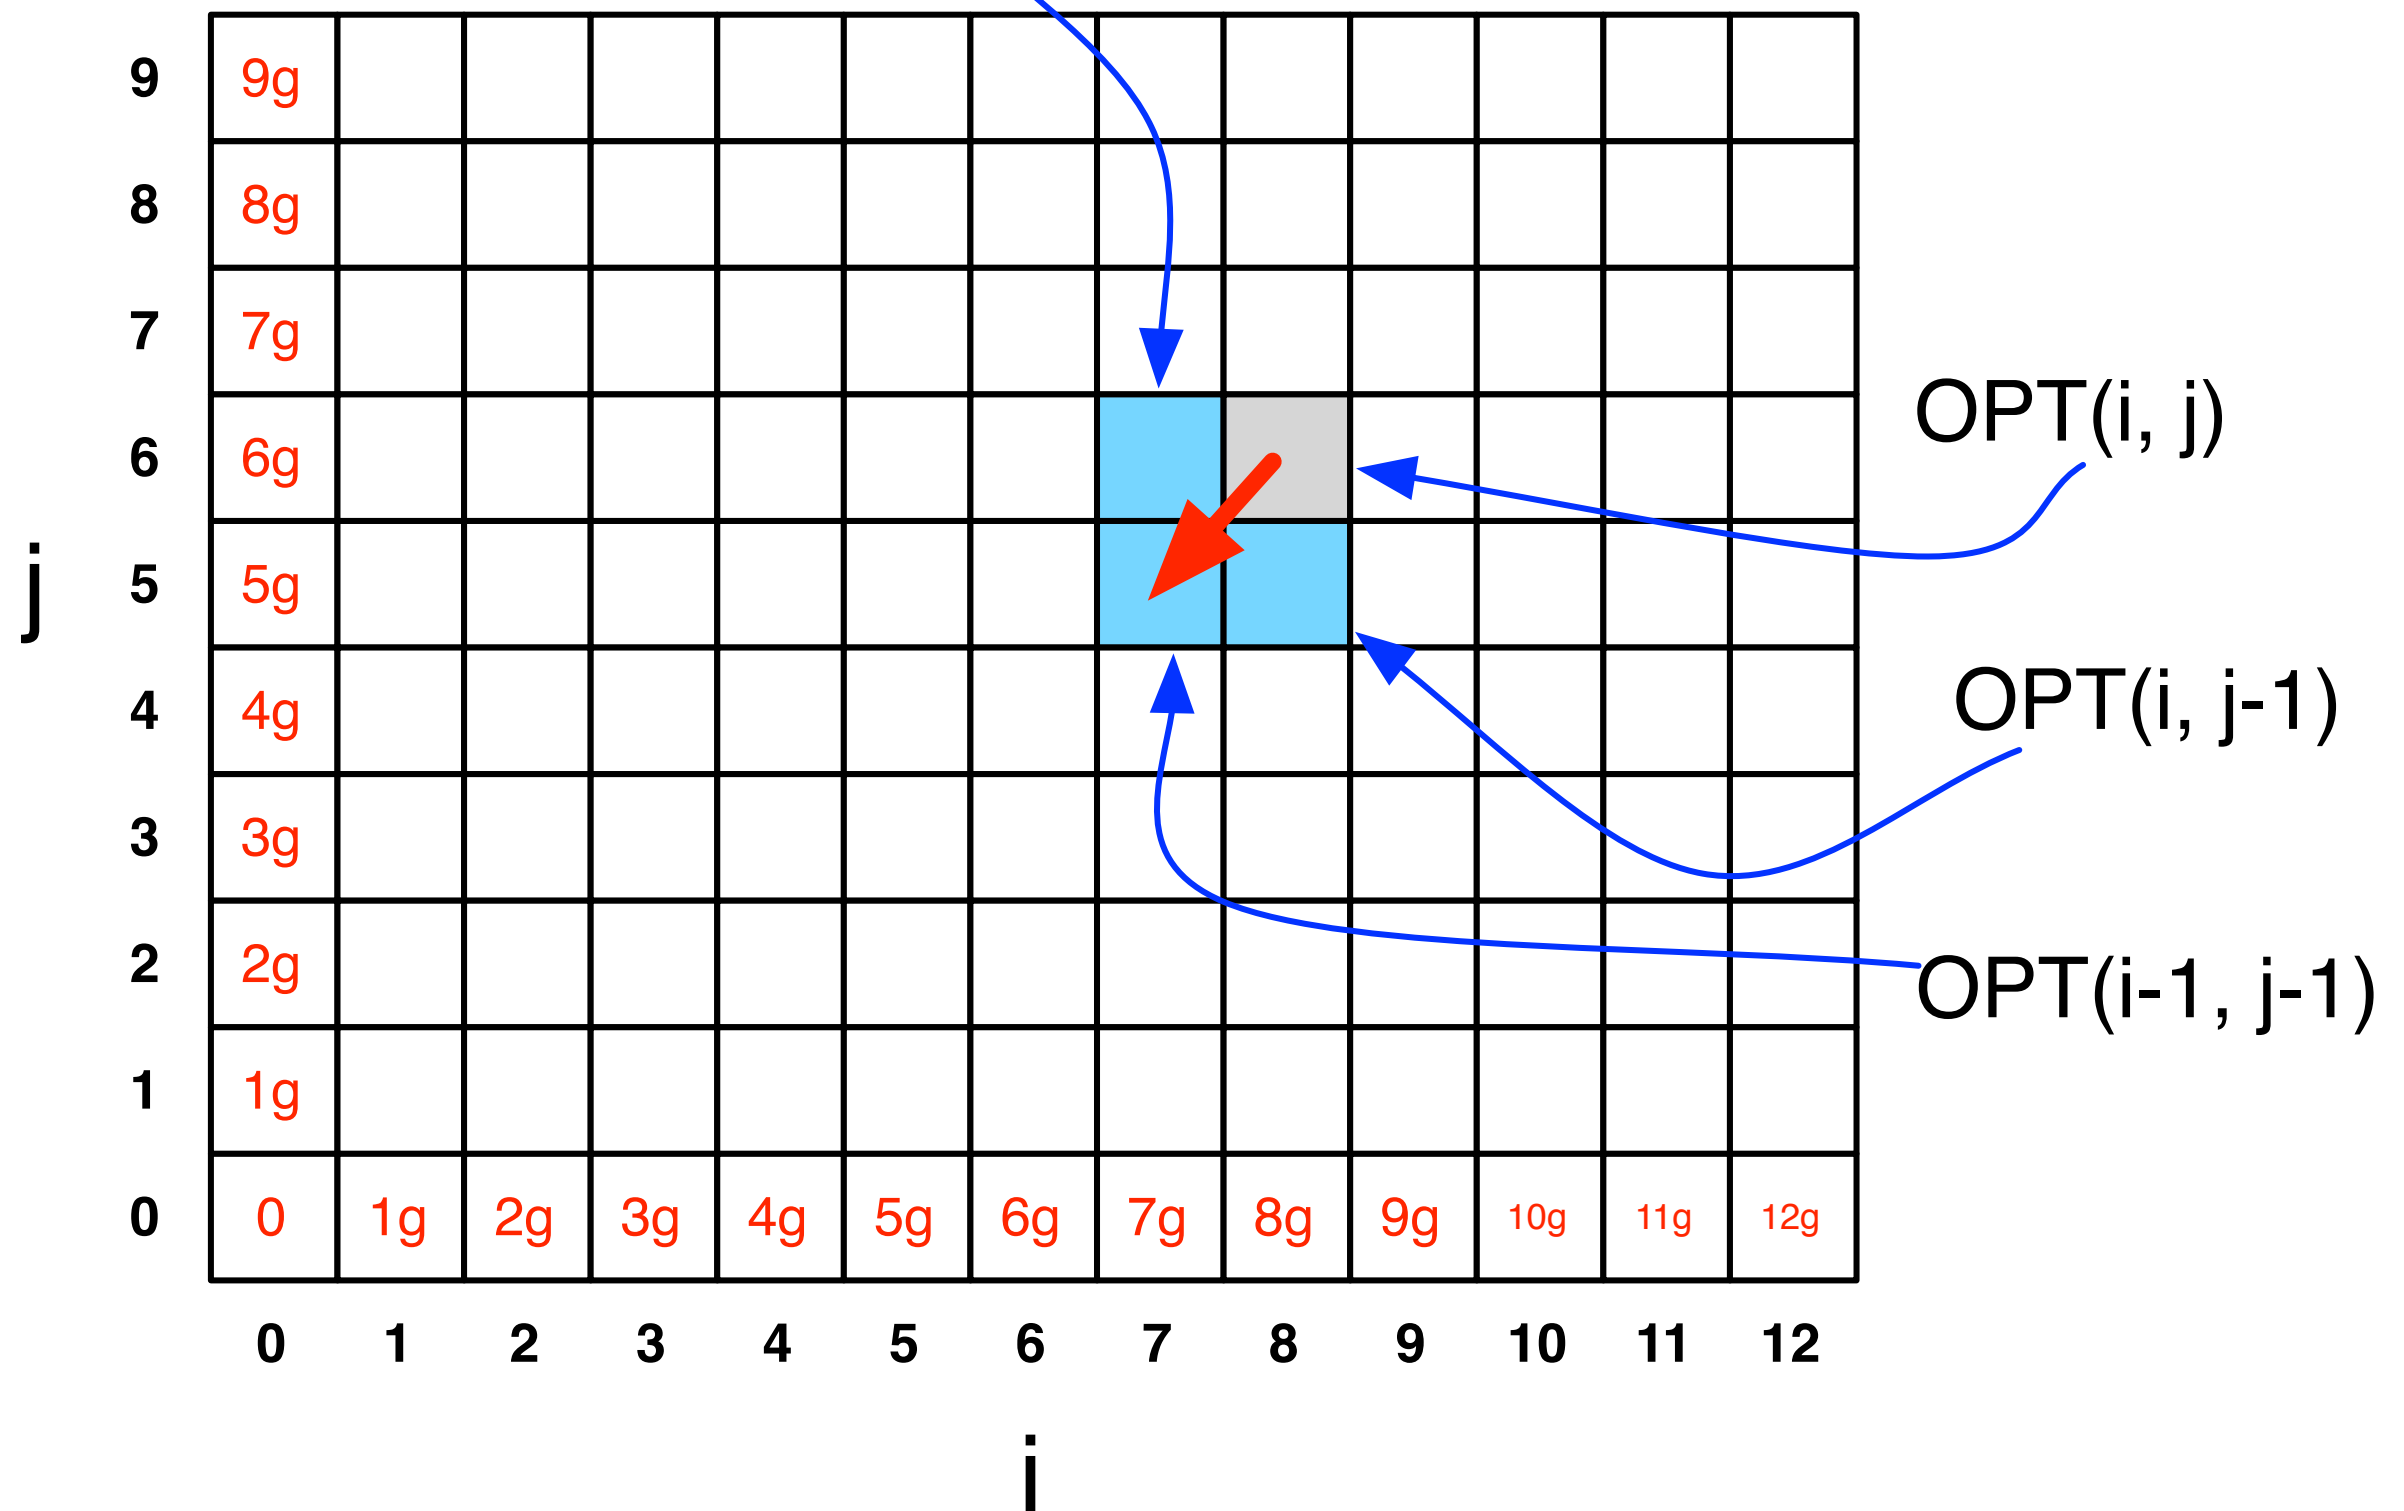
Running Time

Number of entries in array = $O(m \times n)$, where m and n are the lengths of the 2 strings.

Filling in each entry takes constant $O(1)$ time.

Total running time is $O(mn)$.

Finding the actual alignment

$$\text{OPT}(i-1, j)$$


Outputting the Alignment

Build the alignment from right to left.

ACGT

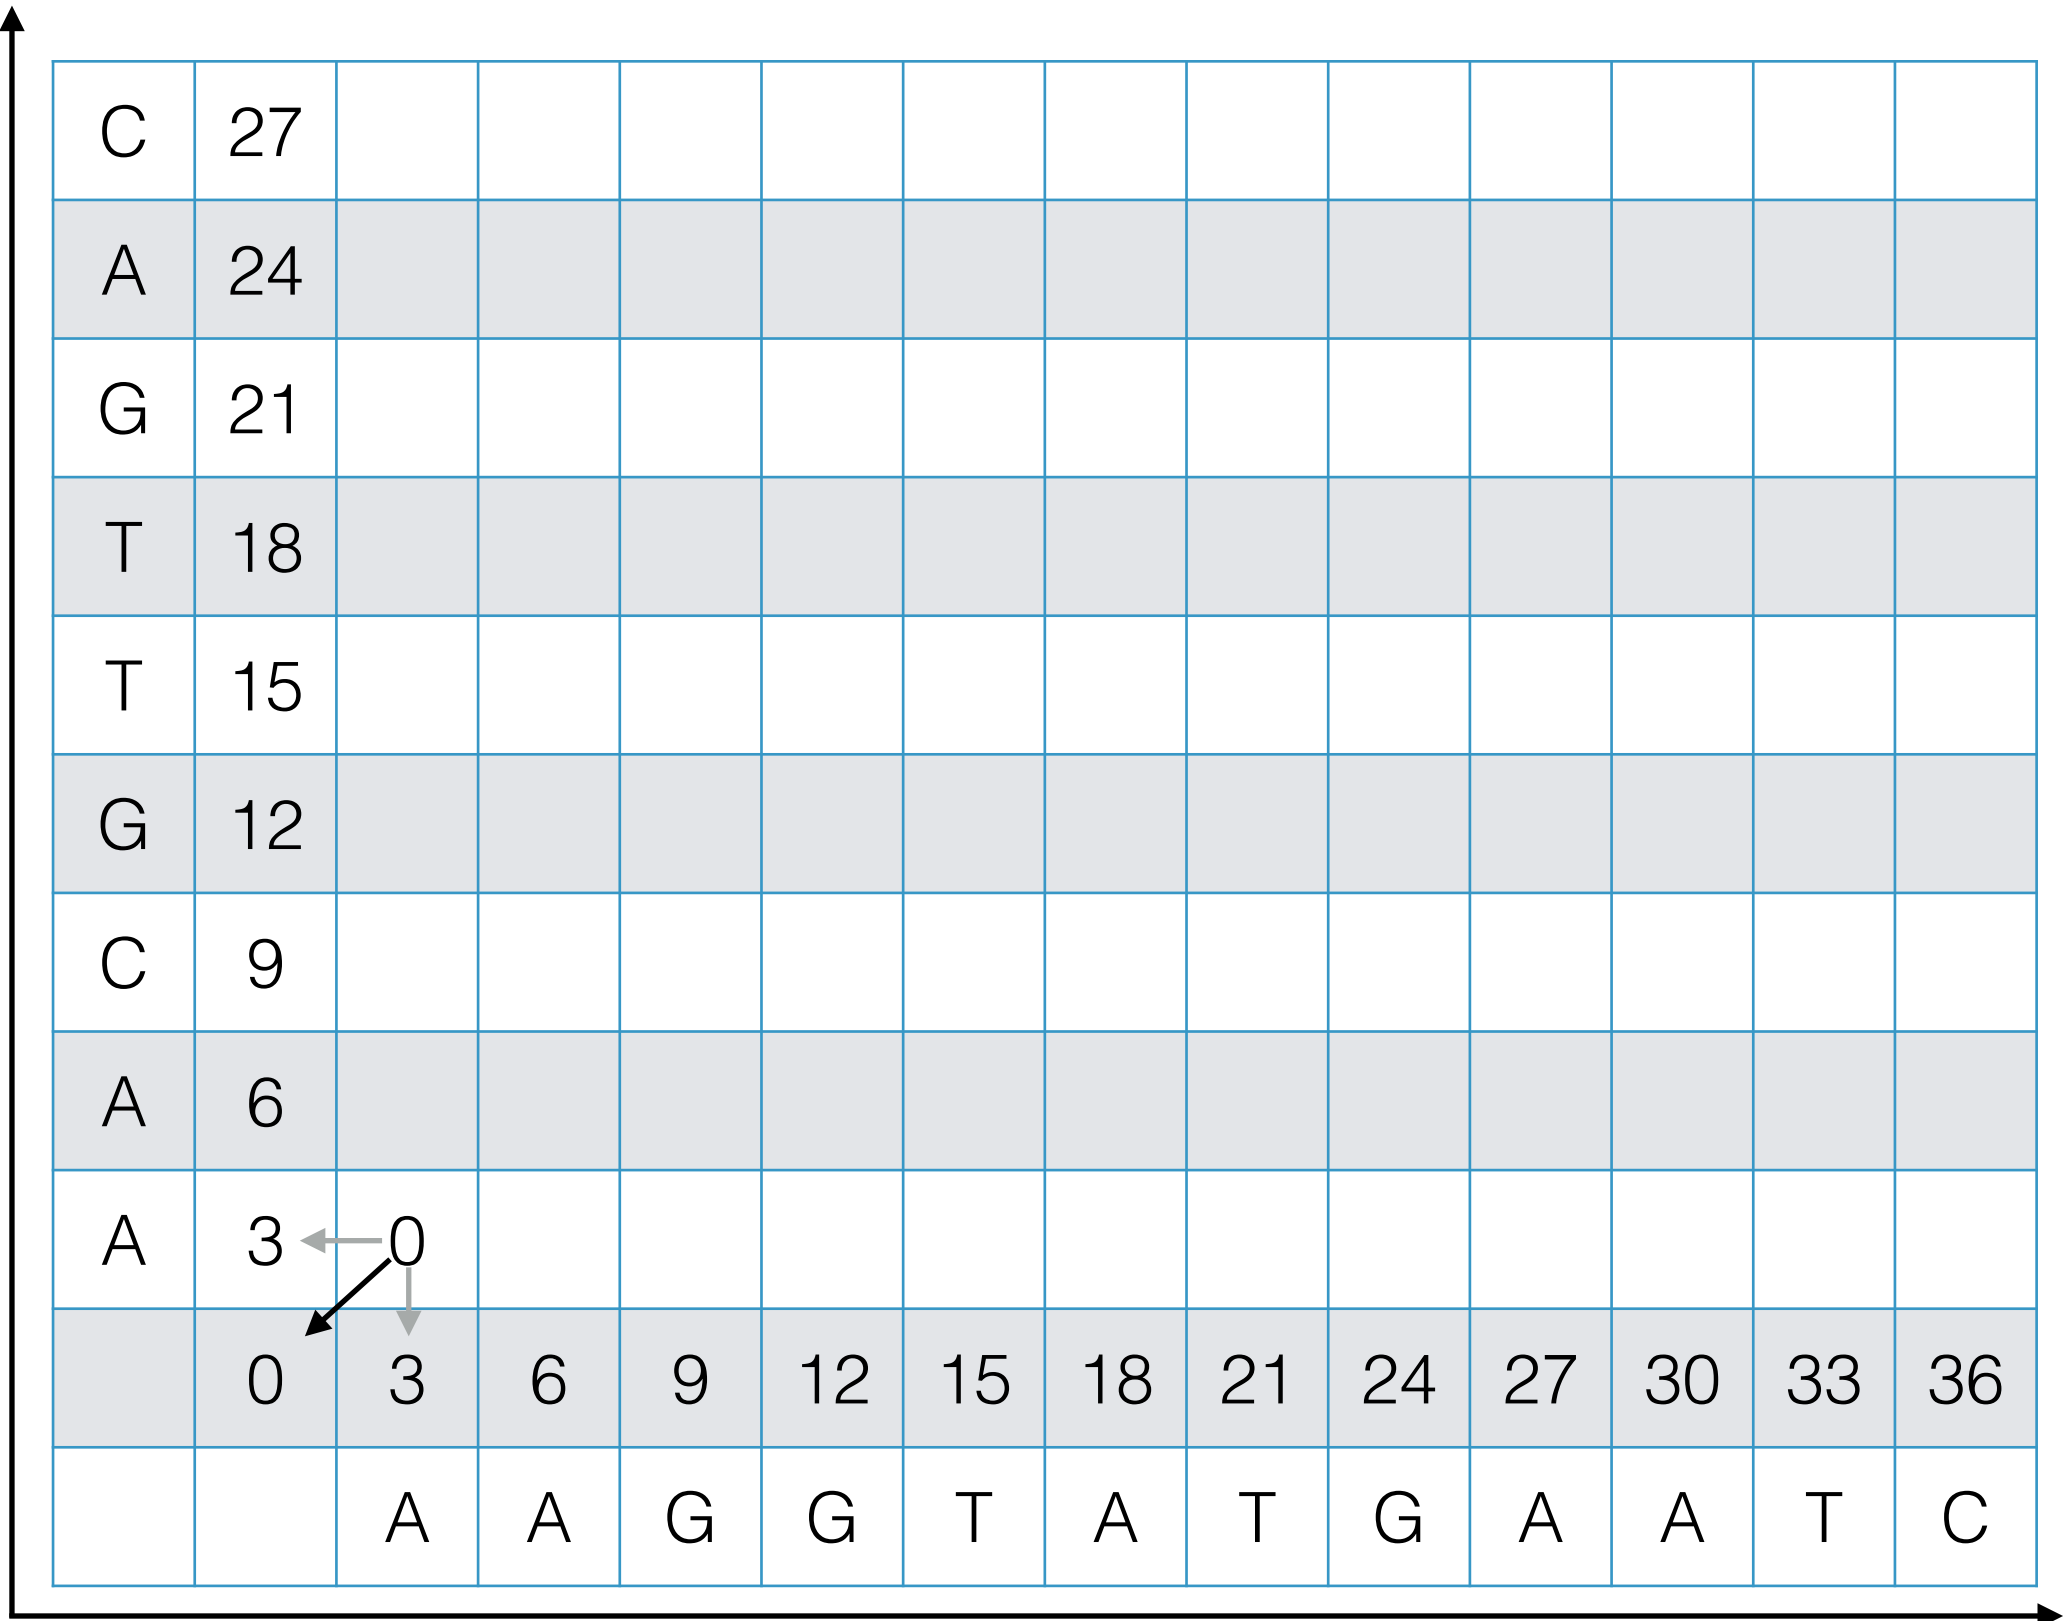
A-GA

Follow the backtrack pointers starting from entry (n,m) .

- If you follow a diagonal pointer, add both characters to the alignment,
- If you follow a left pointer, add a gap to the y-axis string and add the x-axis character
- If you follow a down pointer, add the y-axis character and add a gap to the x-axis string.

gap cost = 3
mismatch cost = 1

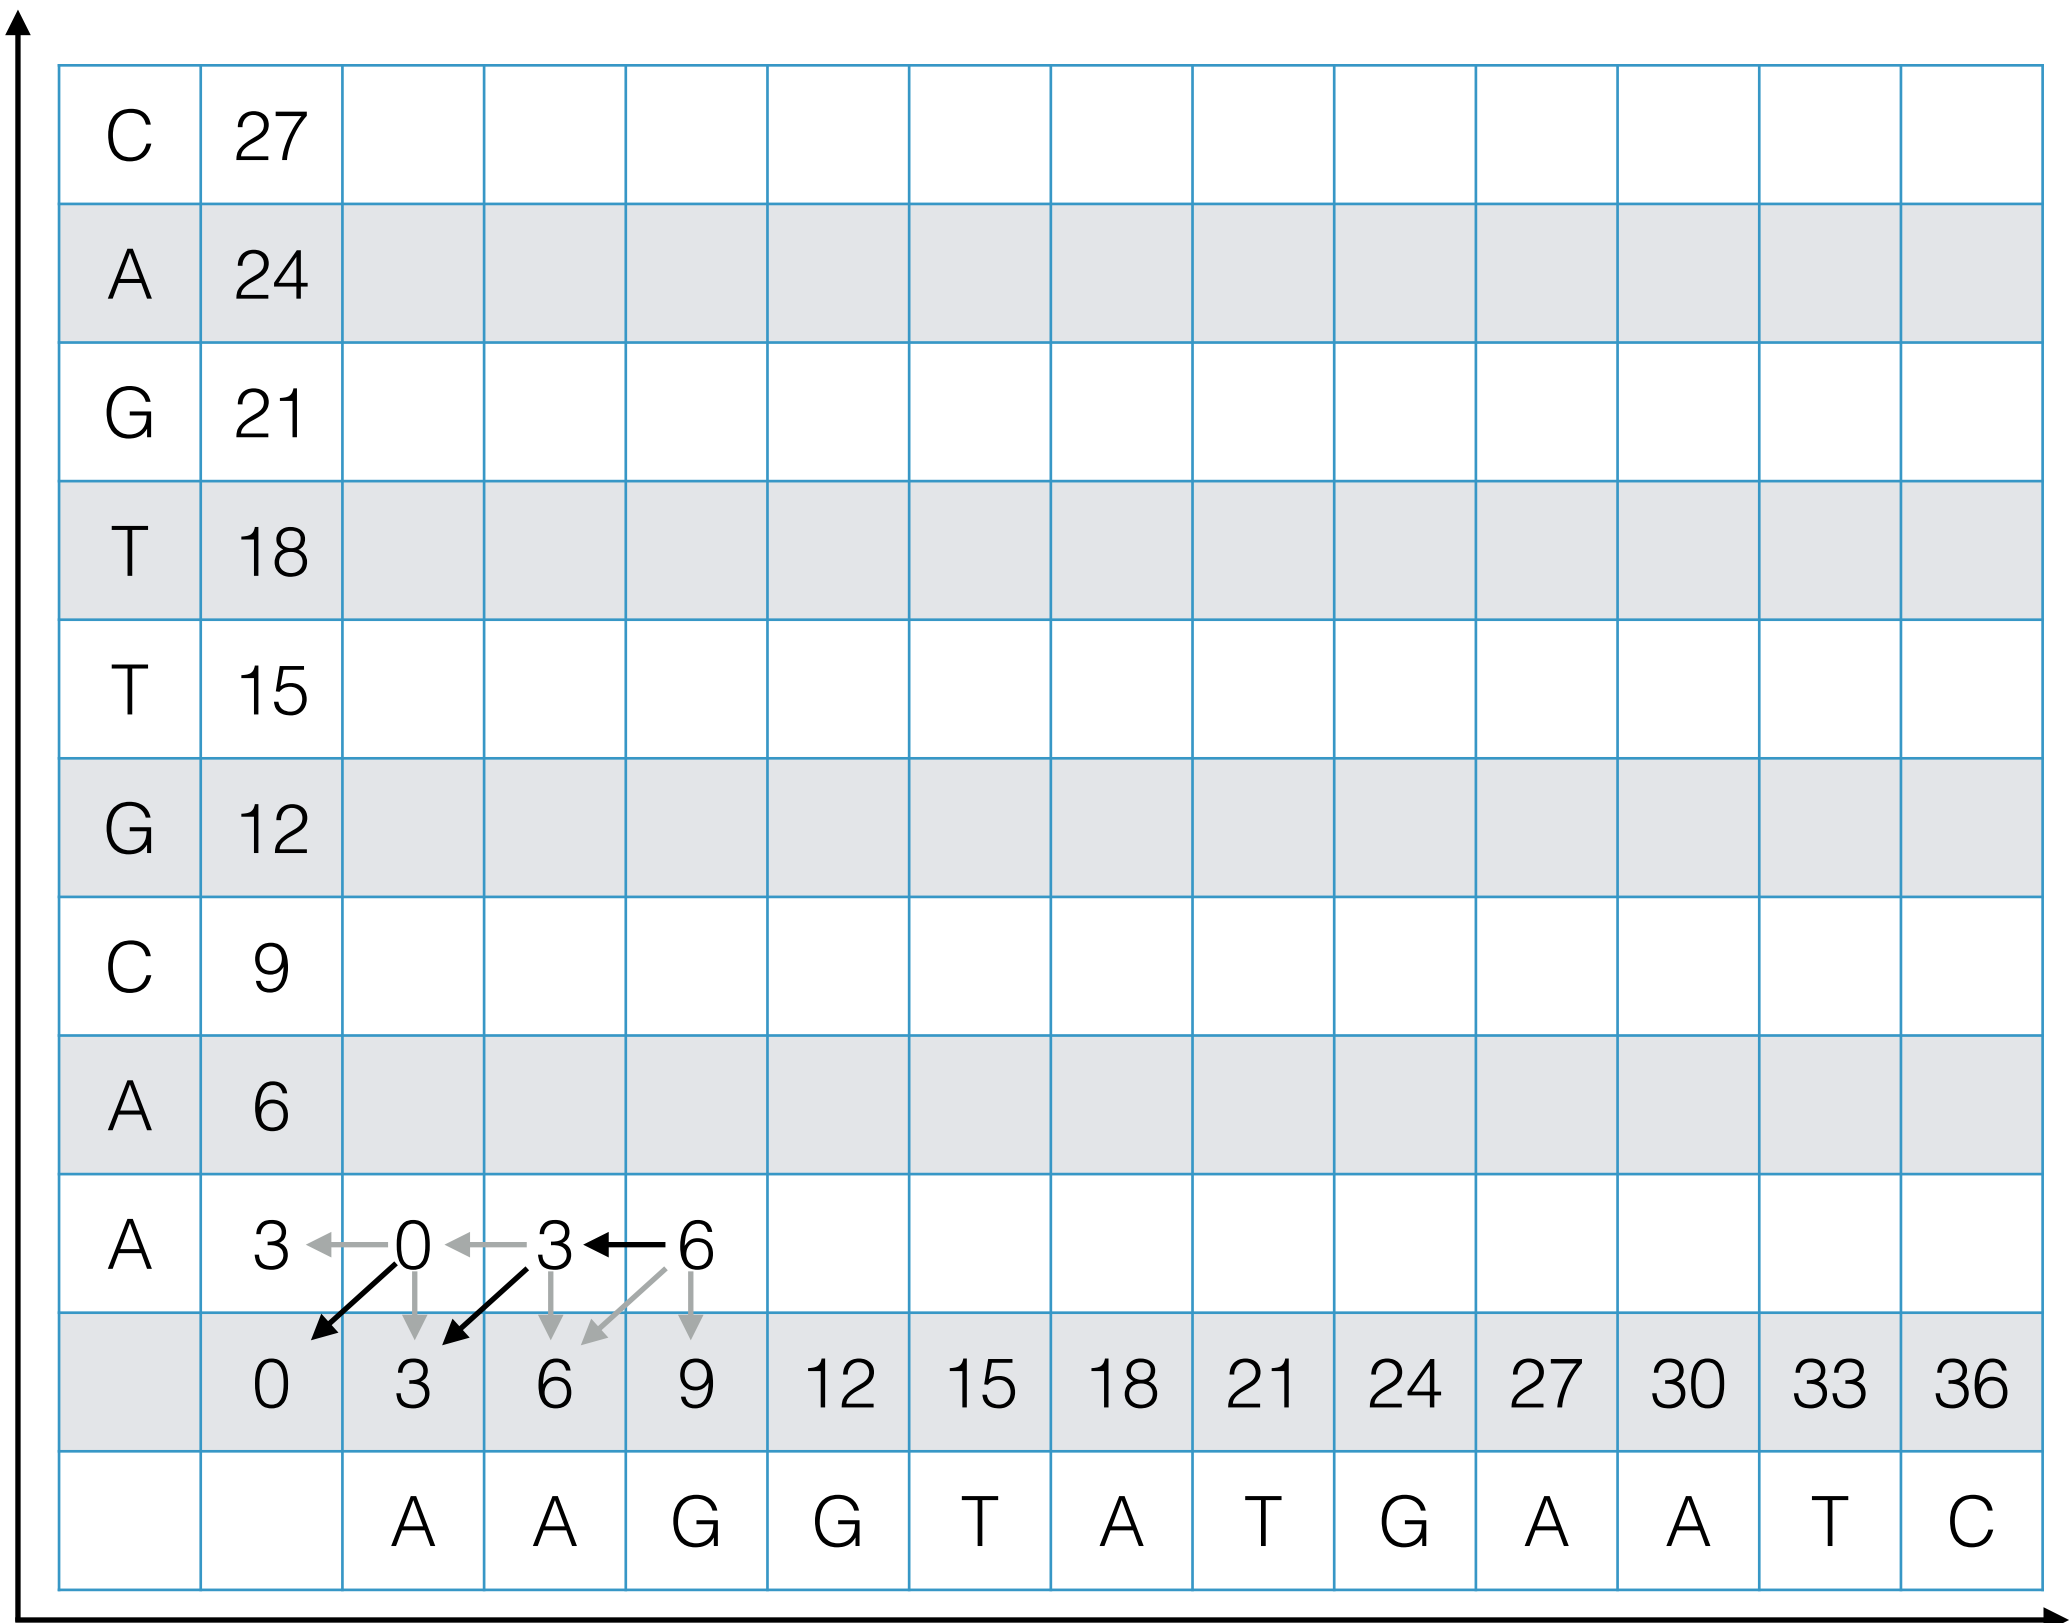
Example



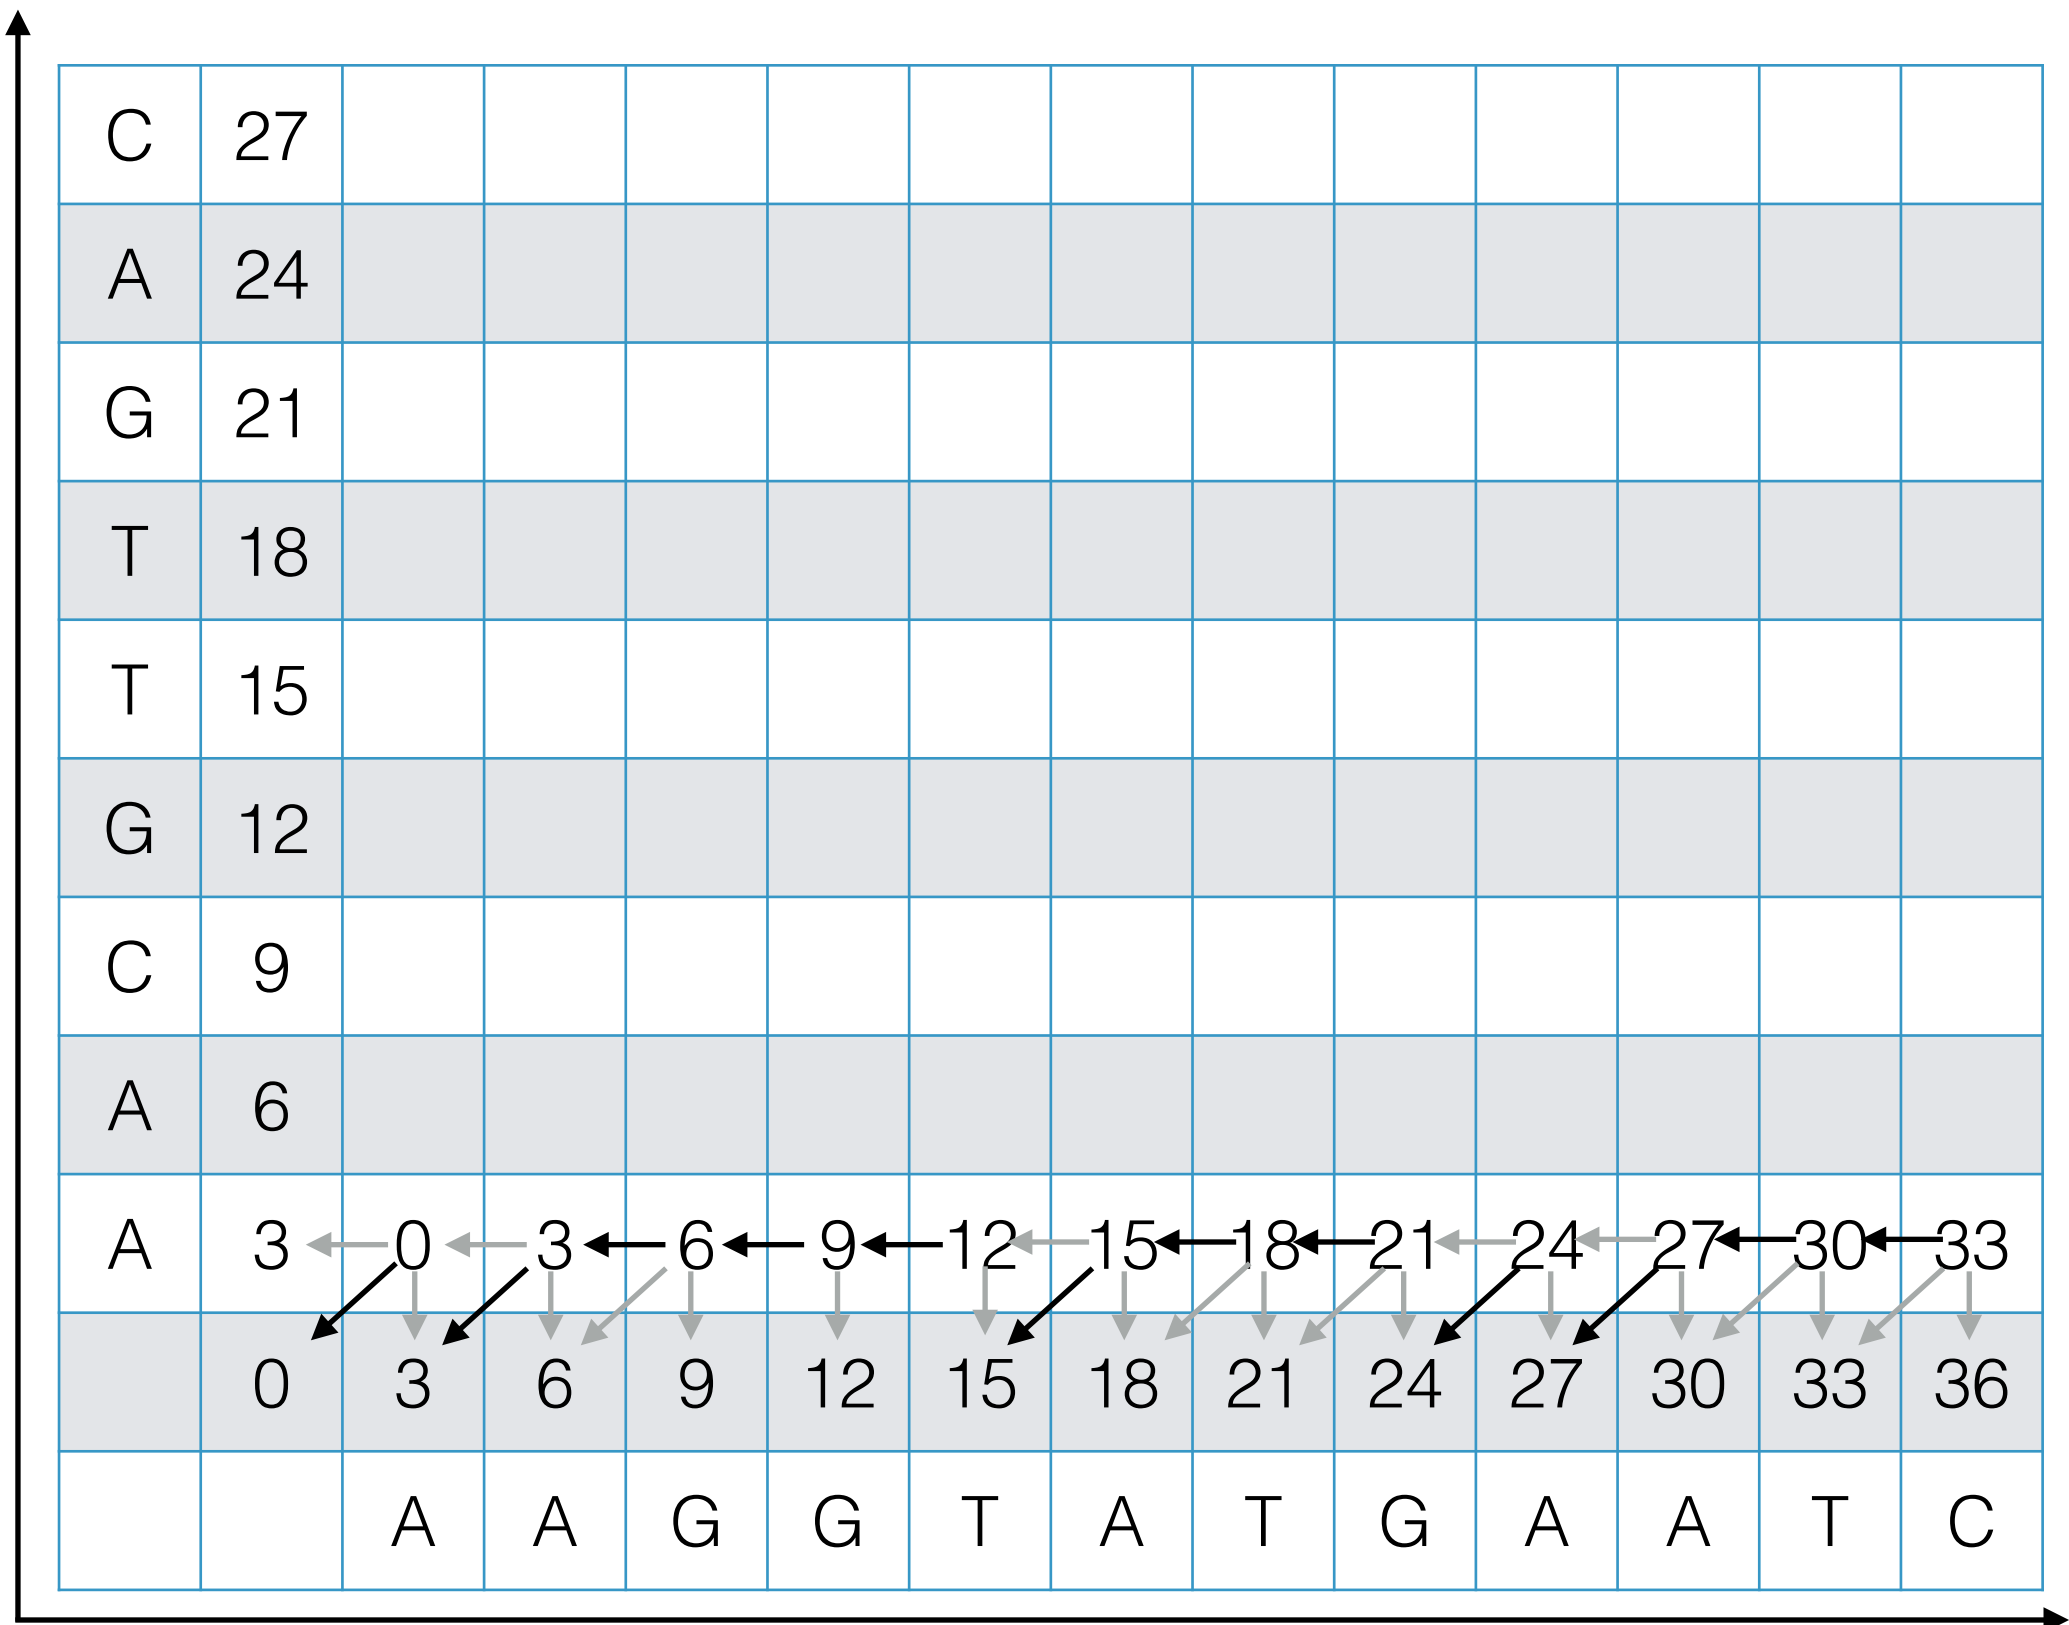
Example

C	27												
A	24												
G	21												
T	18												
T	15												
G	12												
C	9												
A	6												
A	3	←0	←3										
	0	3	6	9	12	15	18	21	24	27	30	33	36
		A	A	G	G	T	A	T	G	A	A	T	C

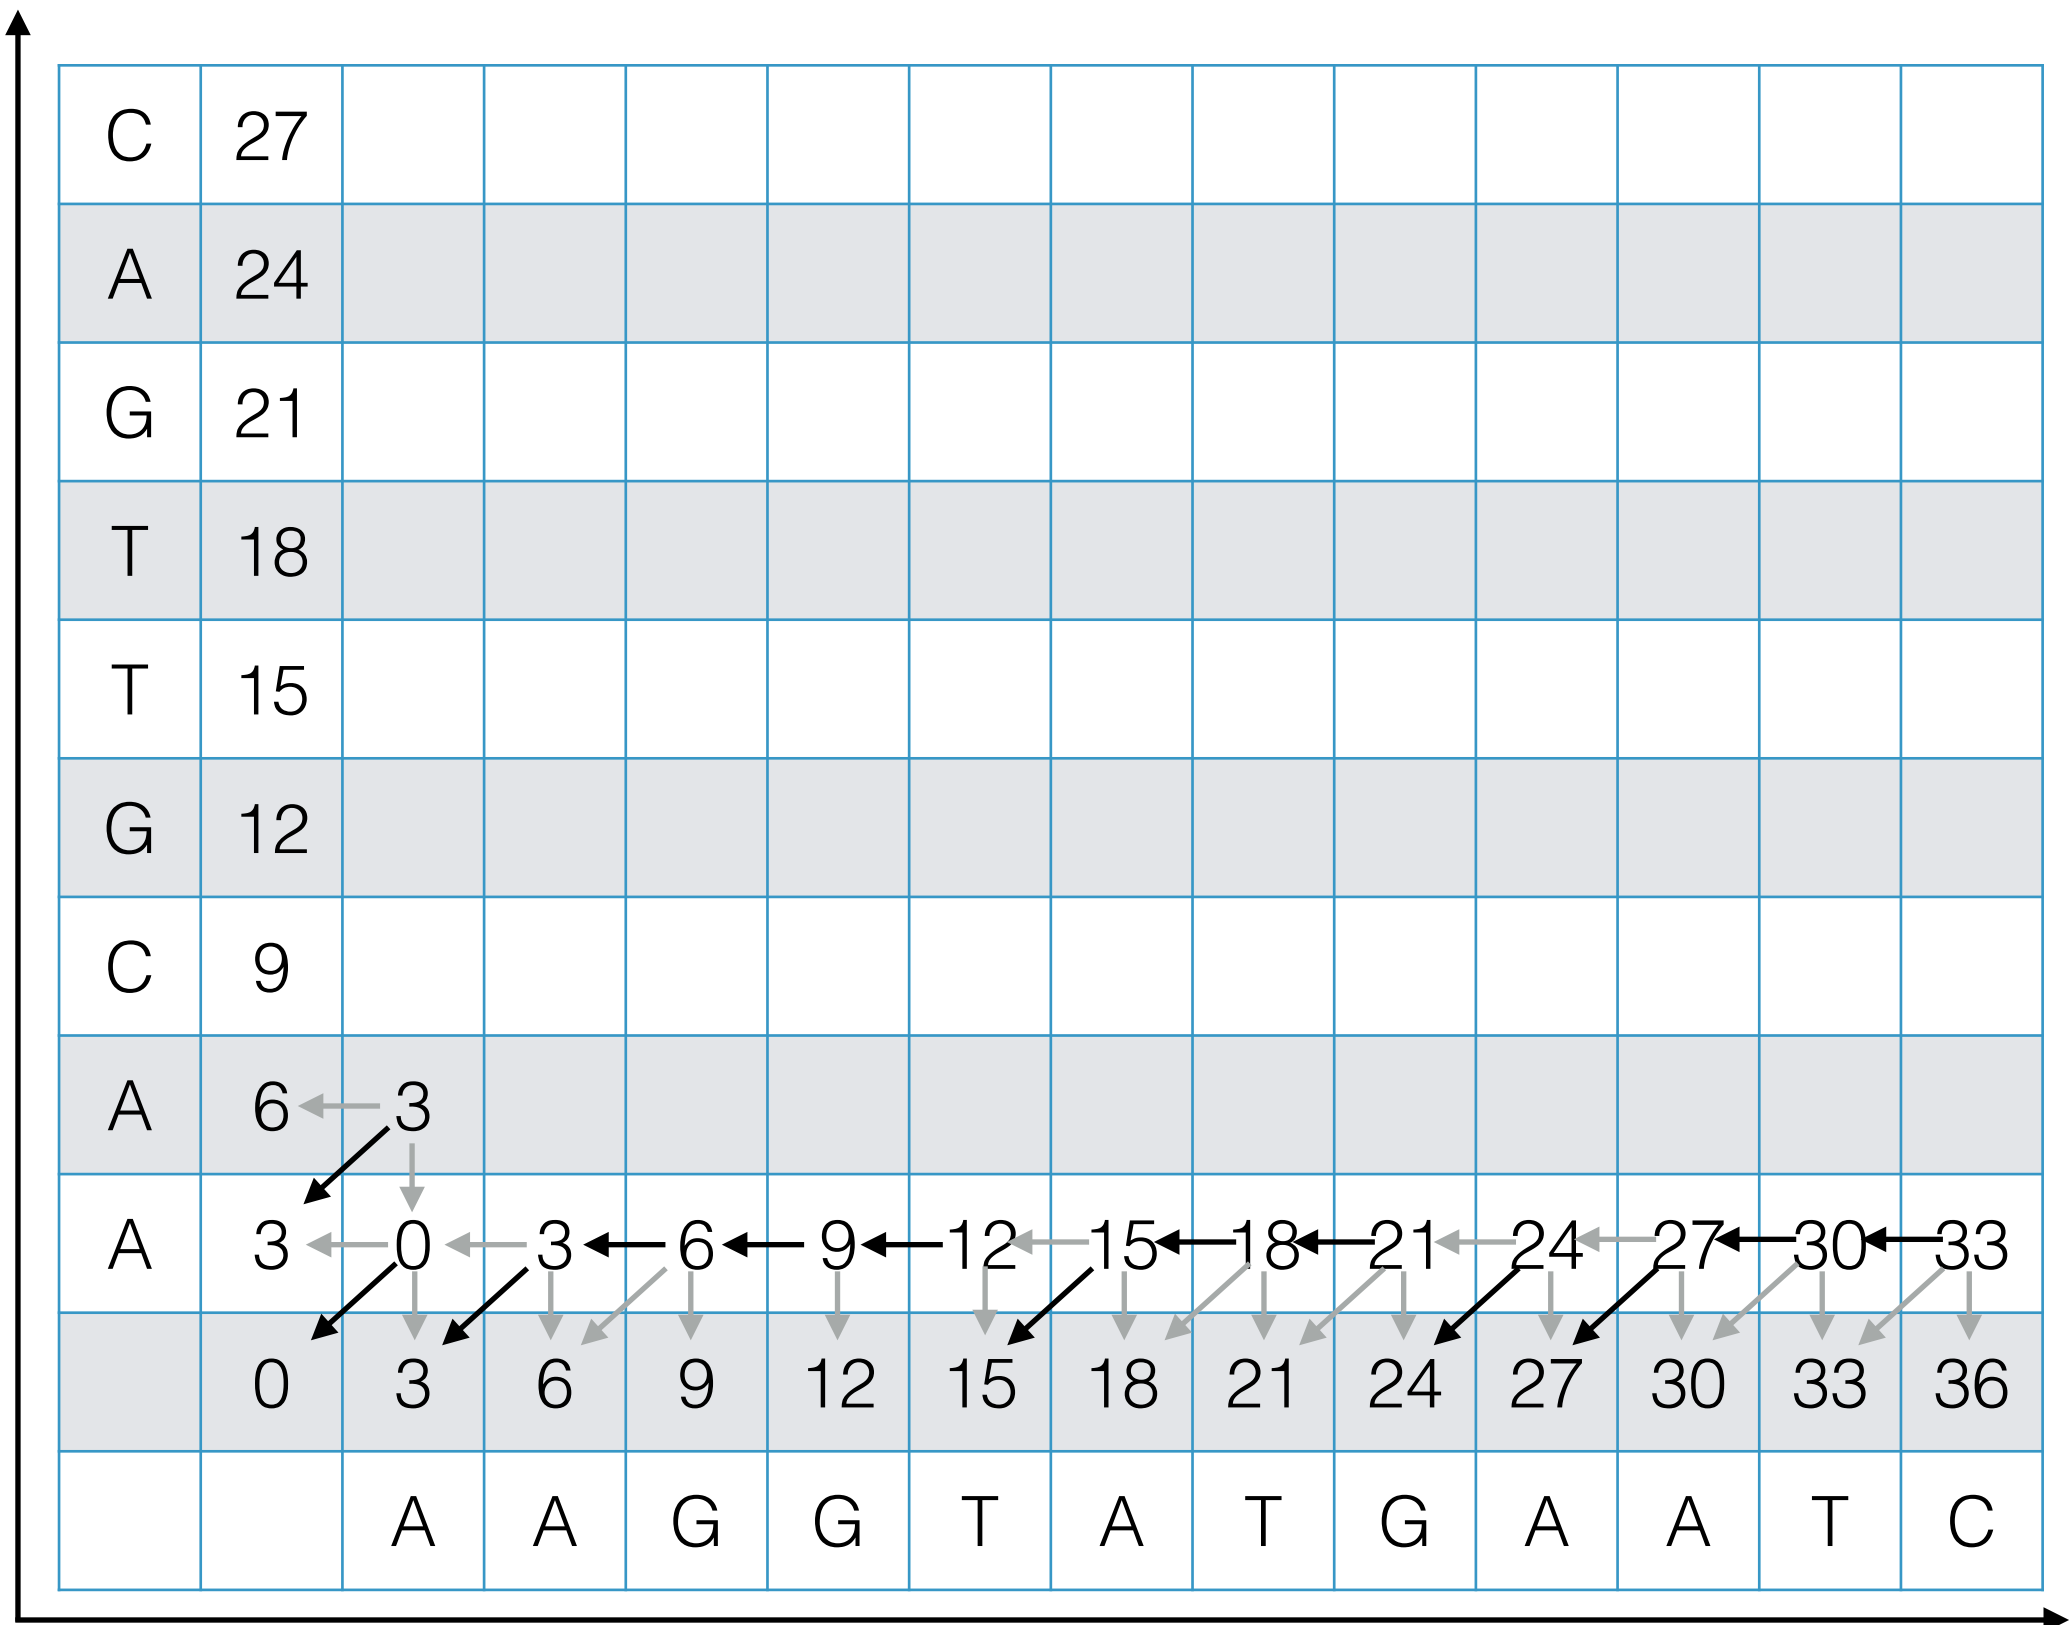
Example



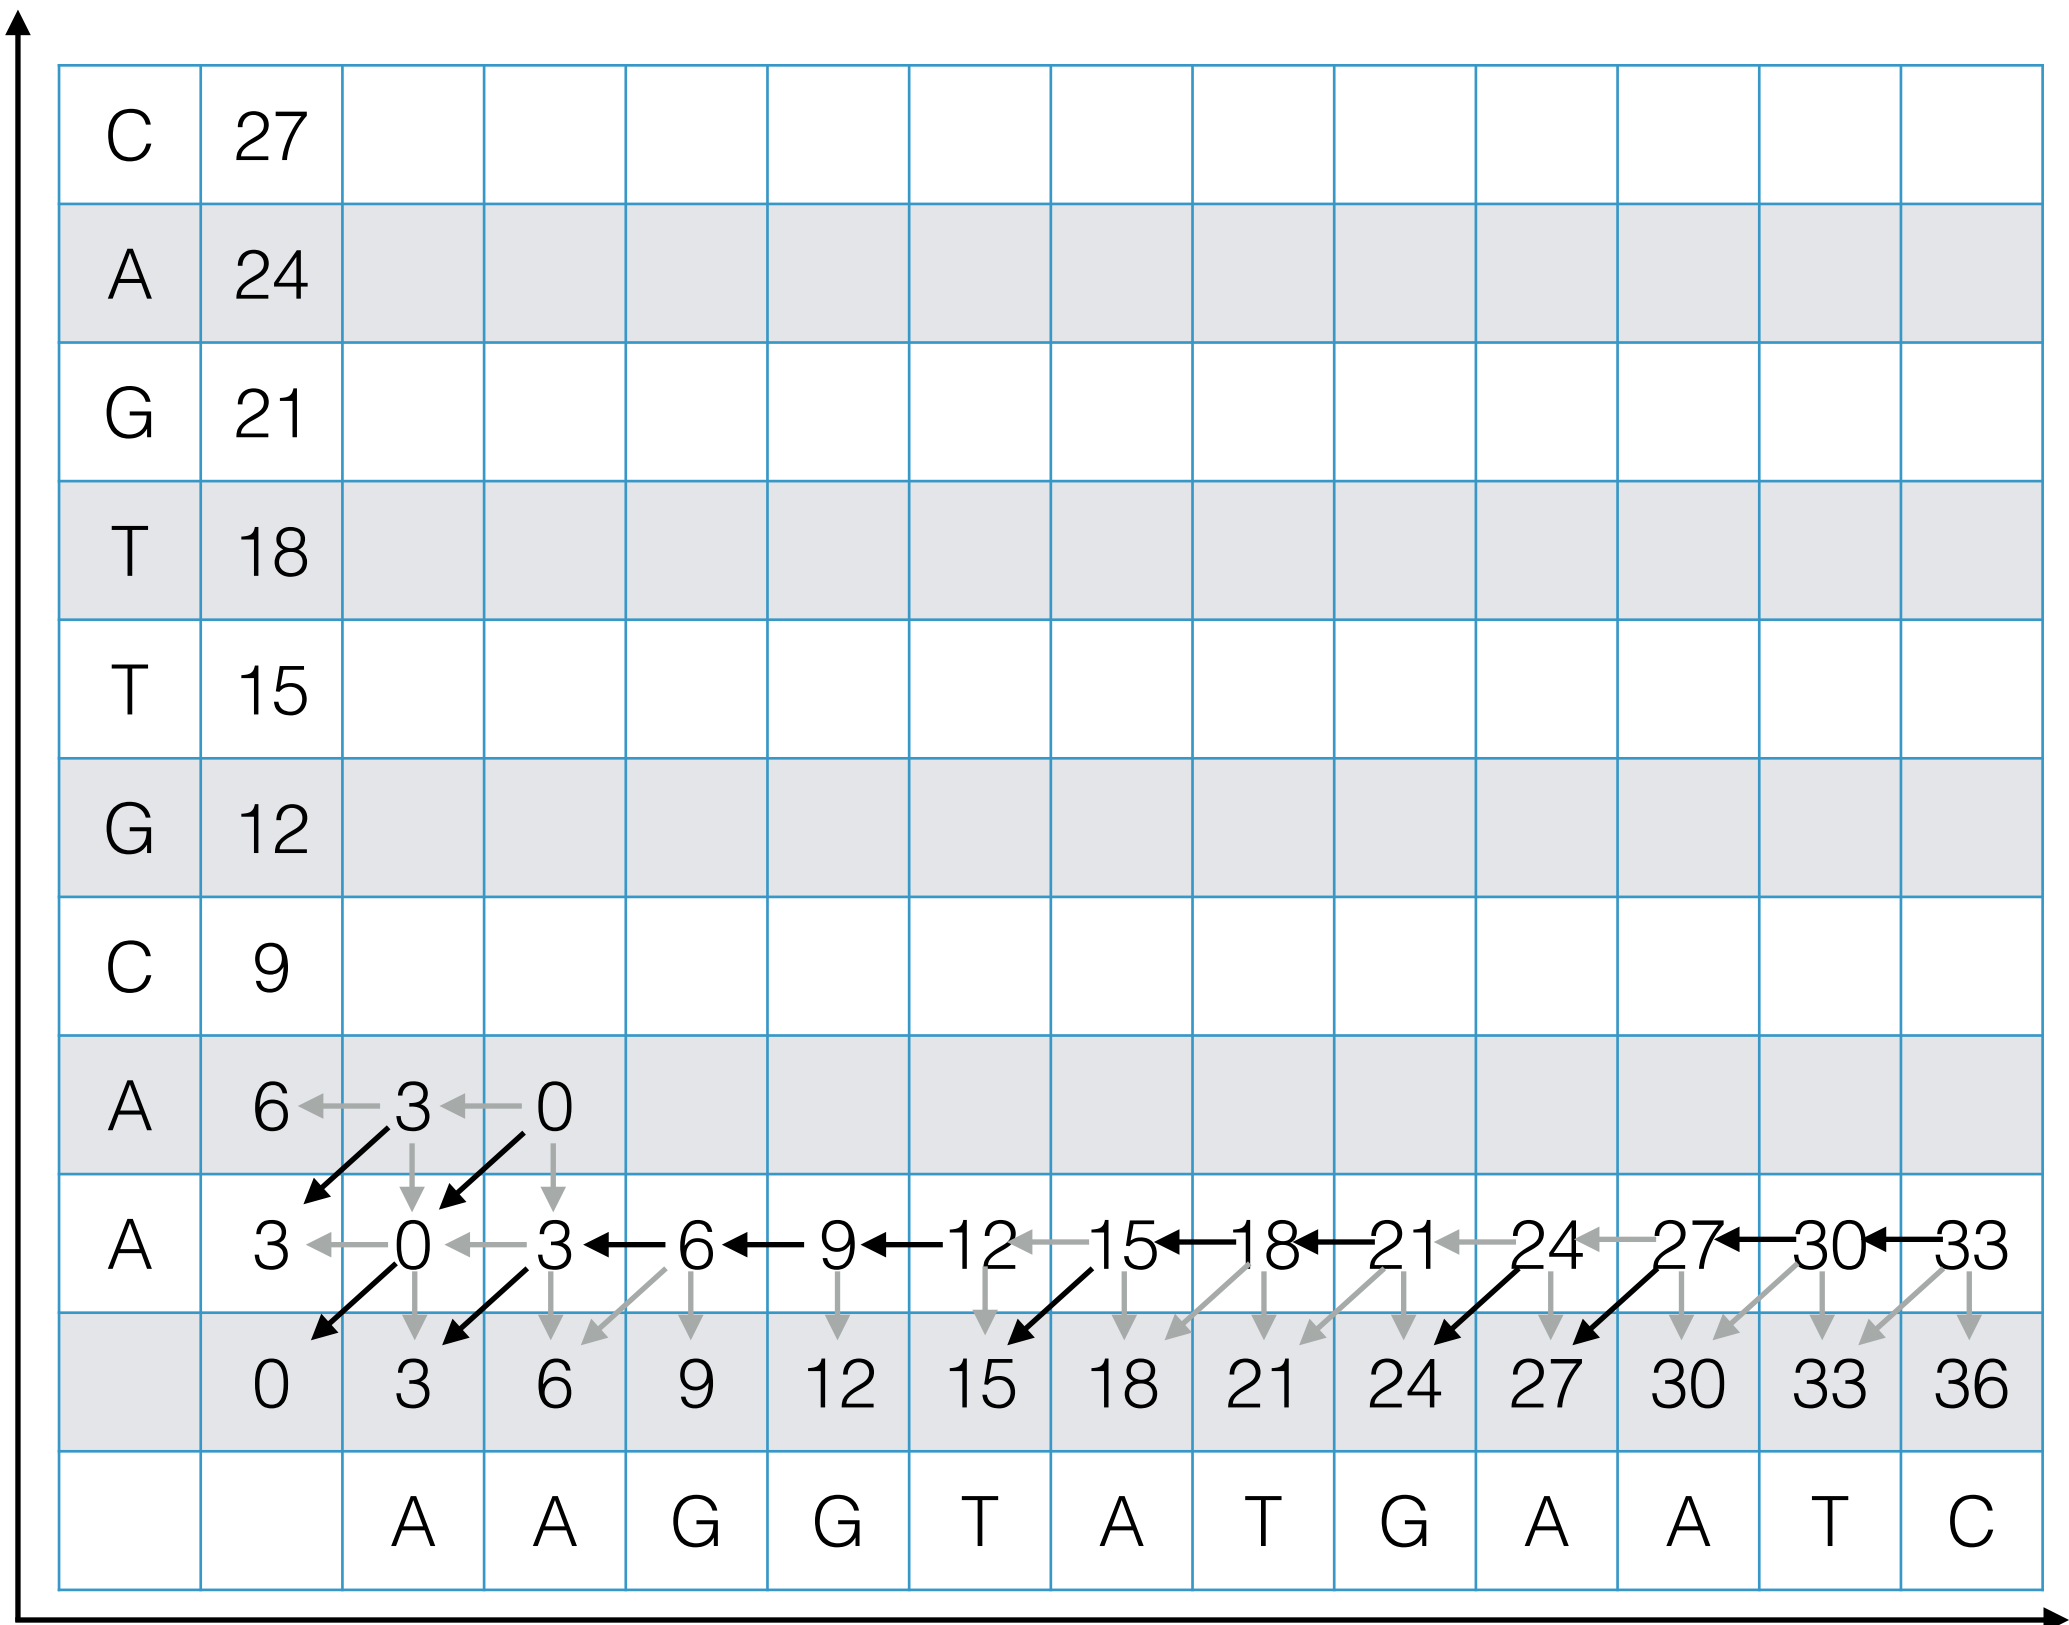
Example



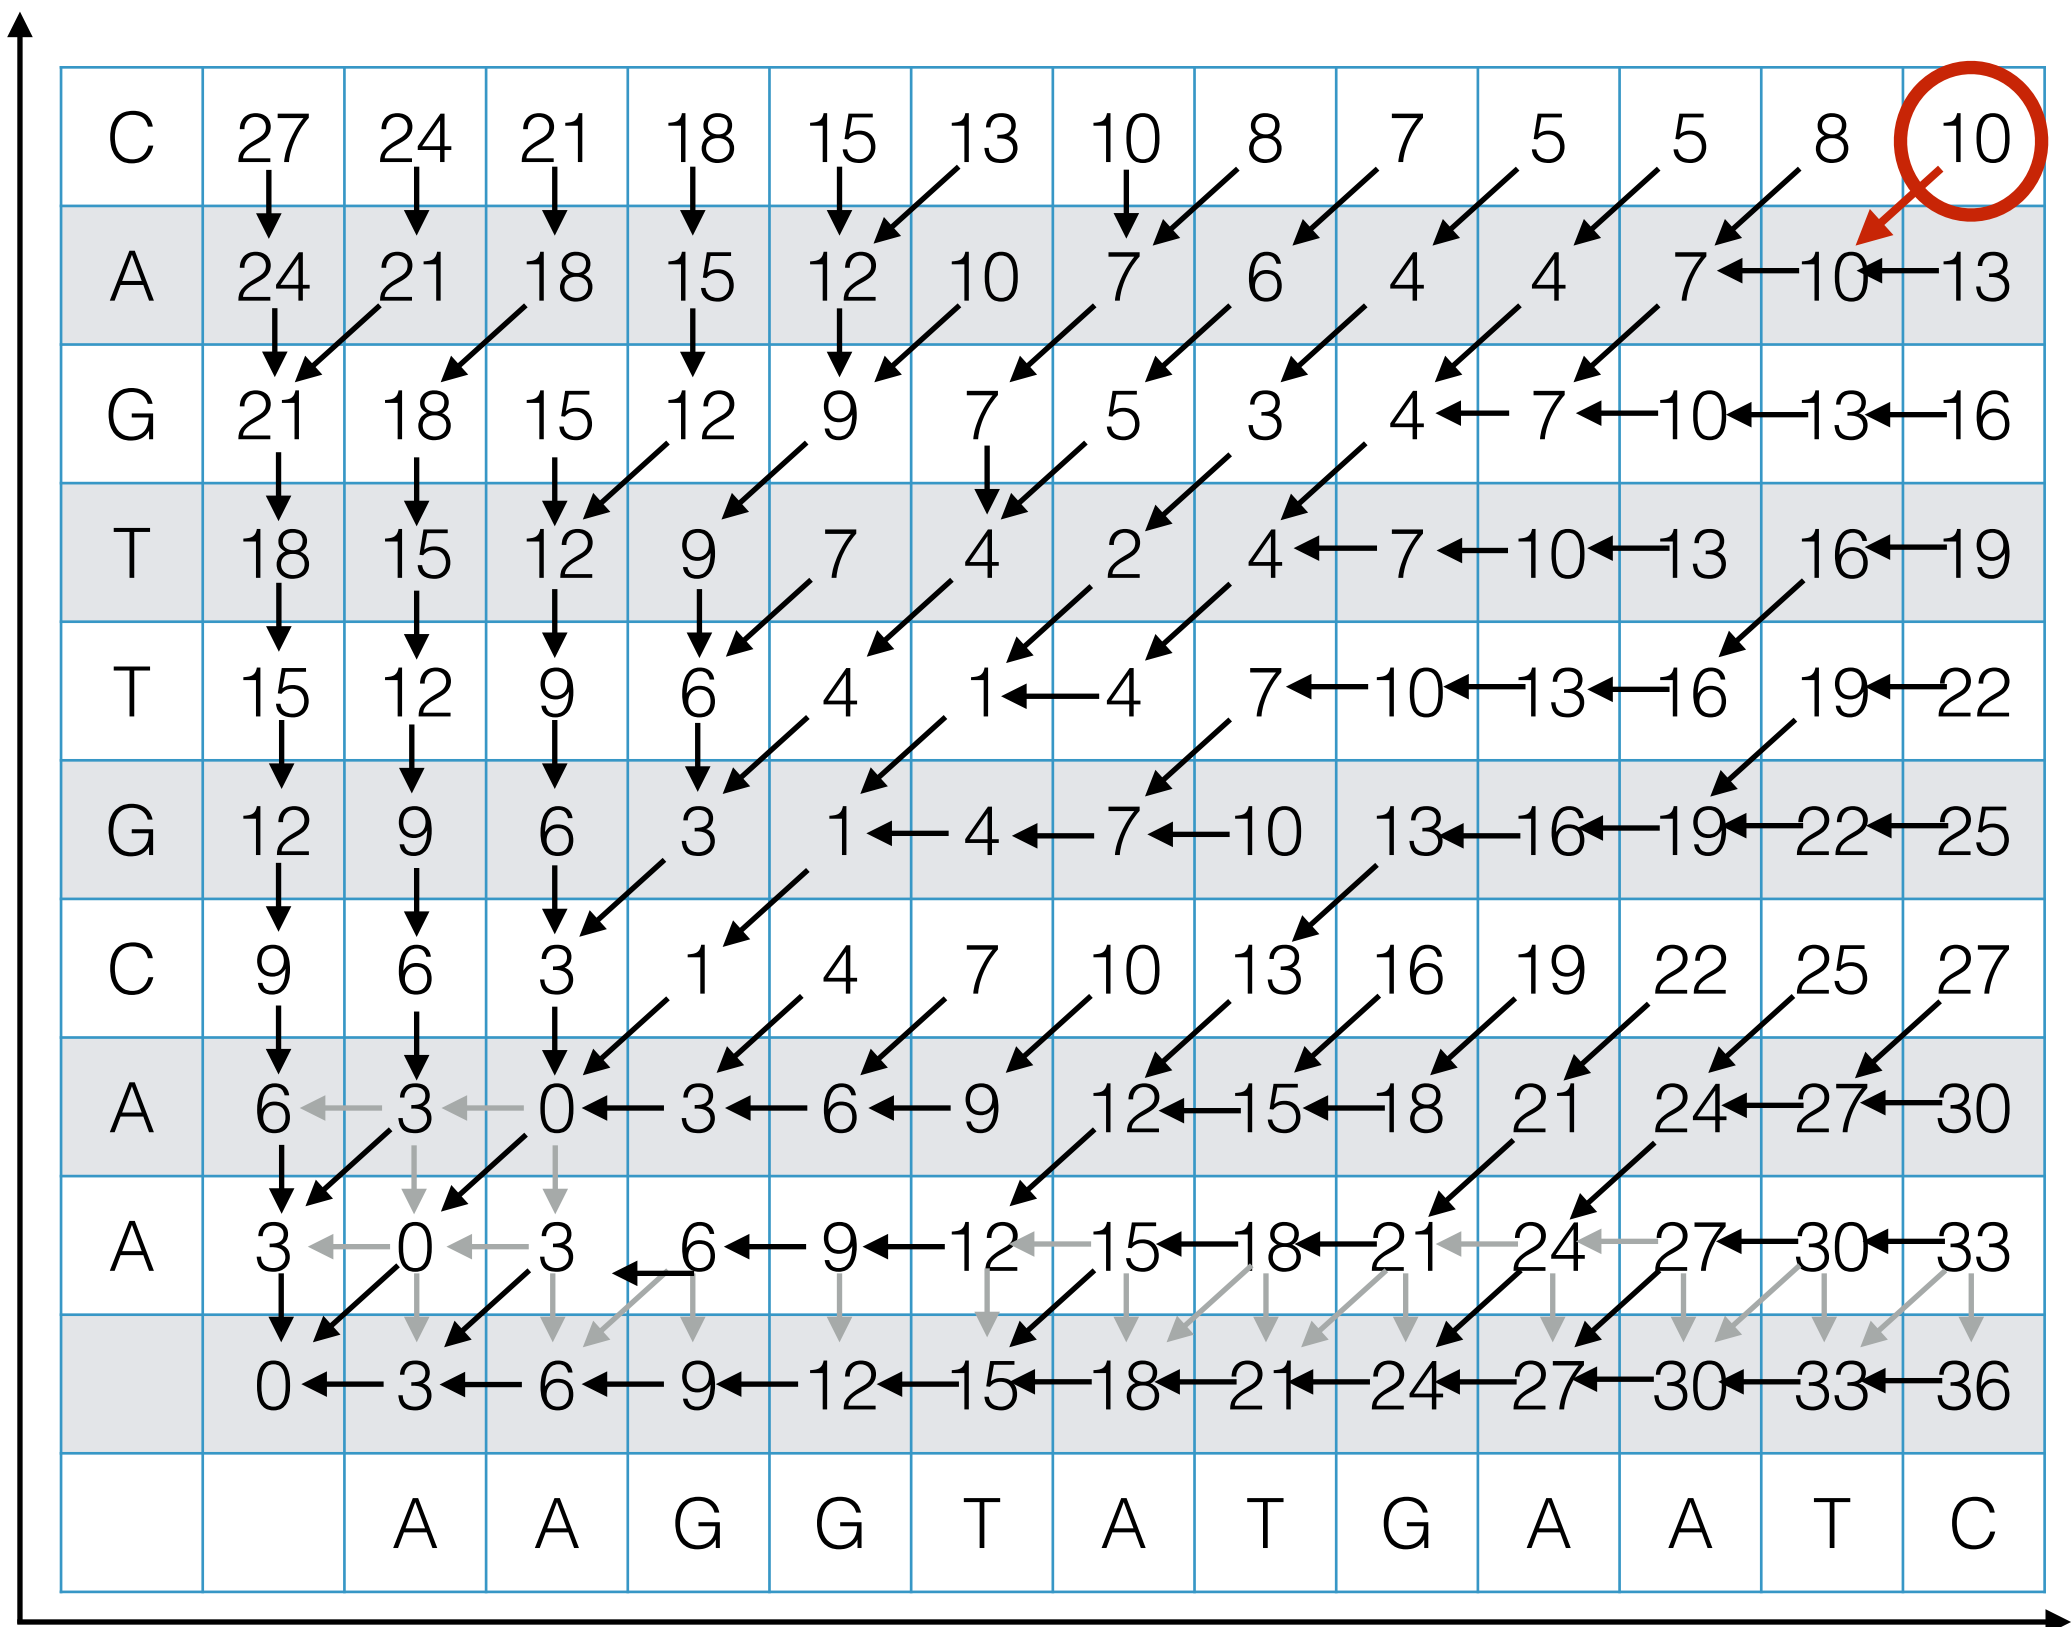
Example



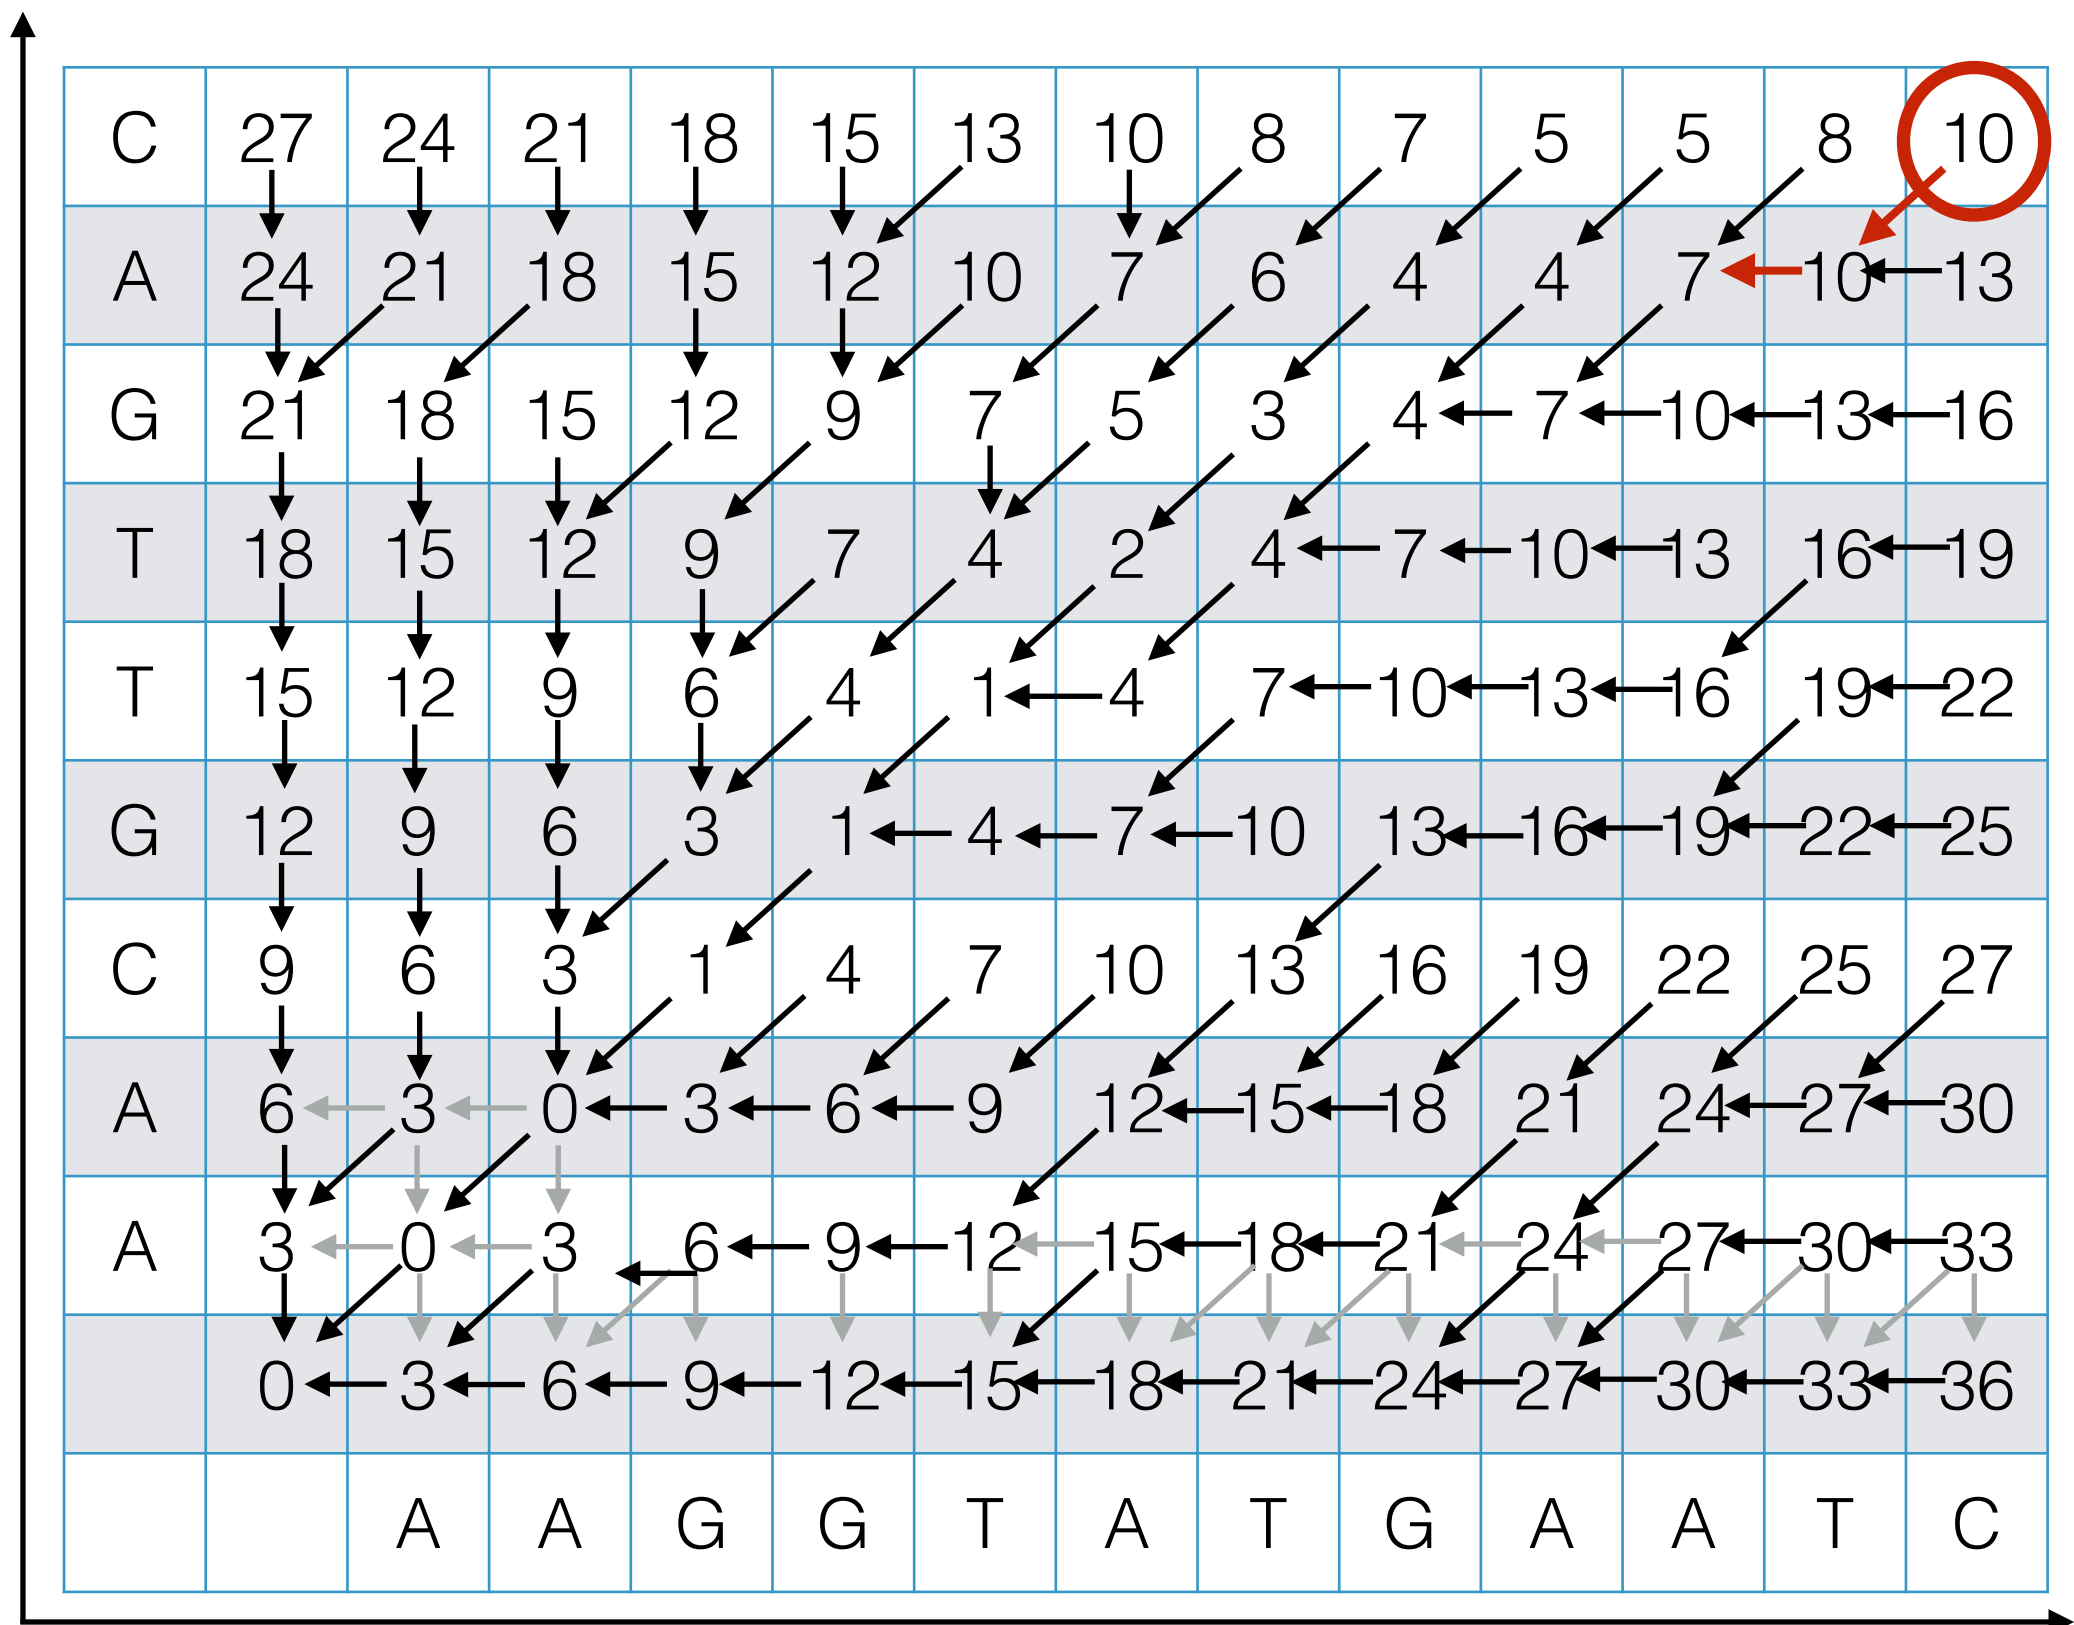
Example



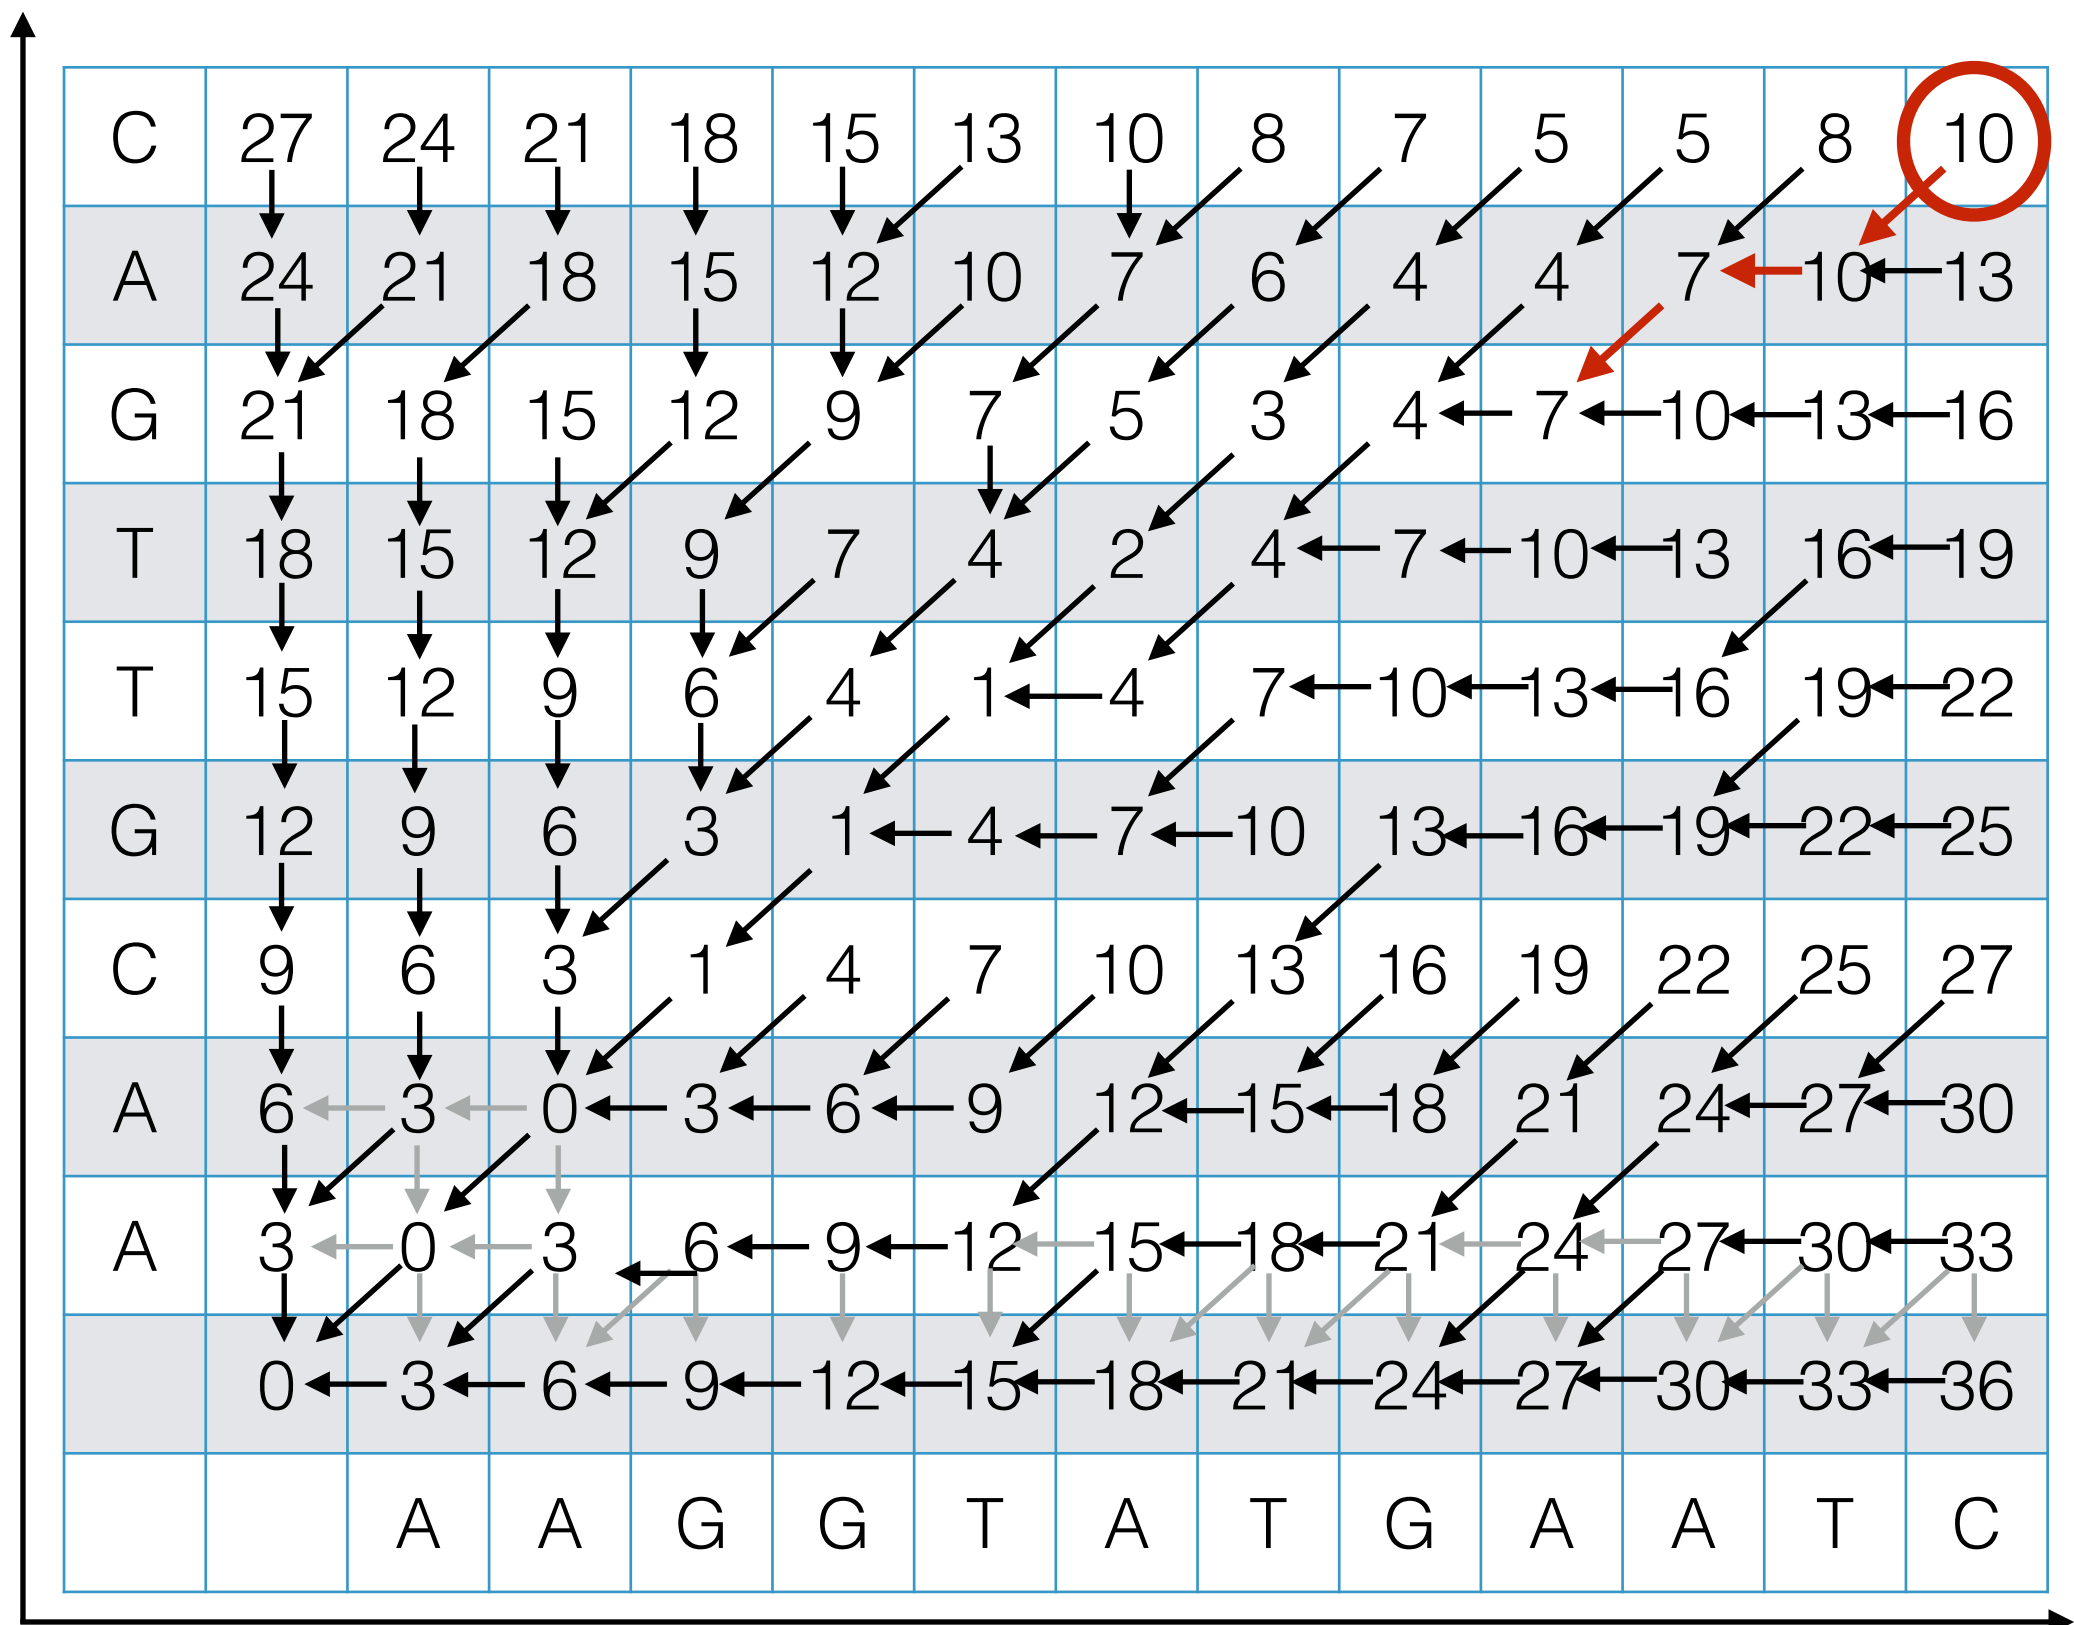
Example



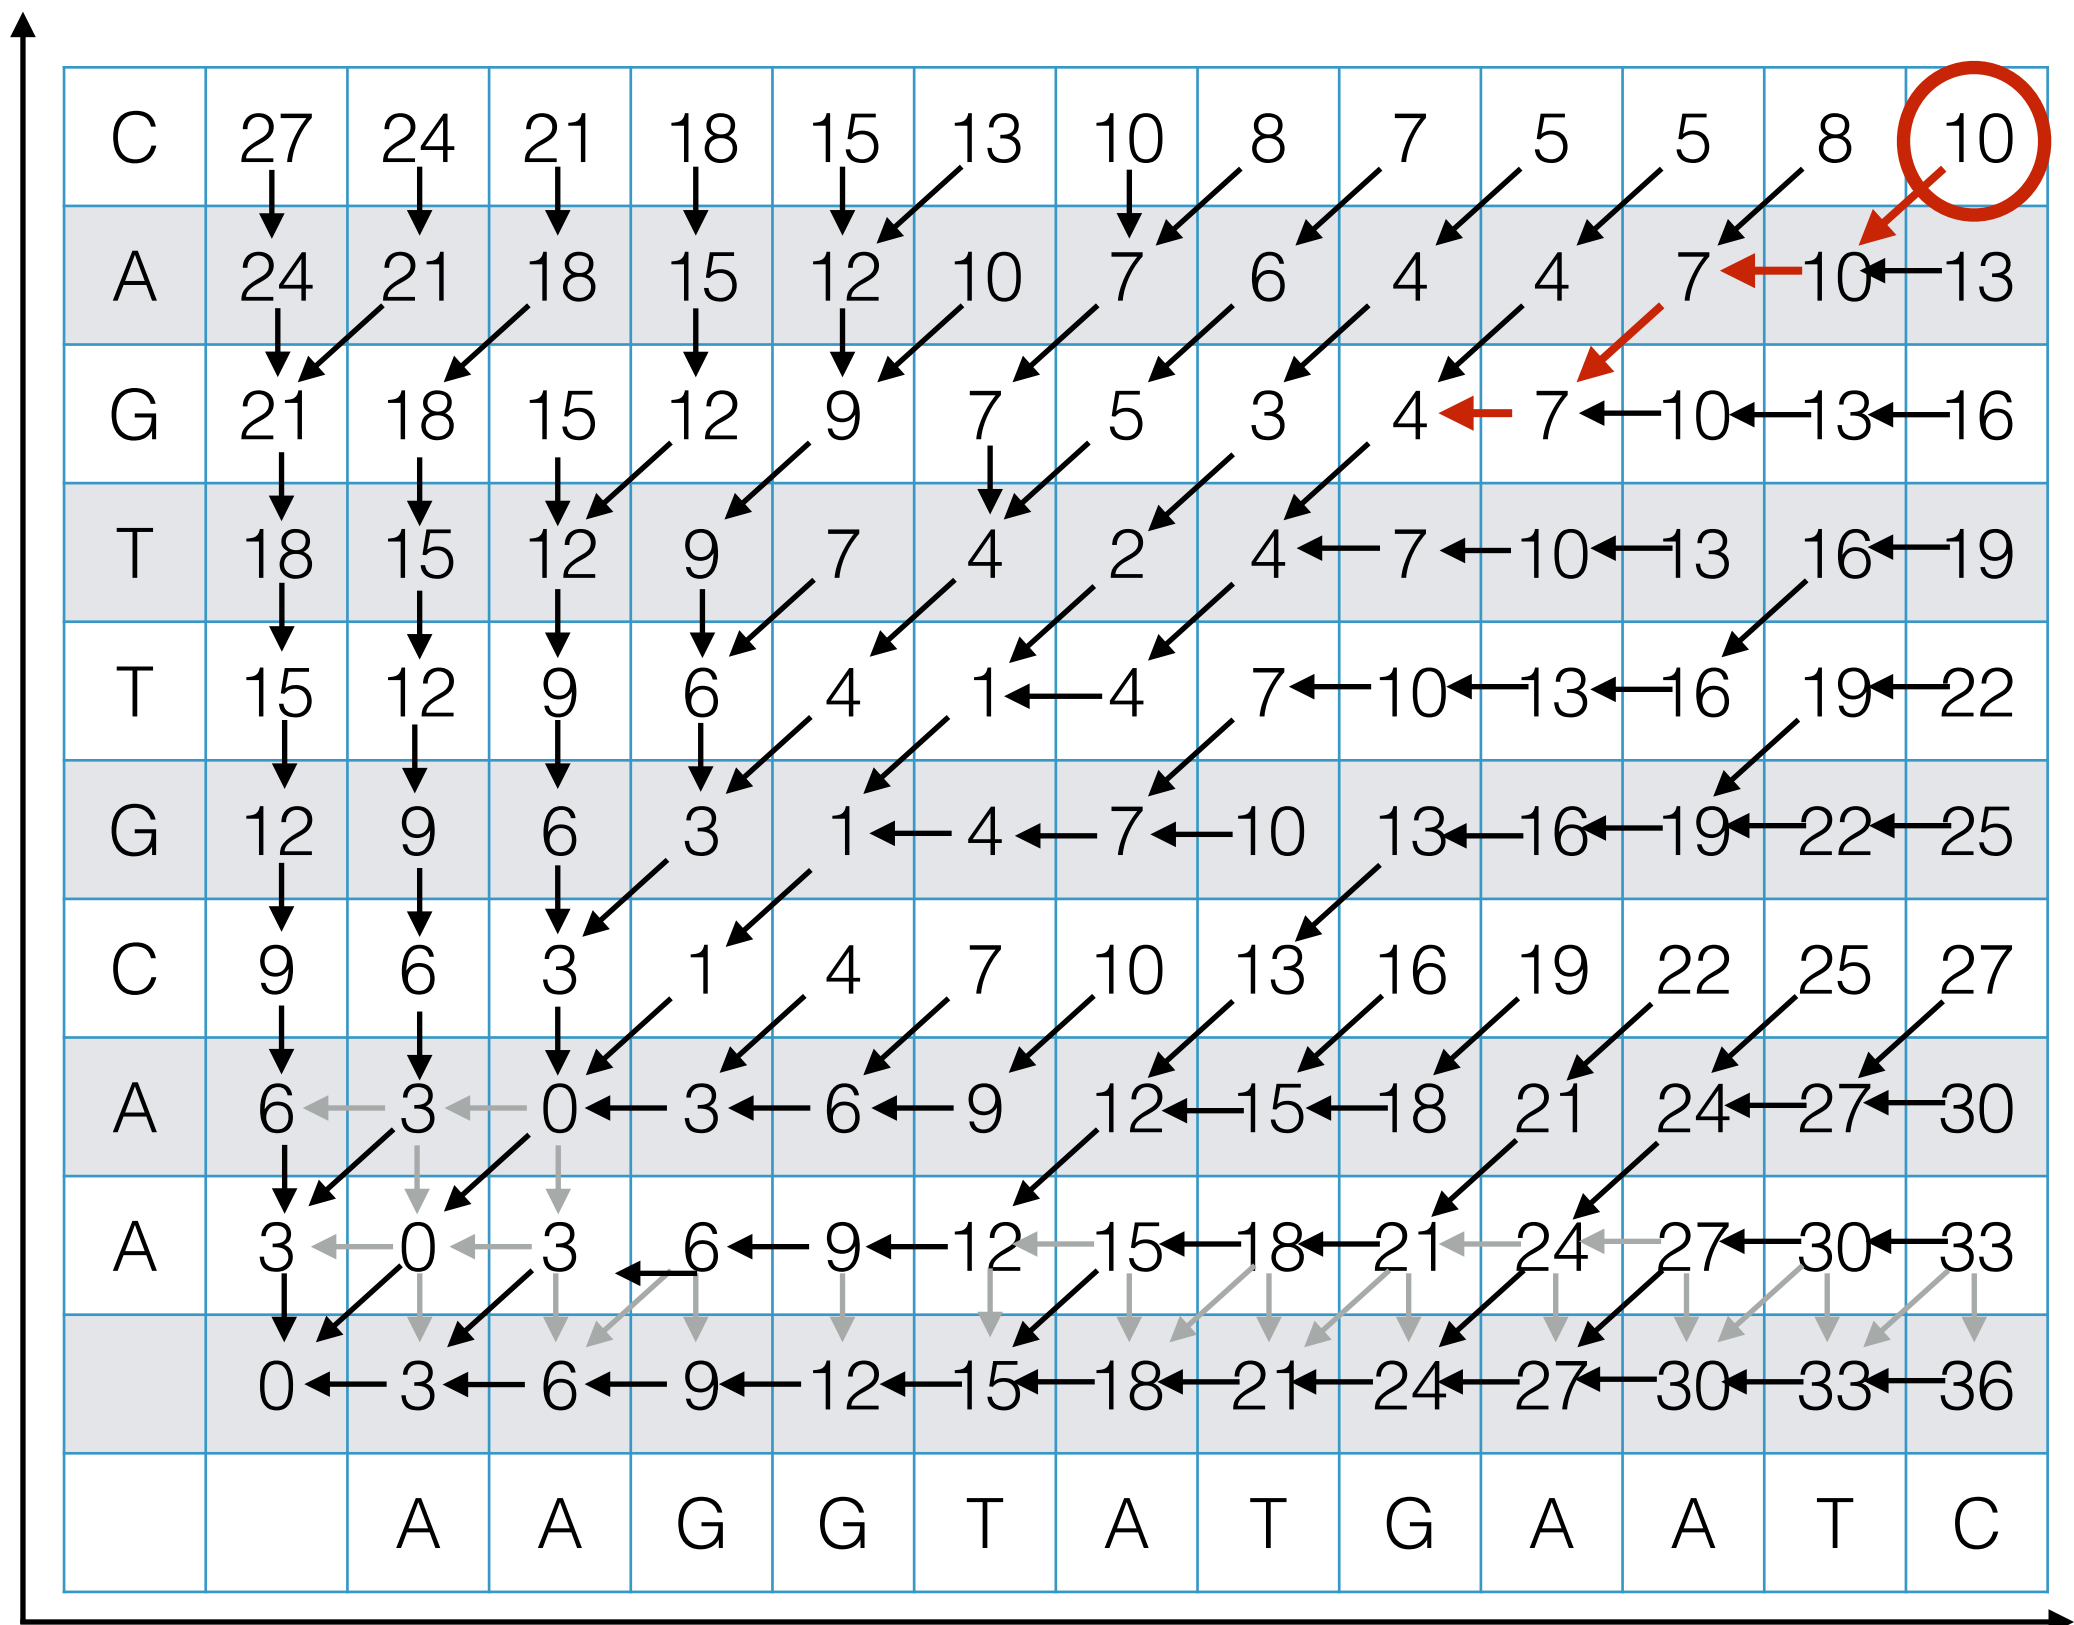
Example



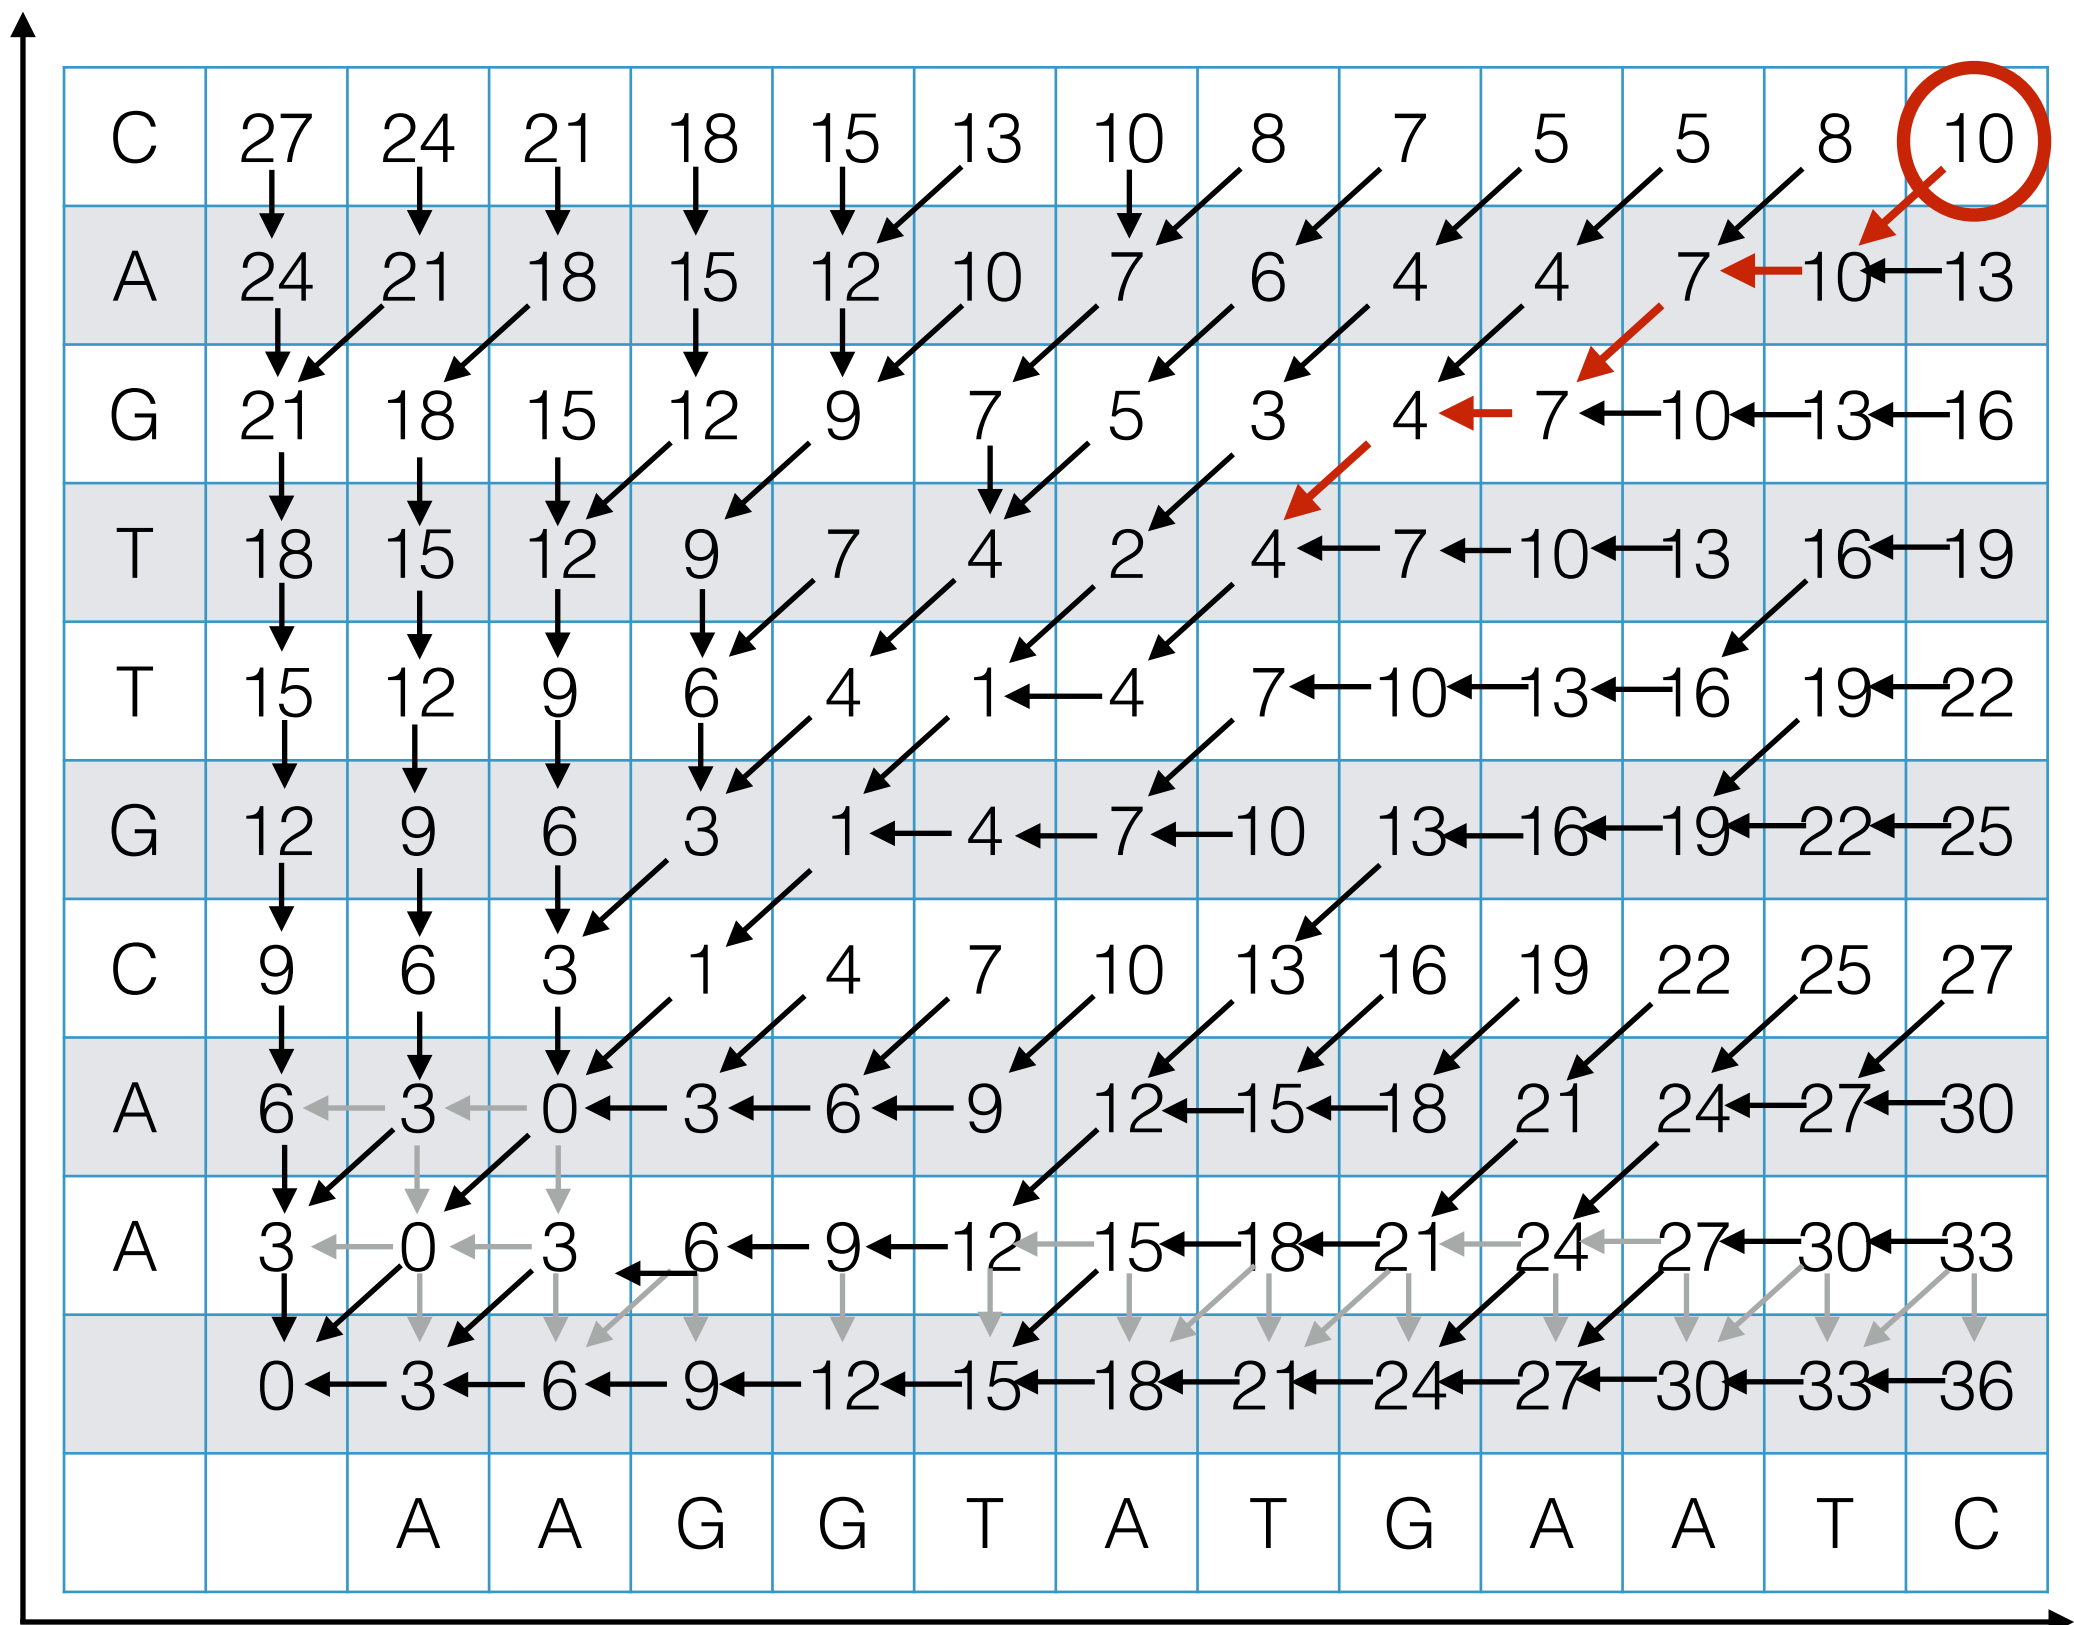
Example



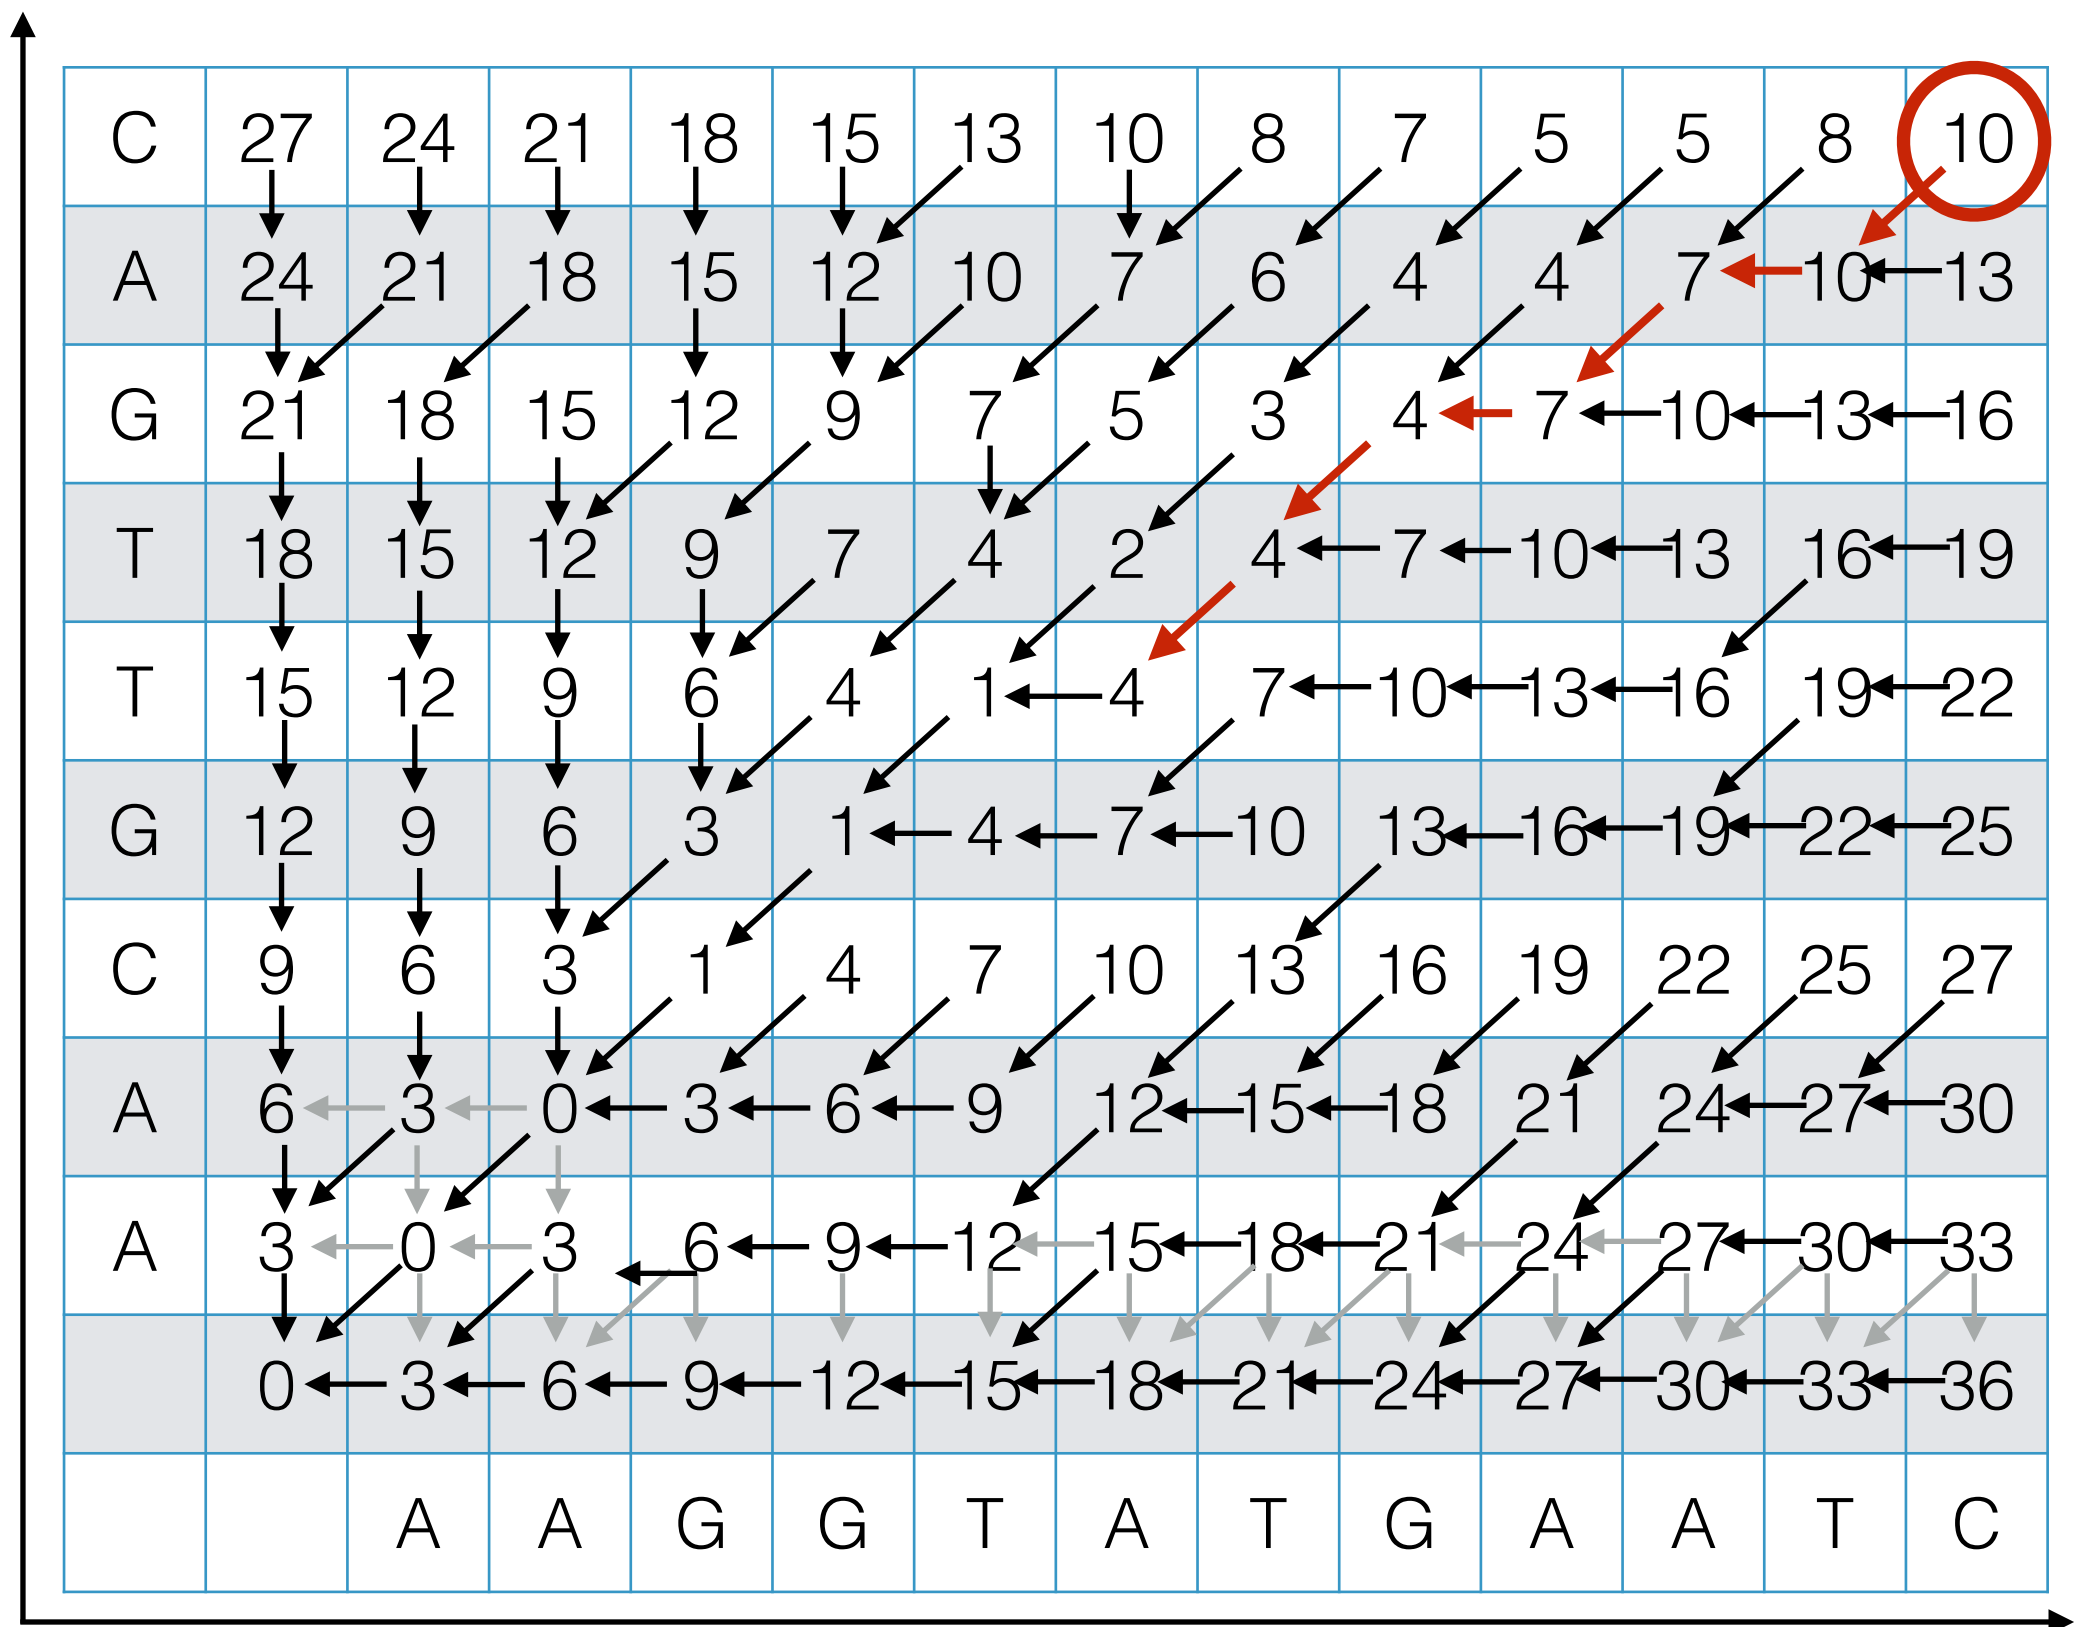
Example



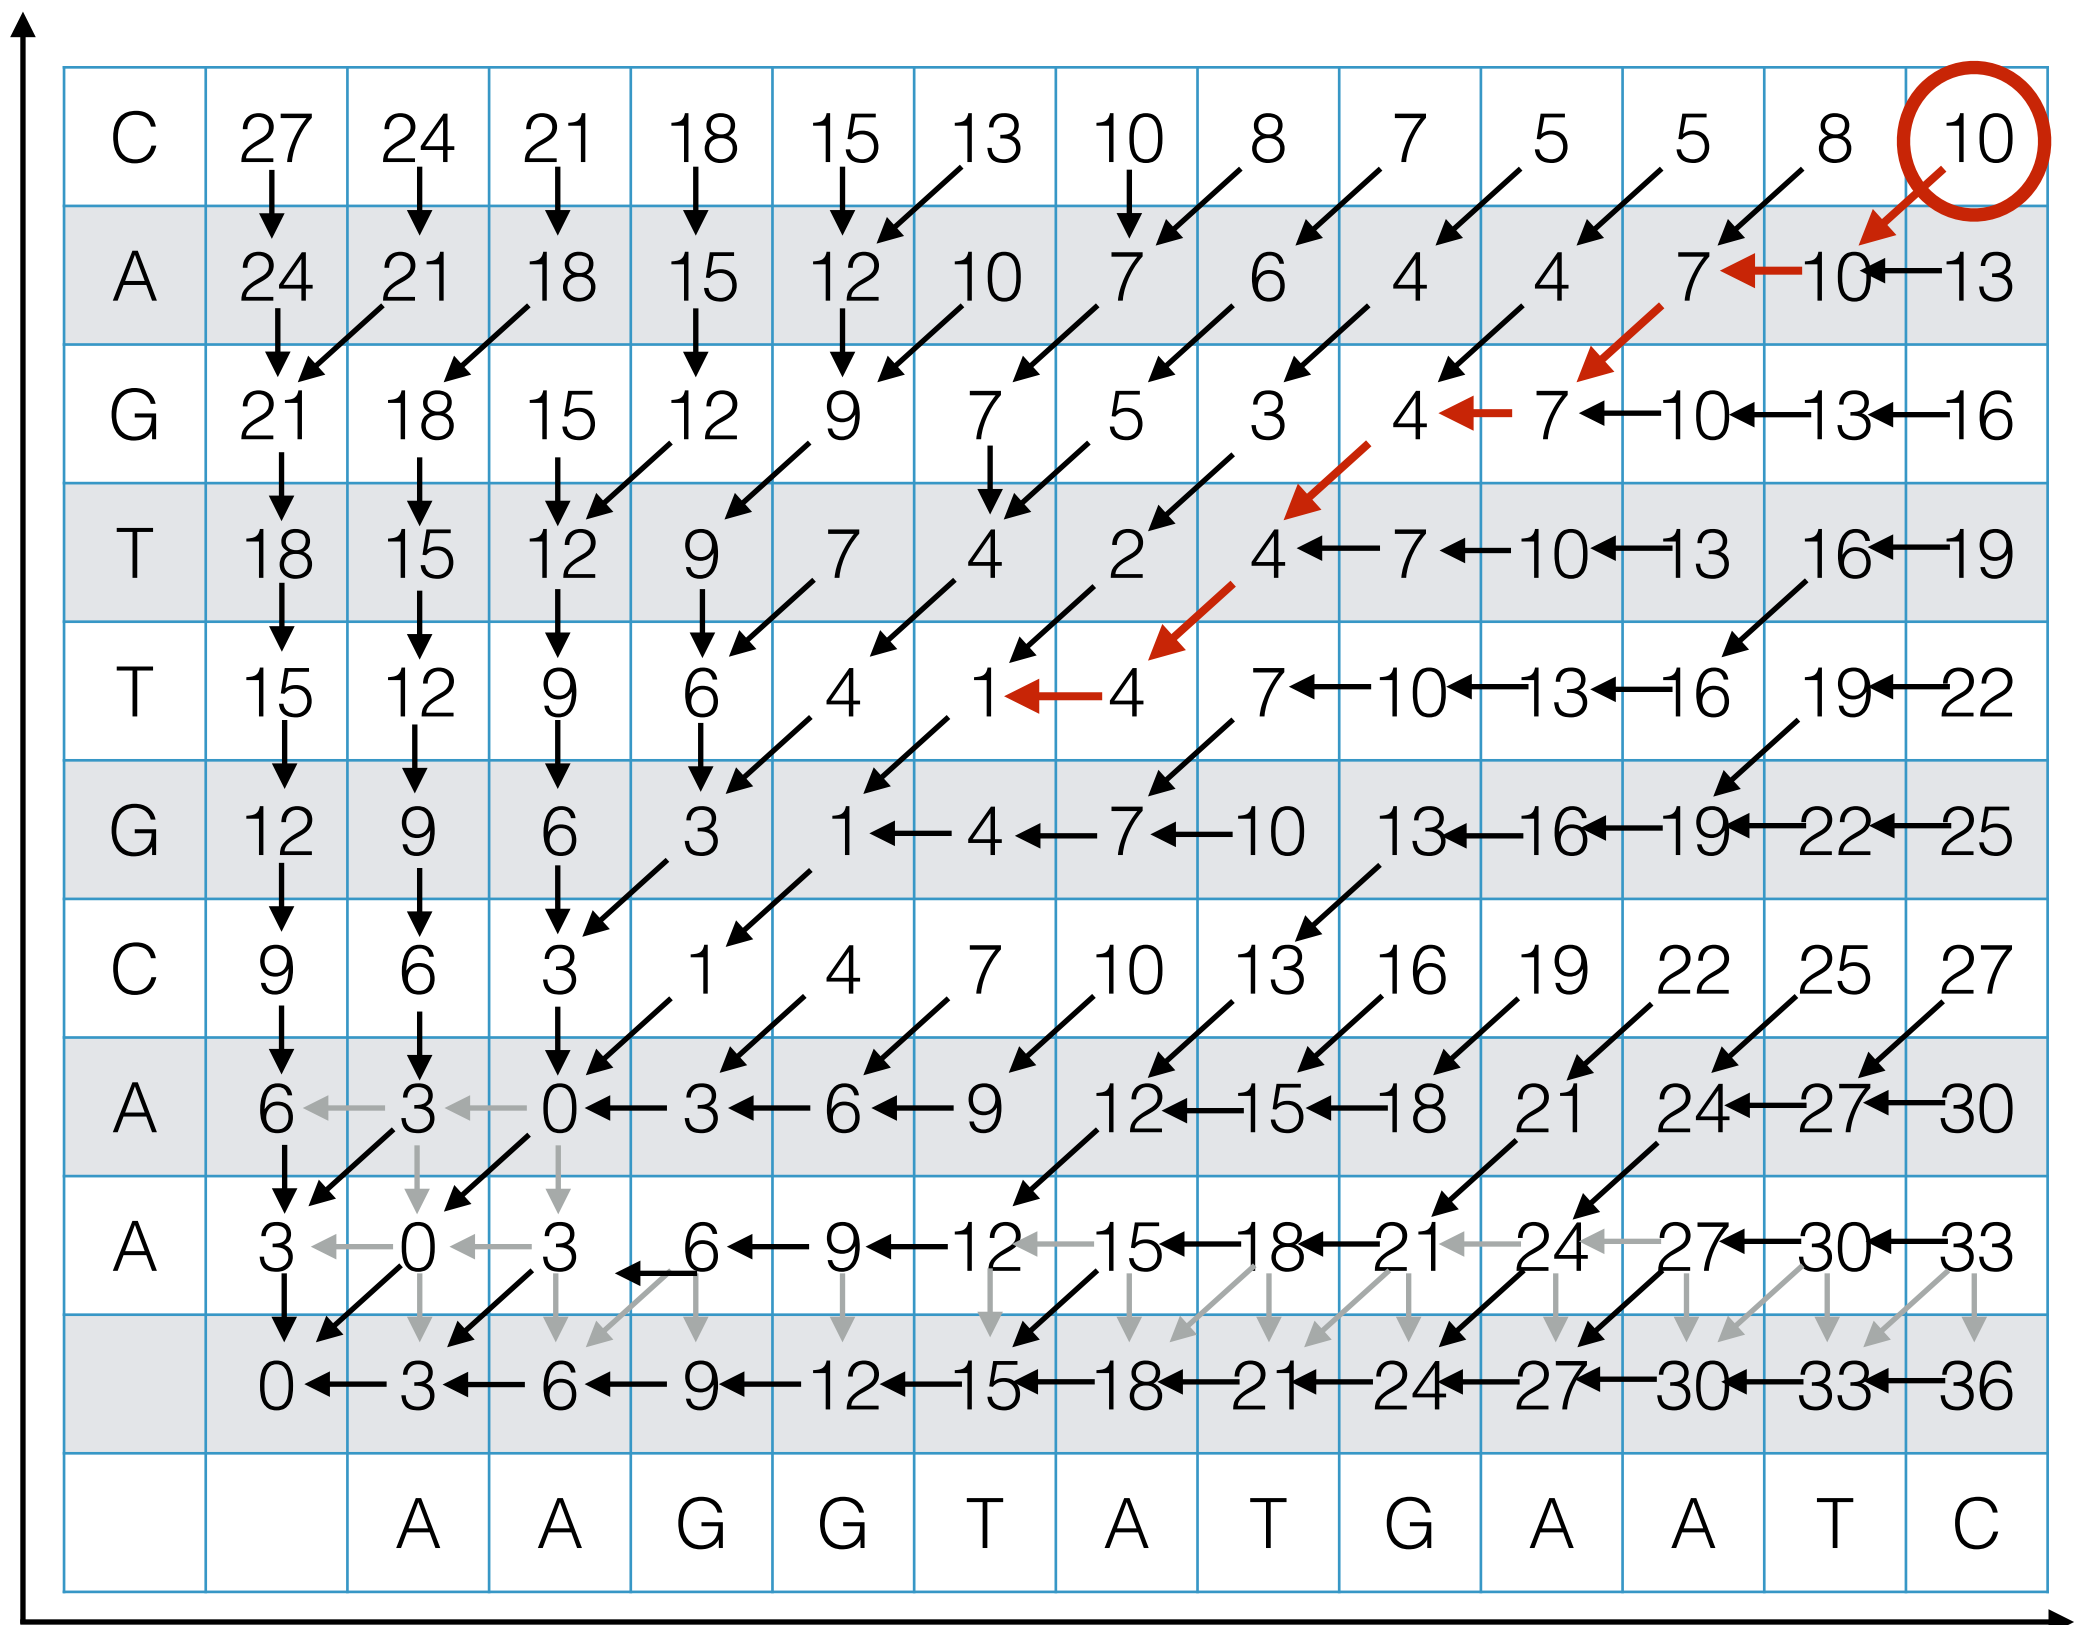
Example



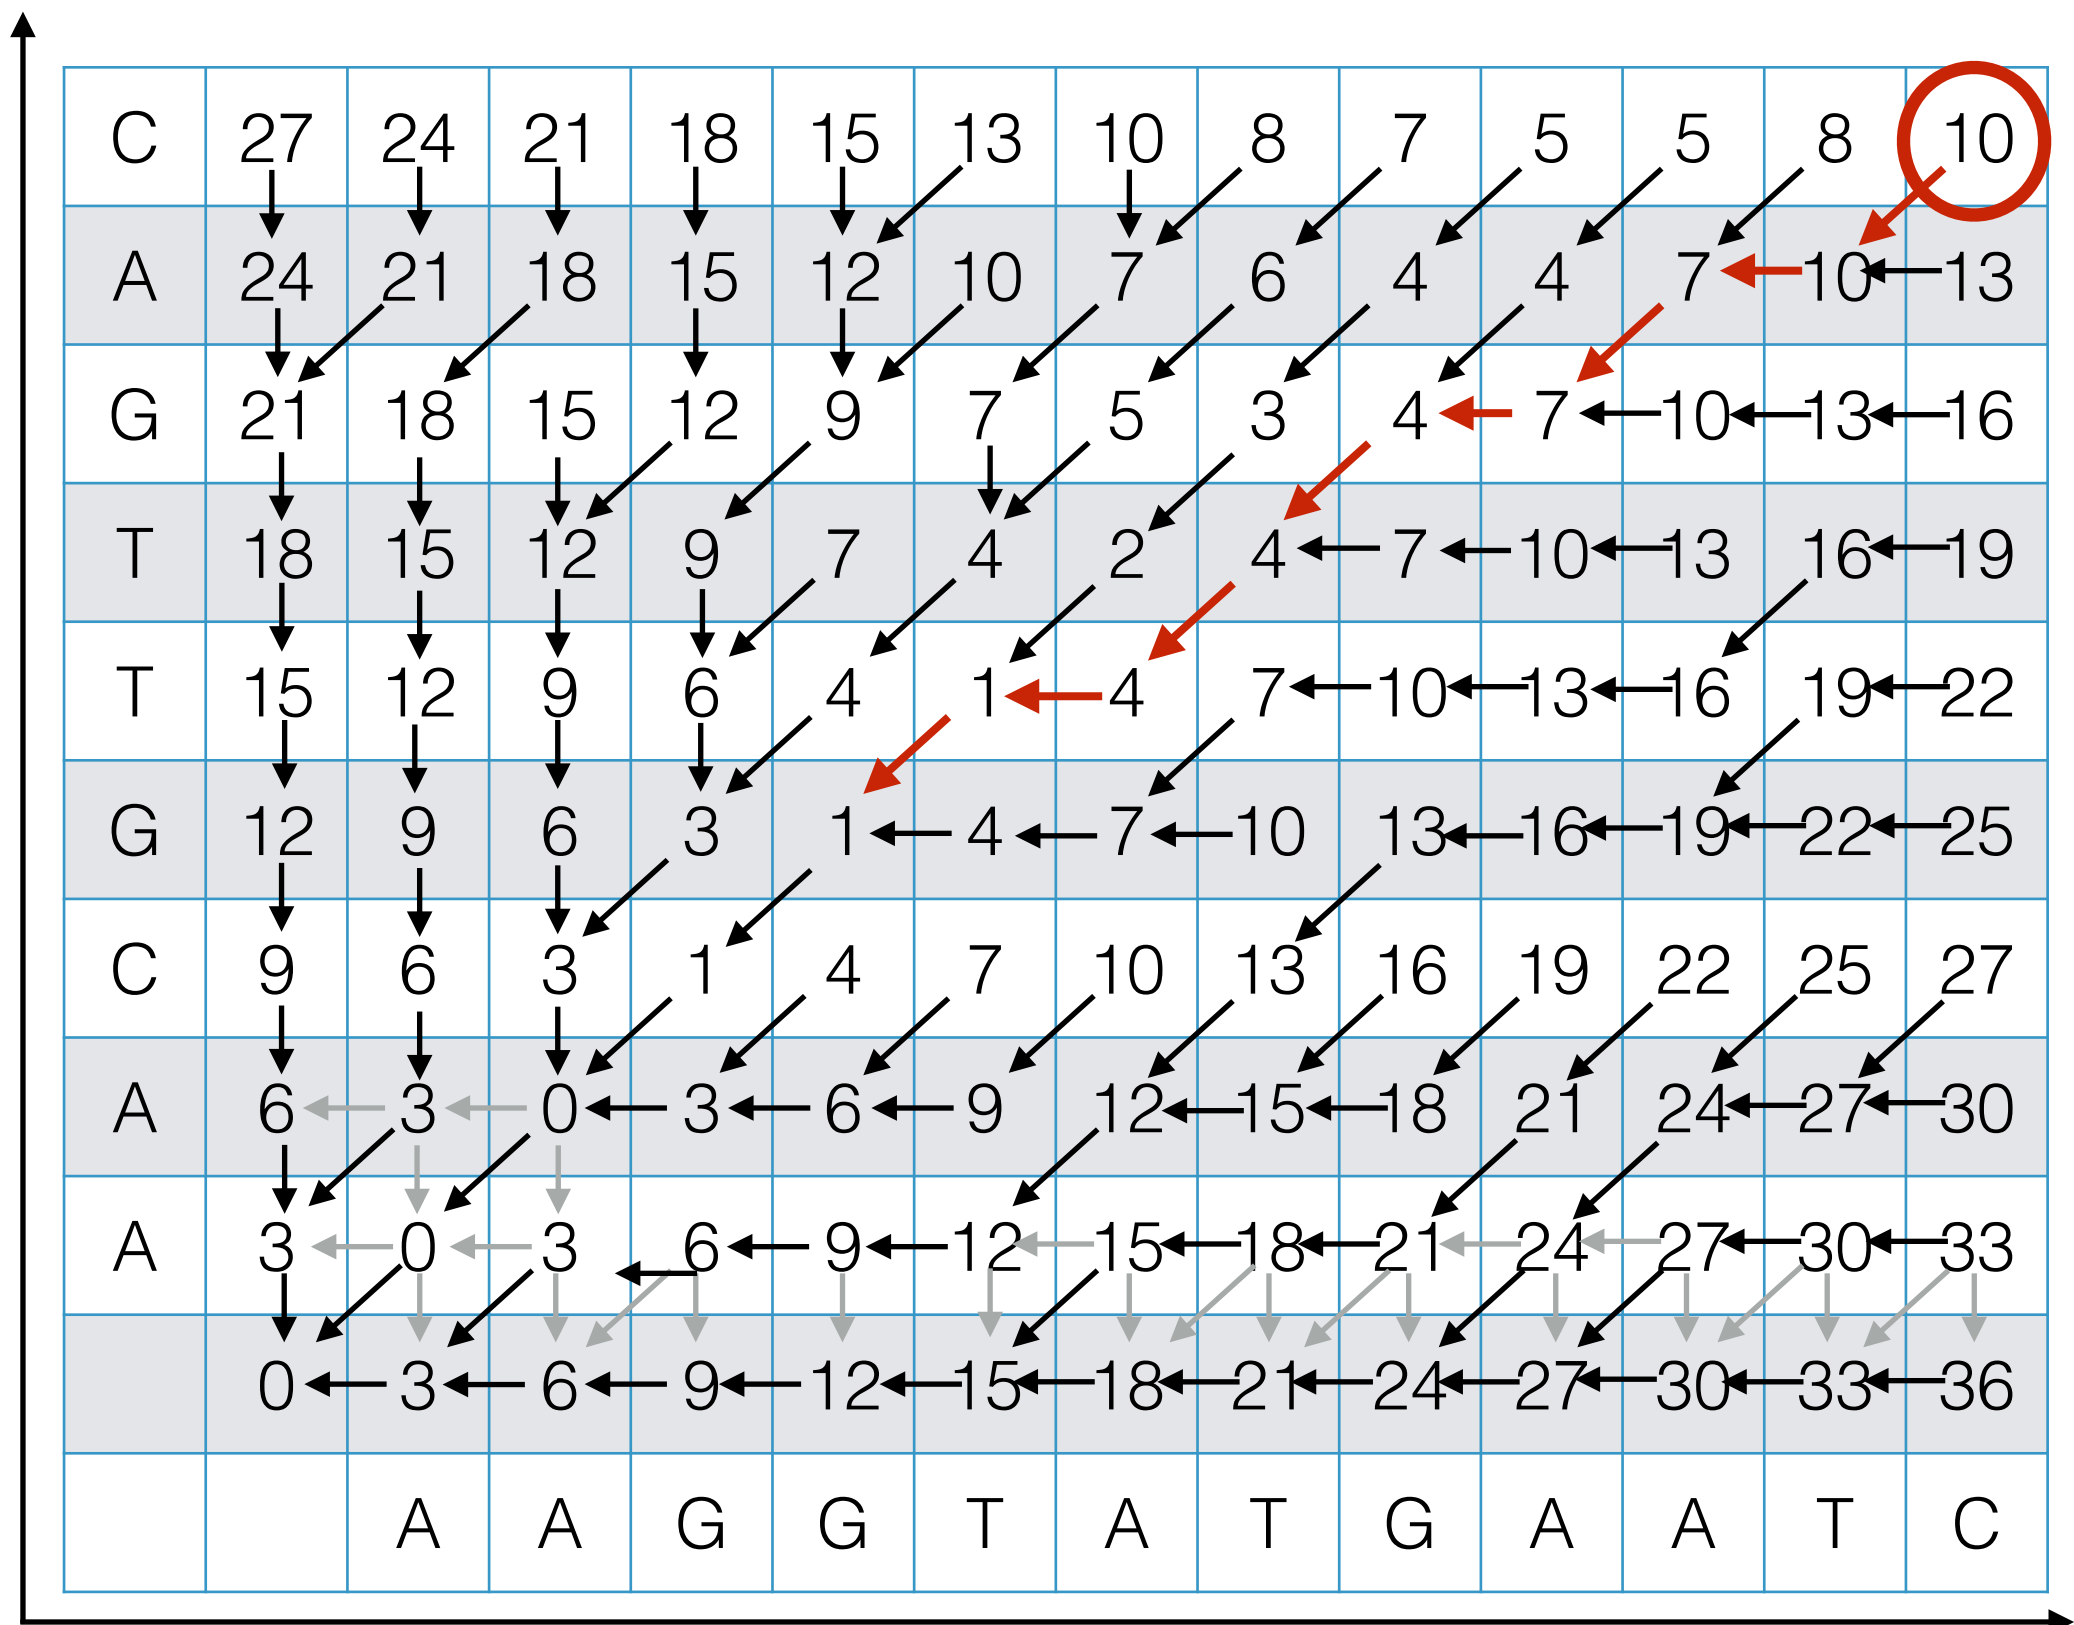
Example



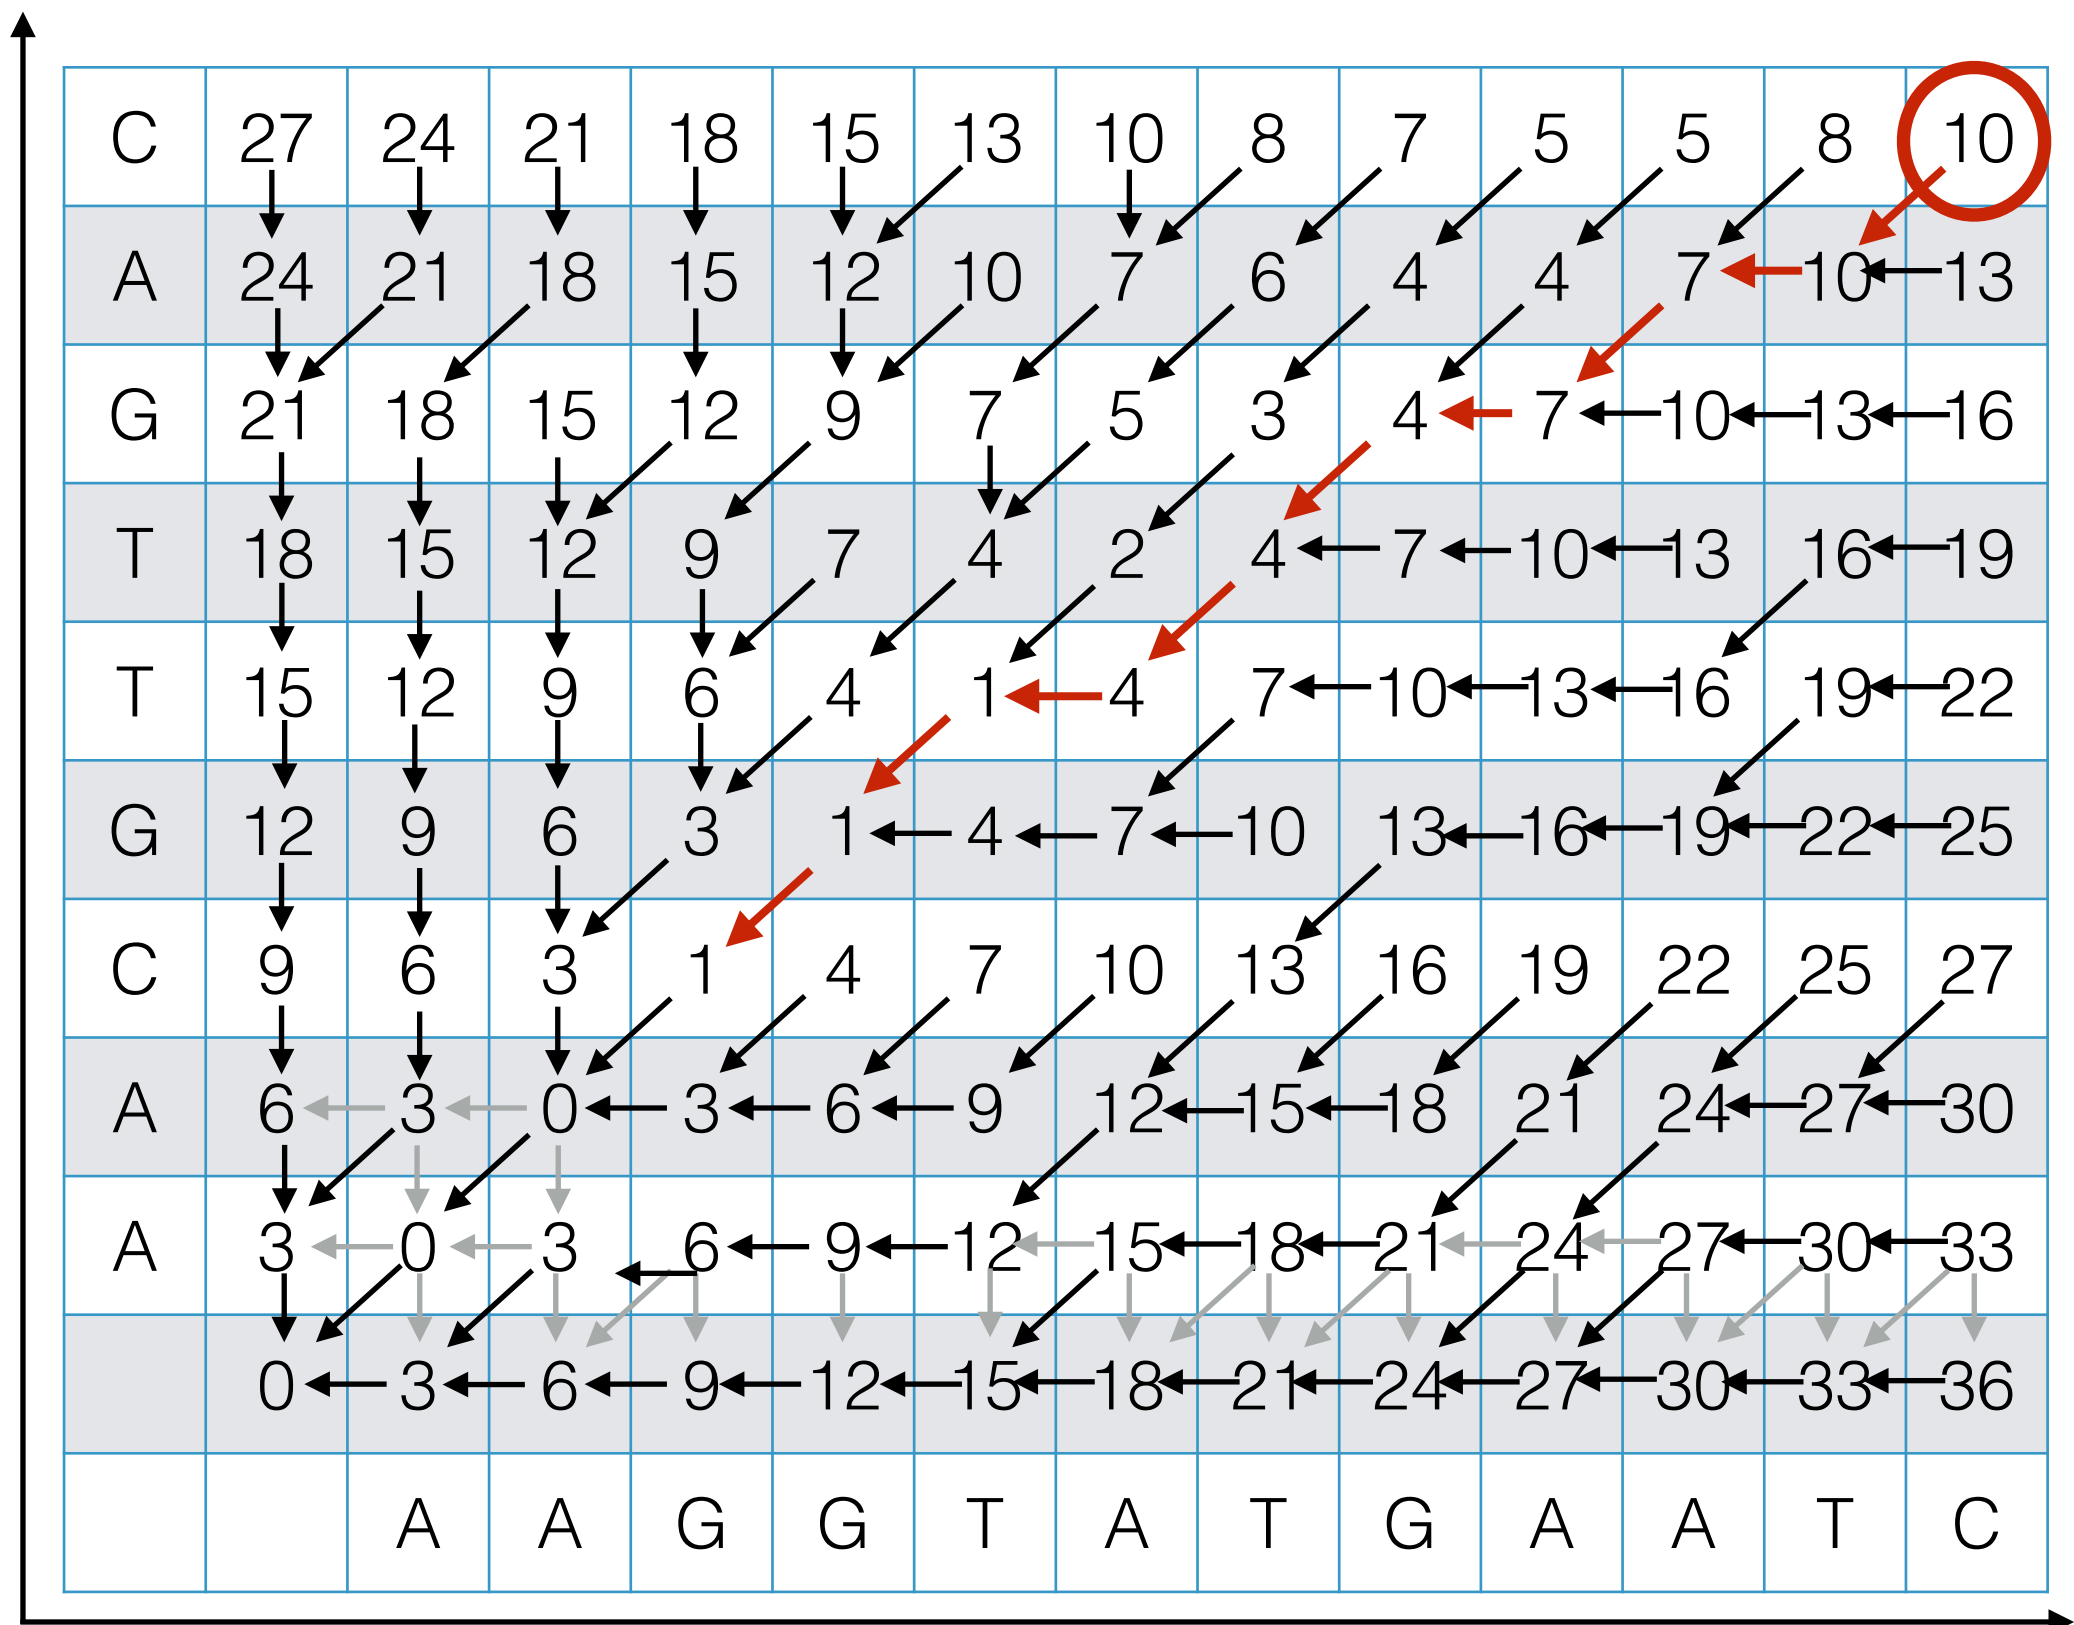
Example



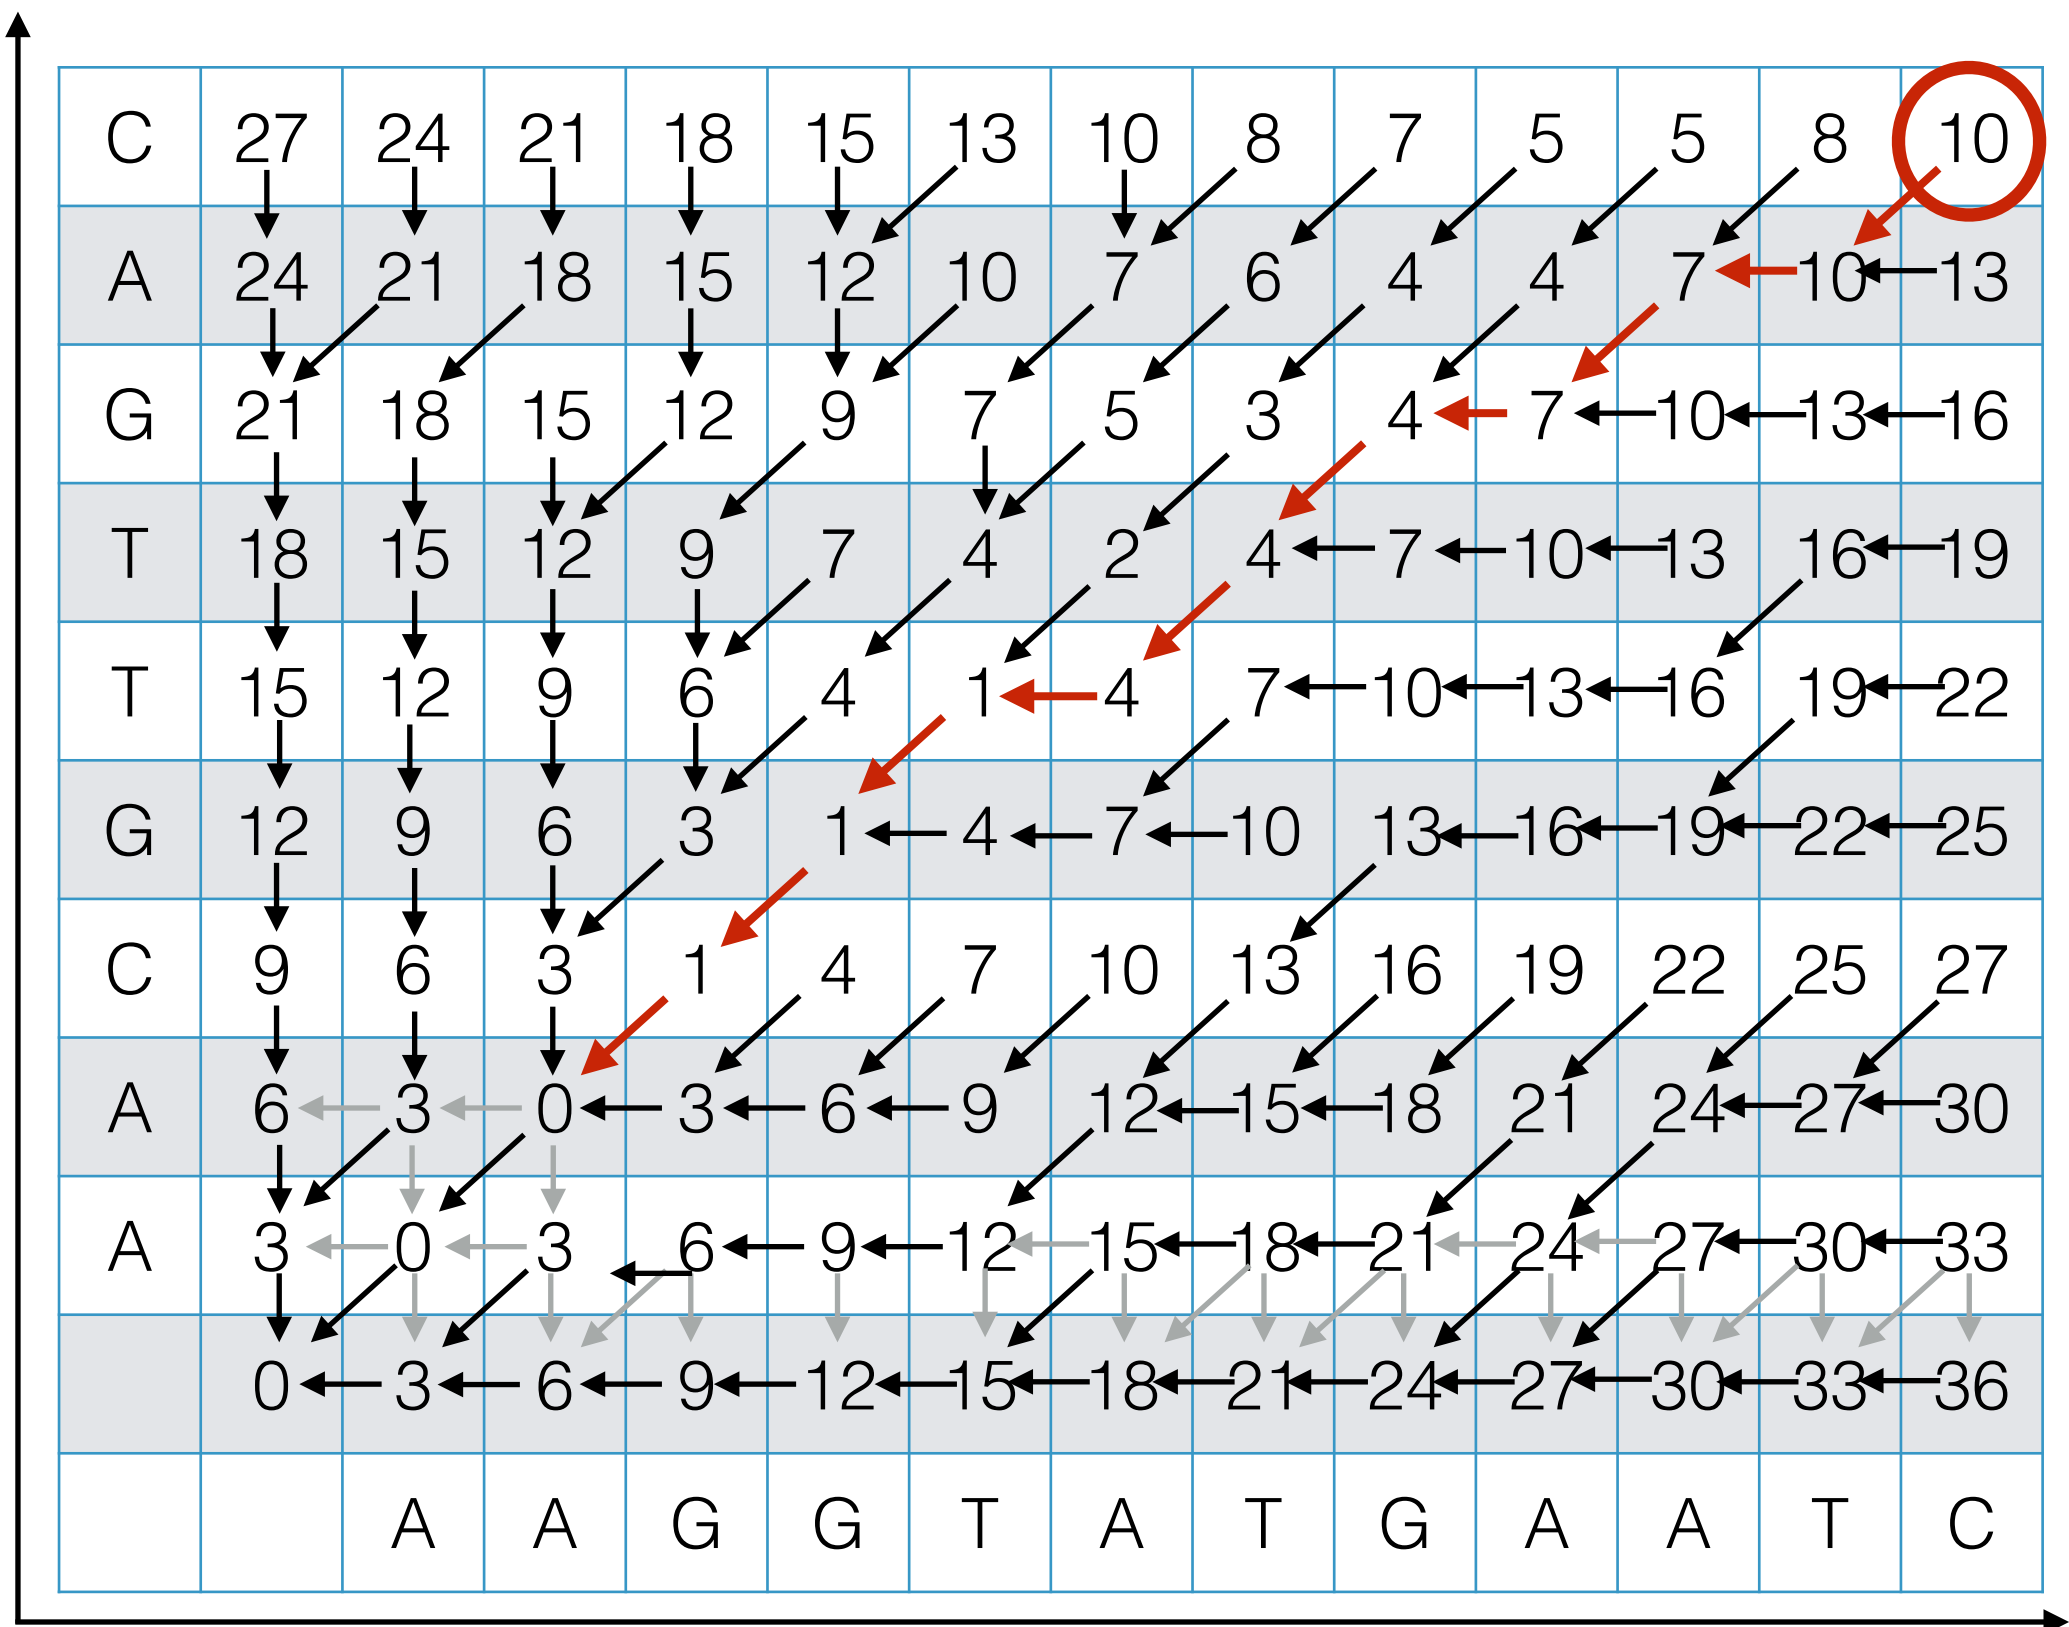
Example



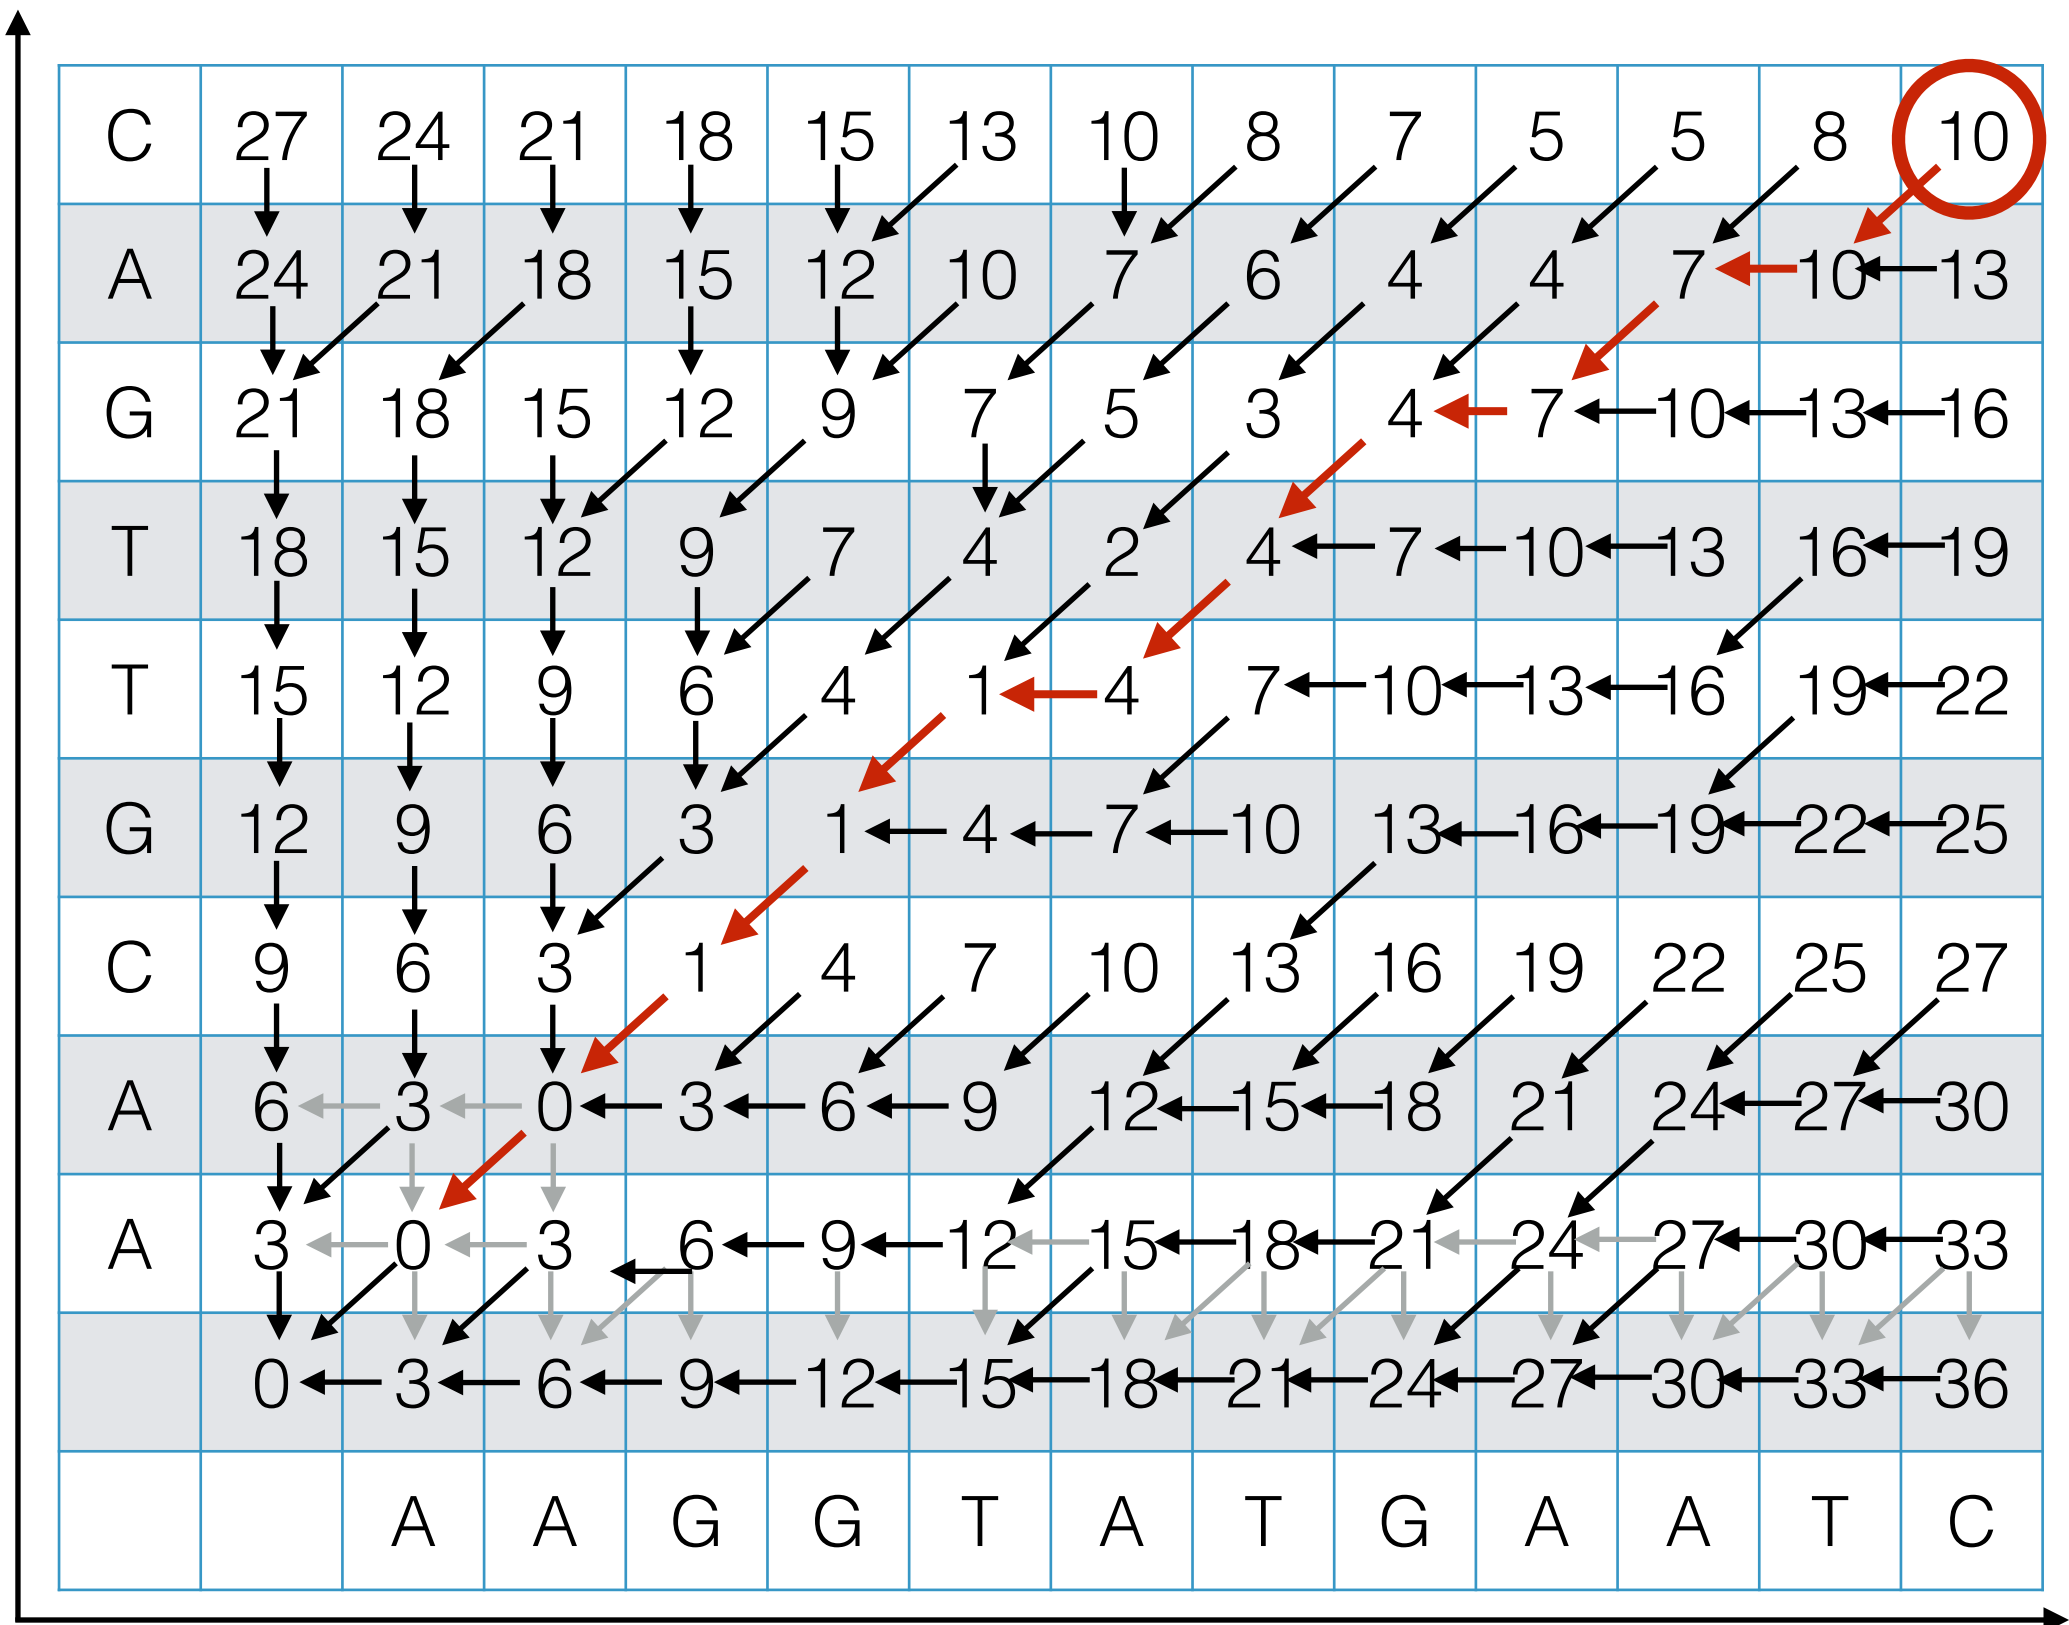
Example



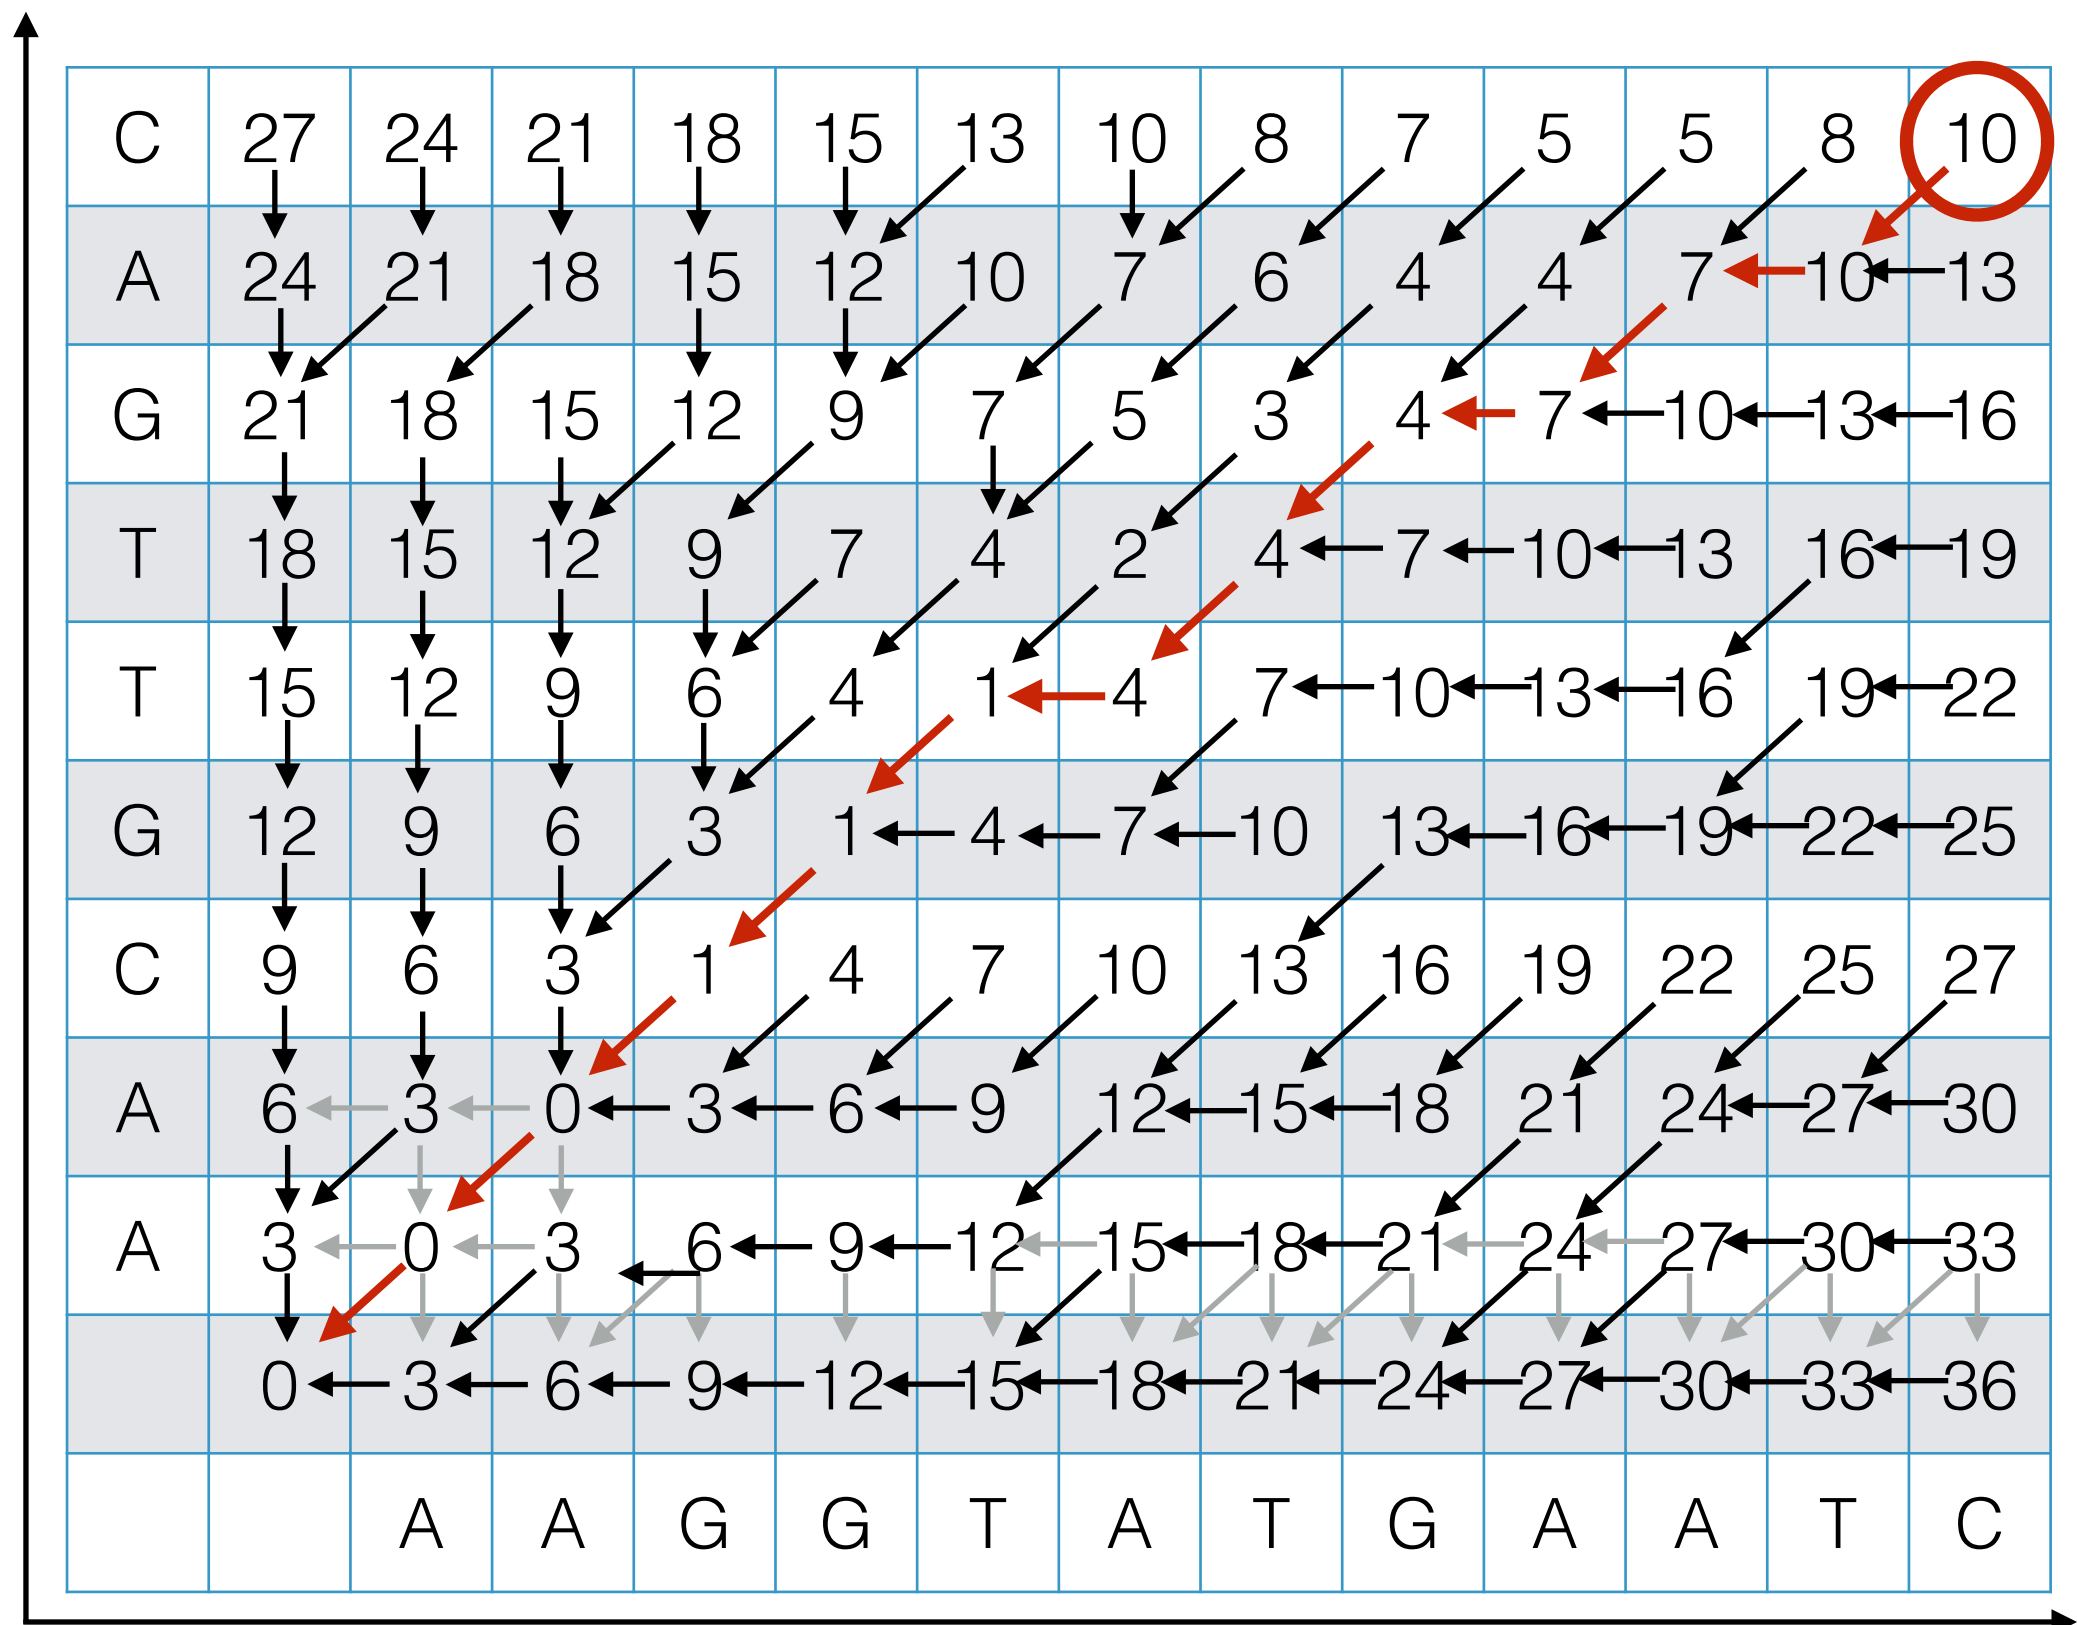
Example



Example



Example



Recap: Dynamic Programming

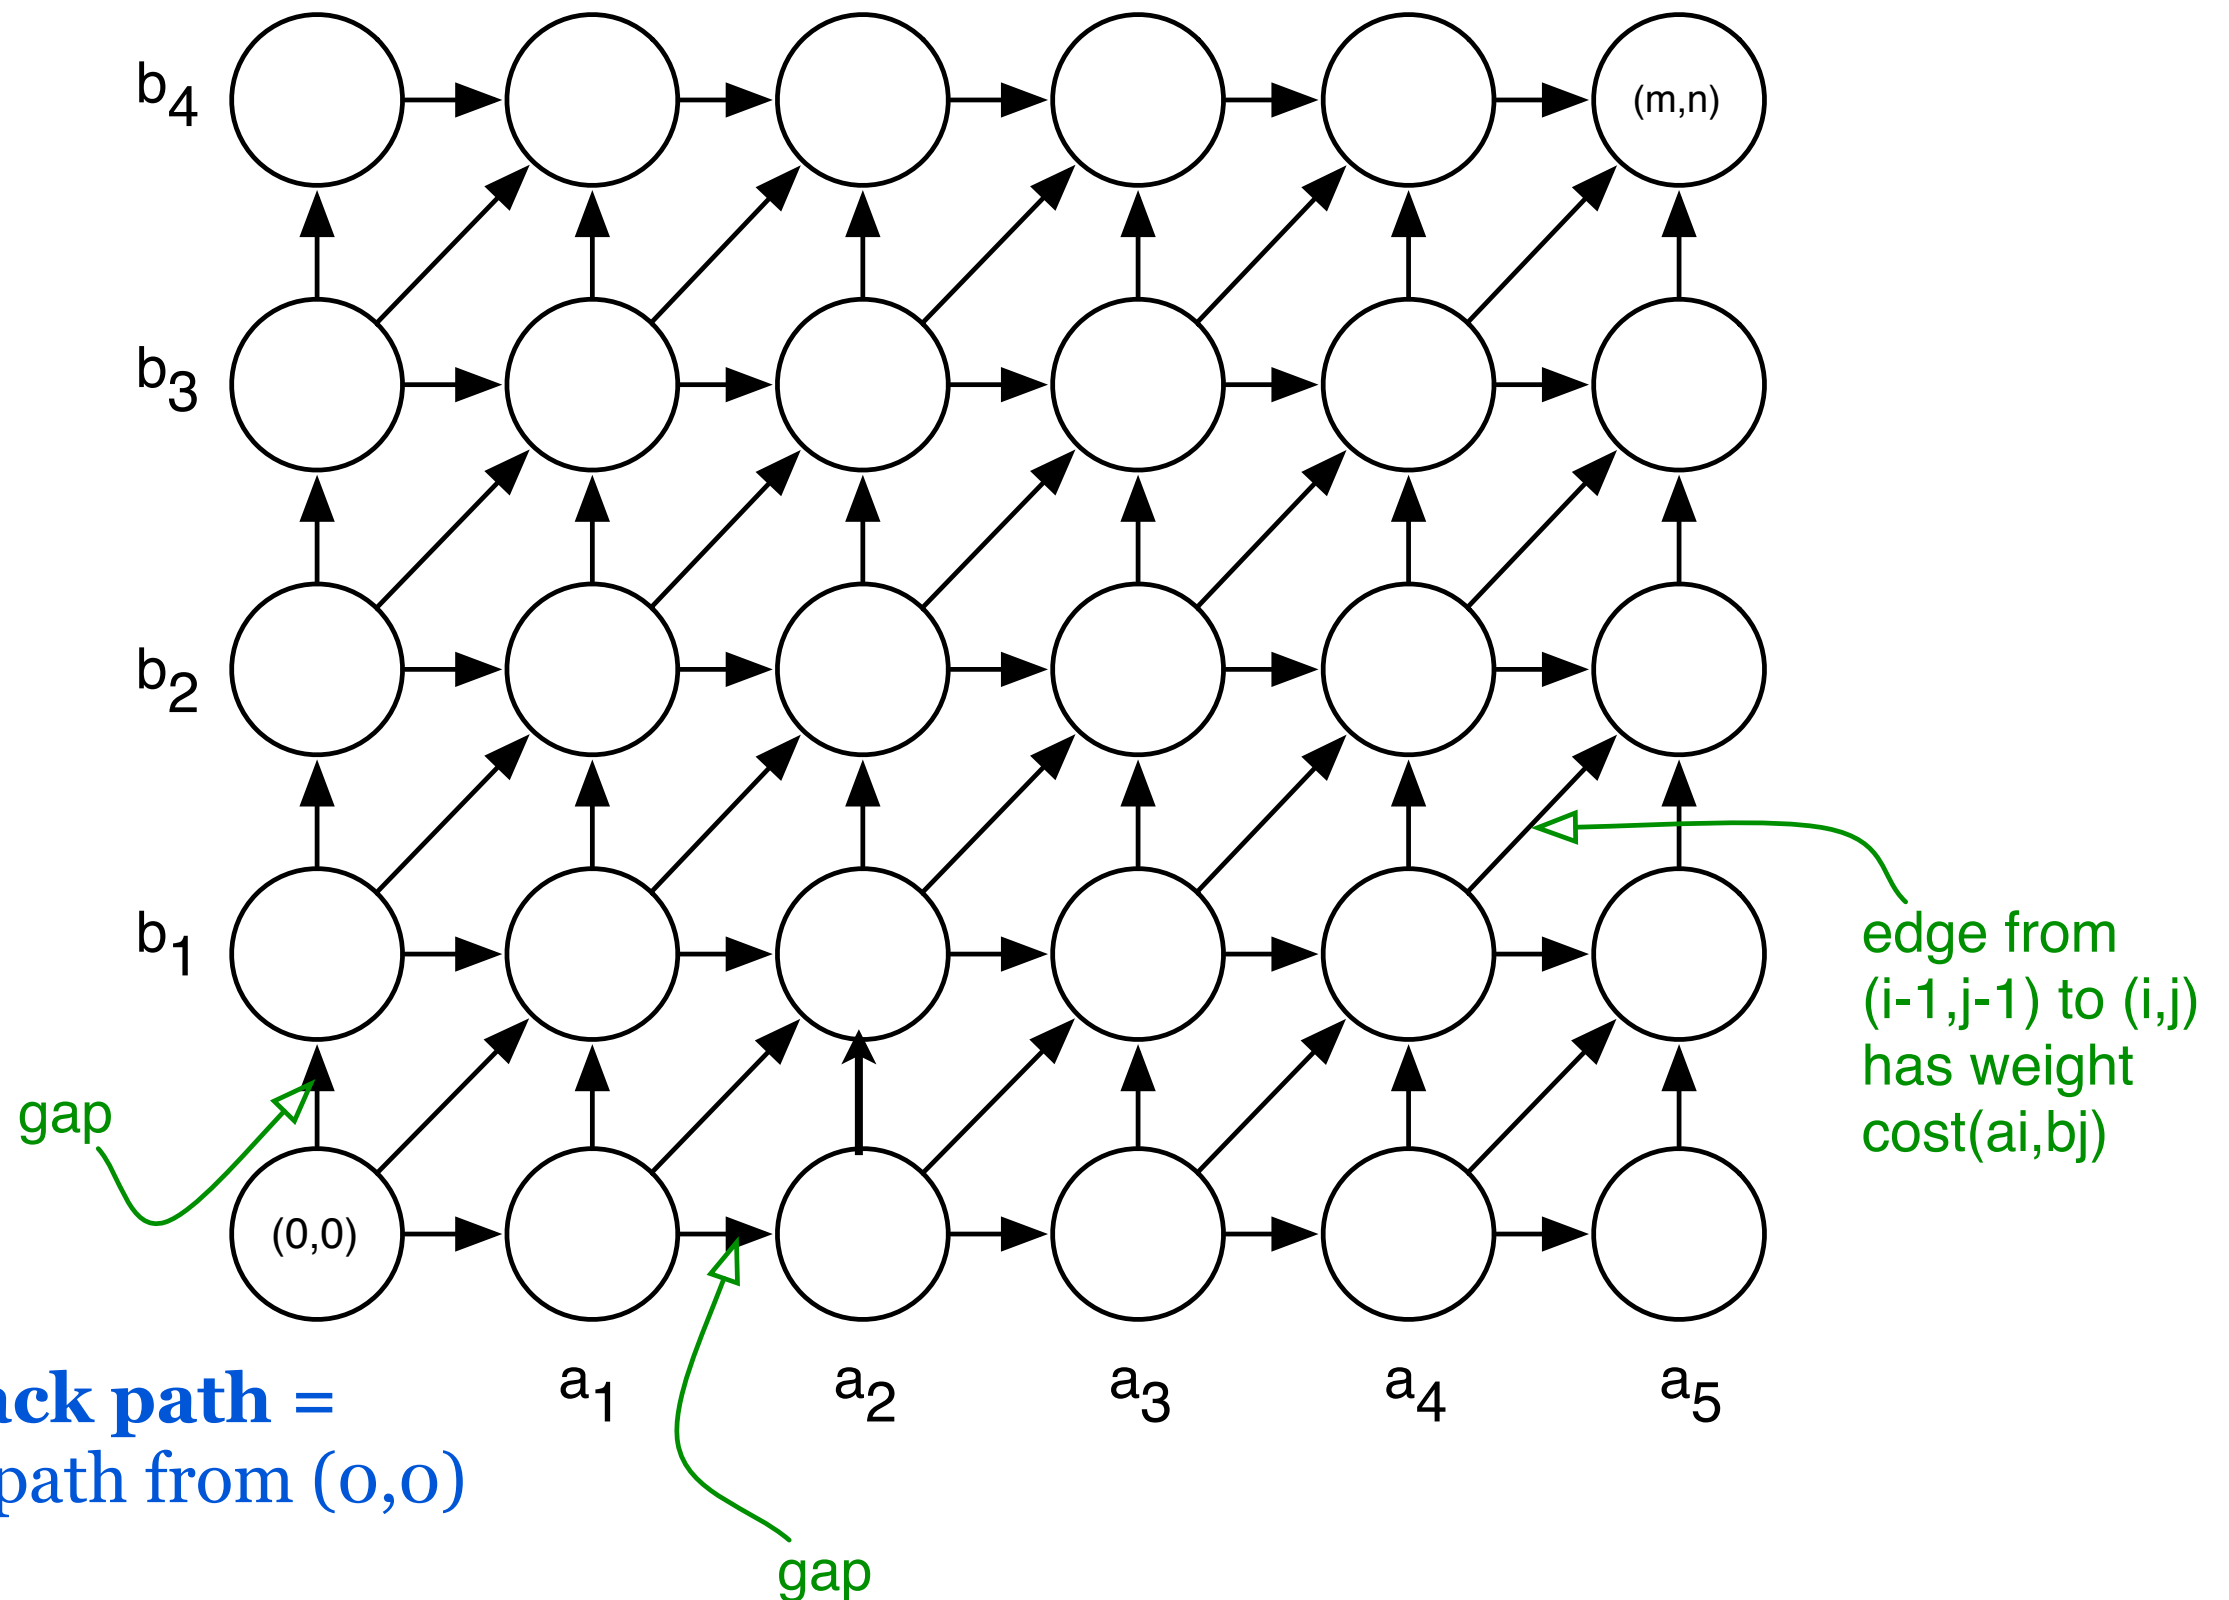
The previous sequence alignment / edit distance algorithm is an example of dynamic programming.

Main idea of dynamic programming: solve the subproblems in an order so that when you need an answer, it's ready.

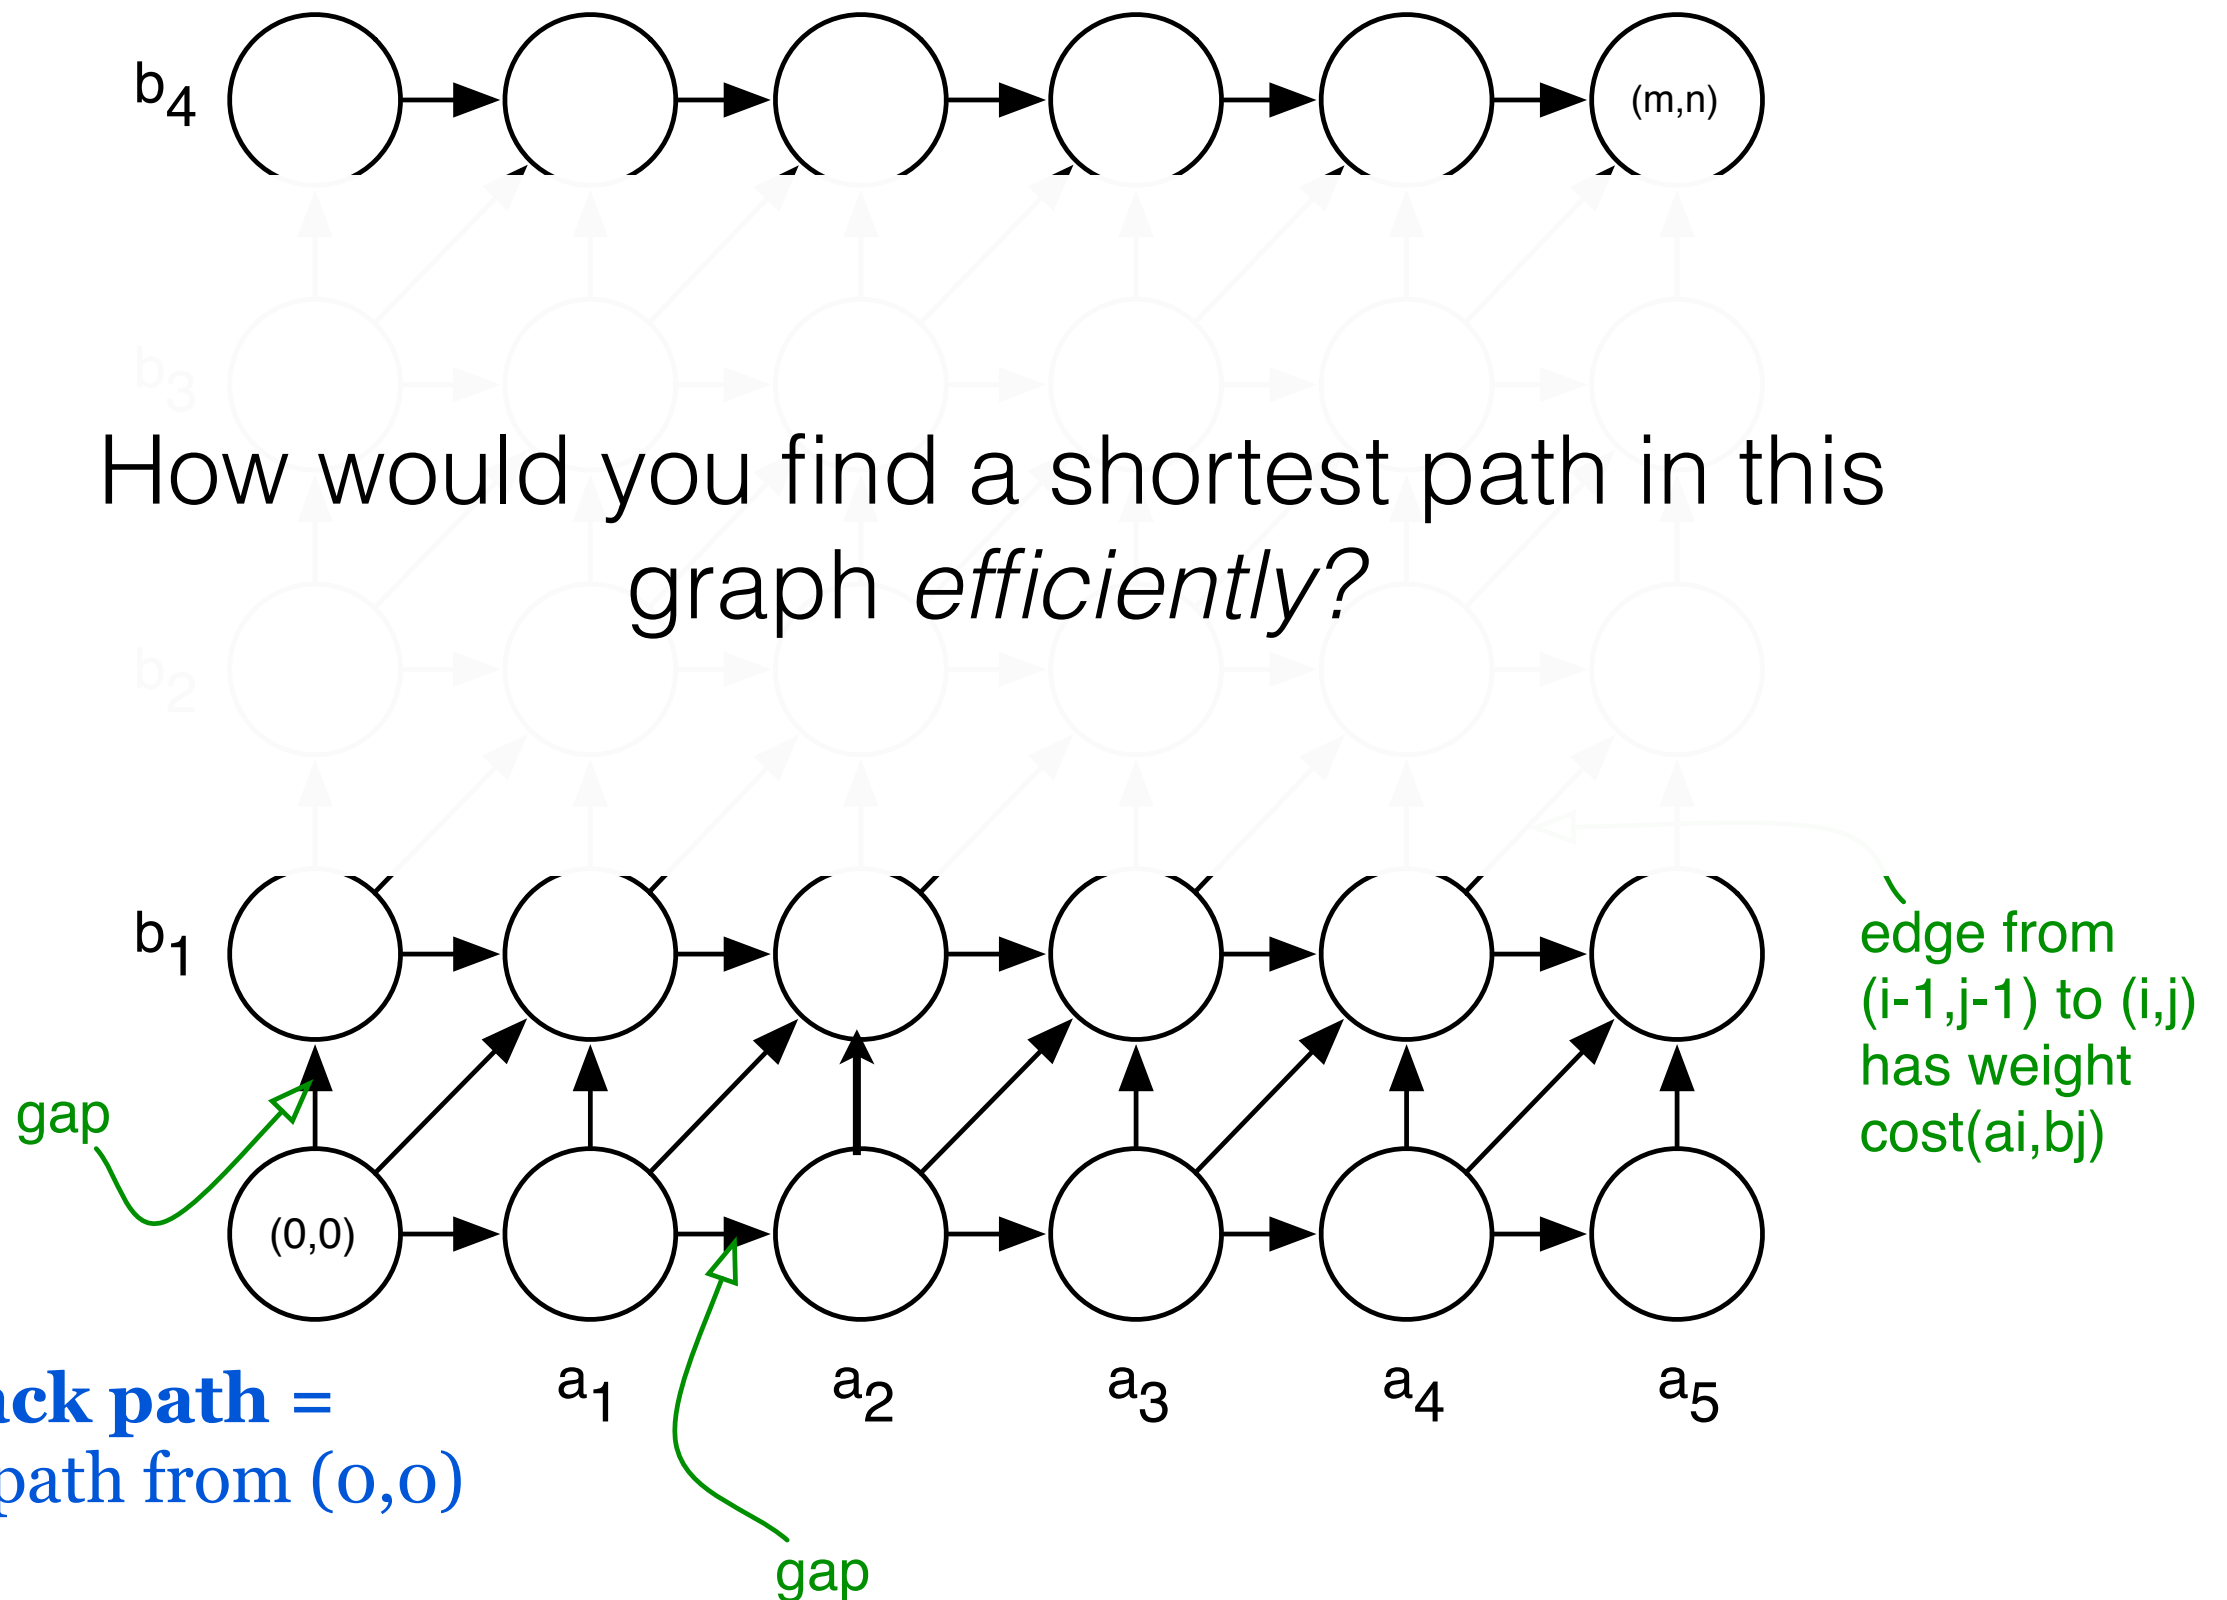
Requirements for DP to apply:

1. Optimal value of the original problem can be computed from some similar subproblems.
2. There are only a polynomial # of subproblems
3. There is a “natural” ordering of subproblems, so that you can solve a subproblem by only looking at **smaller** subproblems.

Another View: Recasting as a Graph



Another View: Recasting as a Graph



Semi-global Alignment Example

Semi-global (glocal): Gaps at the beginning or end of **x** or **y** are free — one useful case is when one string is significantly shorter than the other

sometimes called “cost-free-ends” or “fitting” alignment



We'll discuss the “fitting” variant for in the next few slides for simplicity, but the same basic idea applies for the “overlap” variant as well.

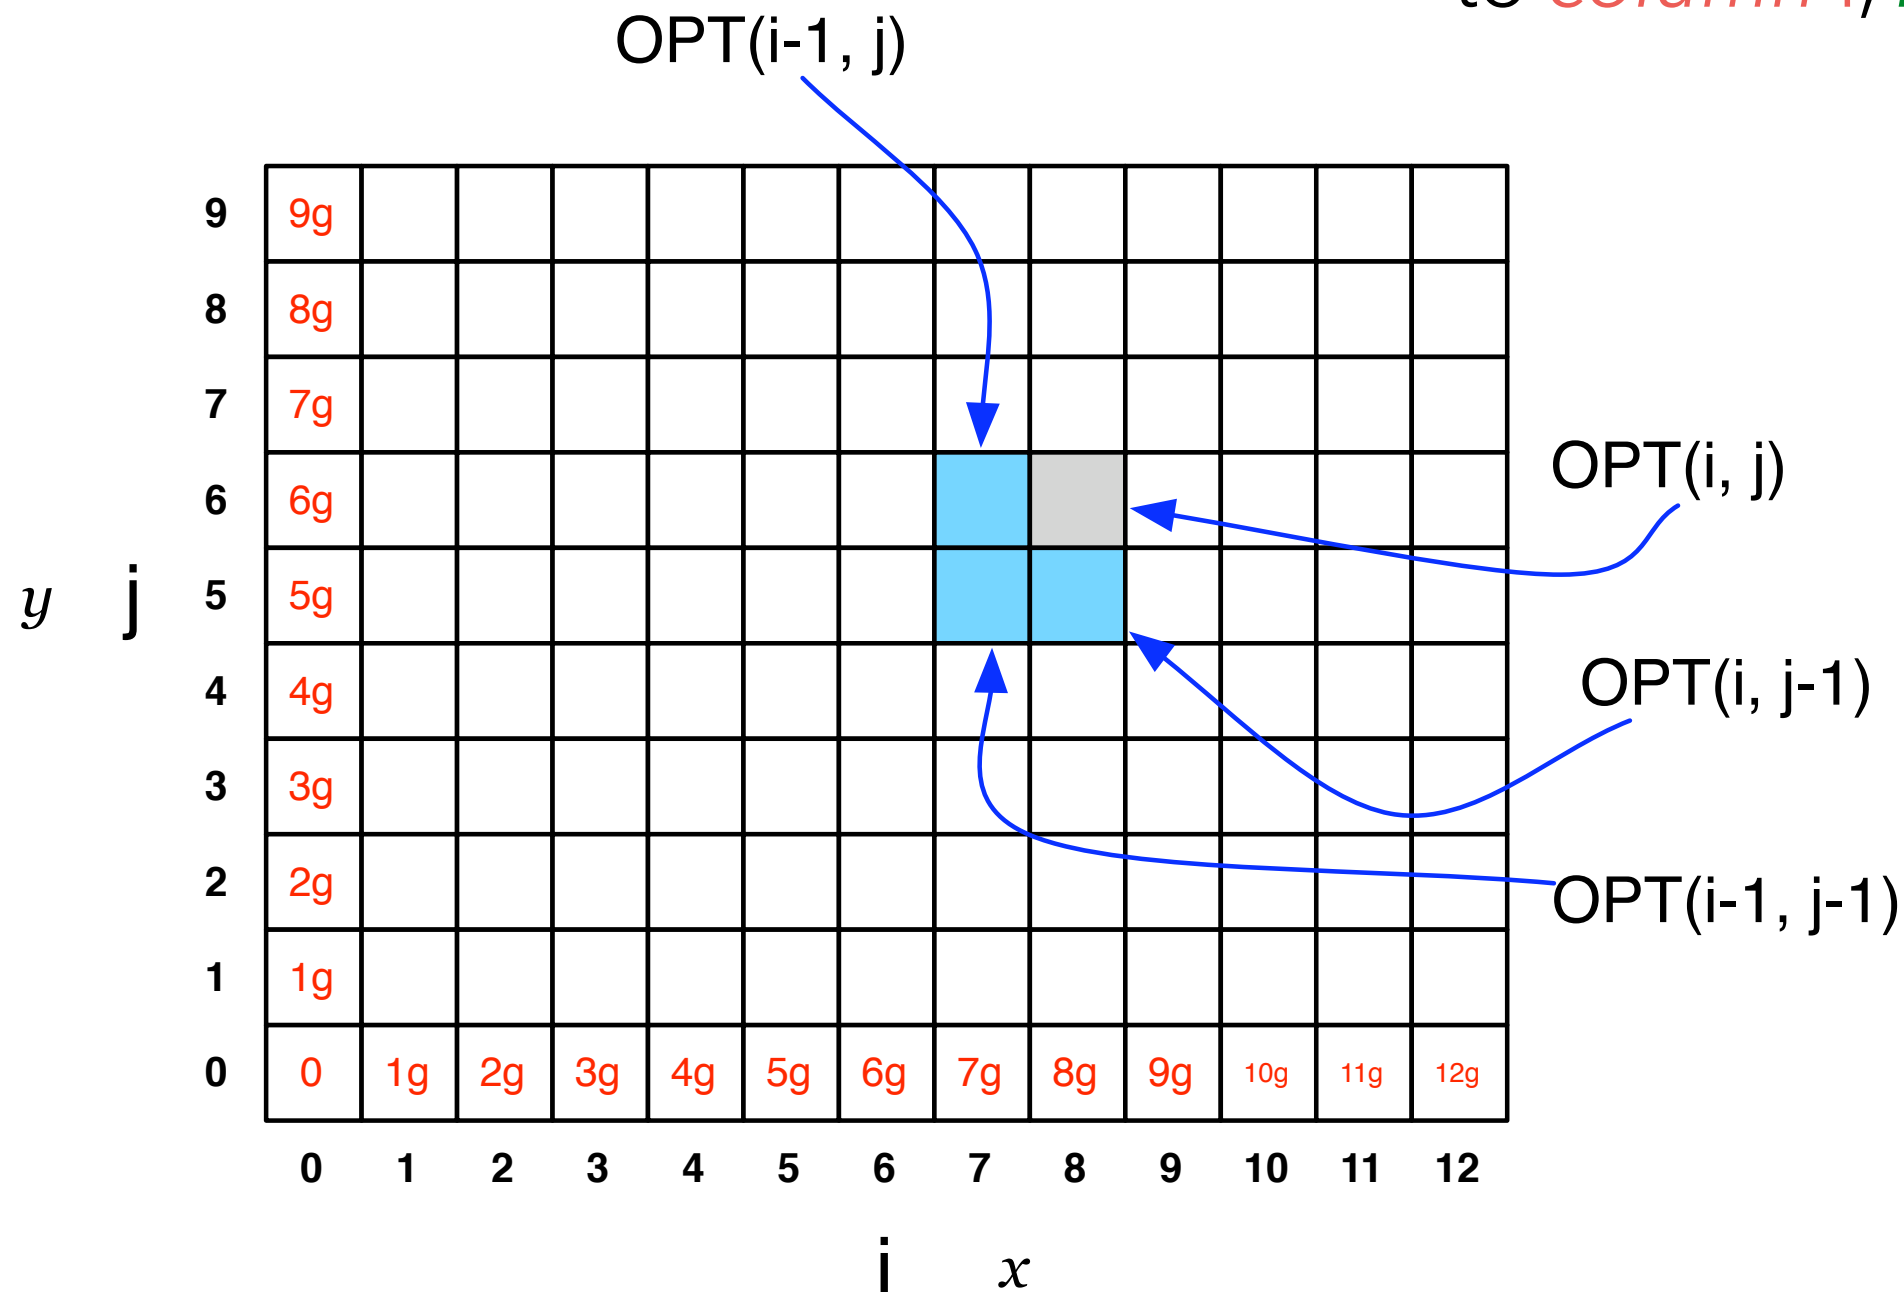
Recall: Global Alignment Matrix

$OPT(i, j)$ contains the score for the best alignment between:

the first i characters of string x [**prefix** i of x]

the first j character of string y [**prefix** j of y]

NOTE: observe the non-standard notation here; $OPT(\textcolor{red}{i}, \textcolor{green}{j})$ is referring to *column* i , *row* j of the matrix.



How to do semi-global alignment?

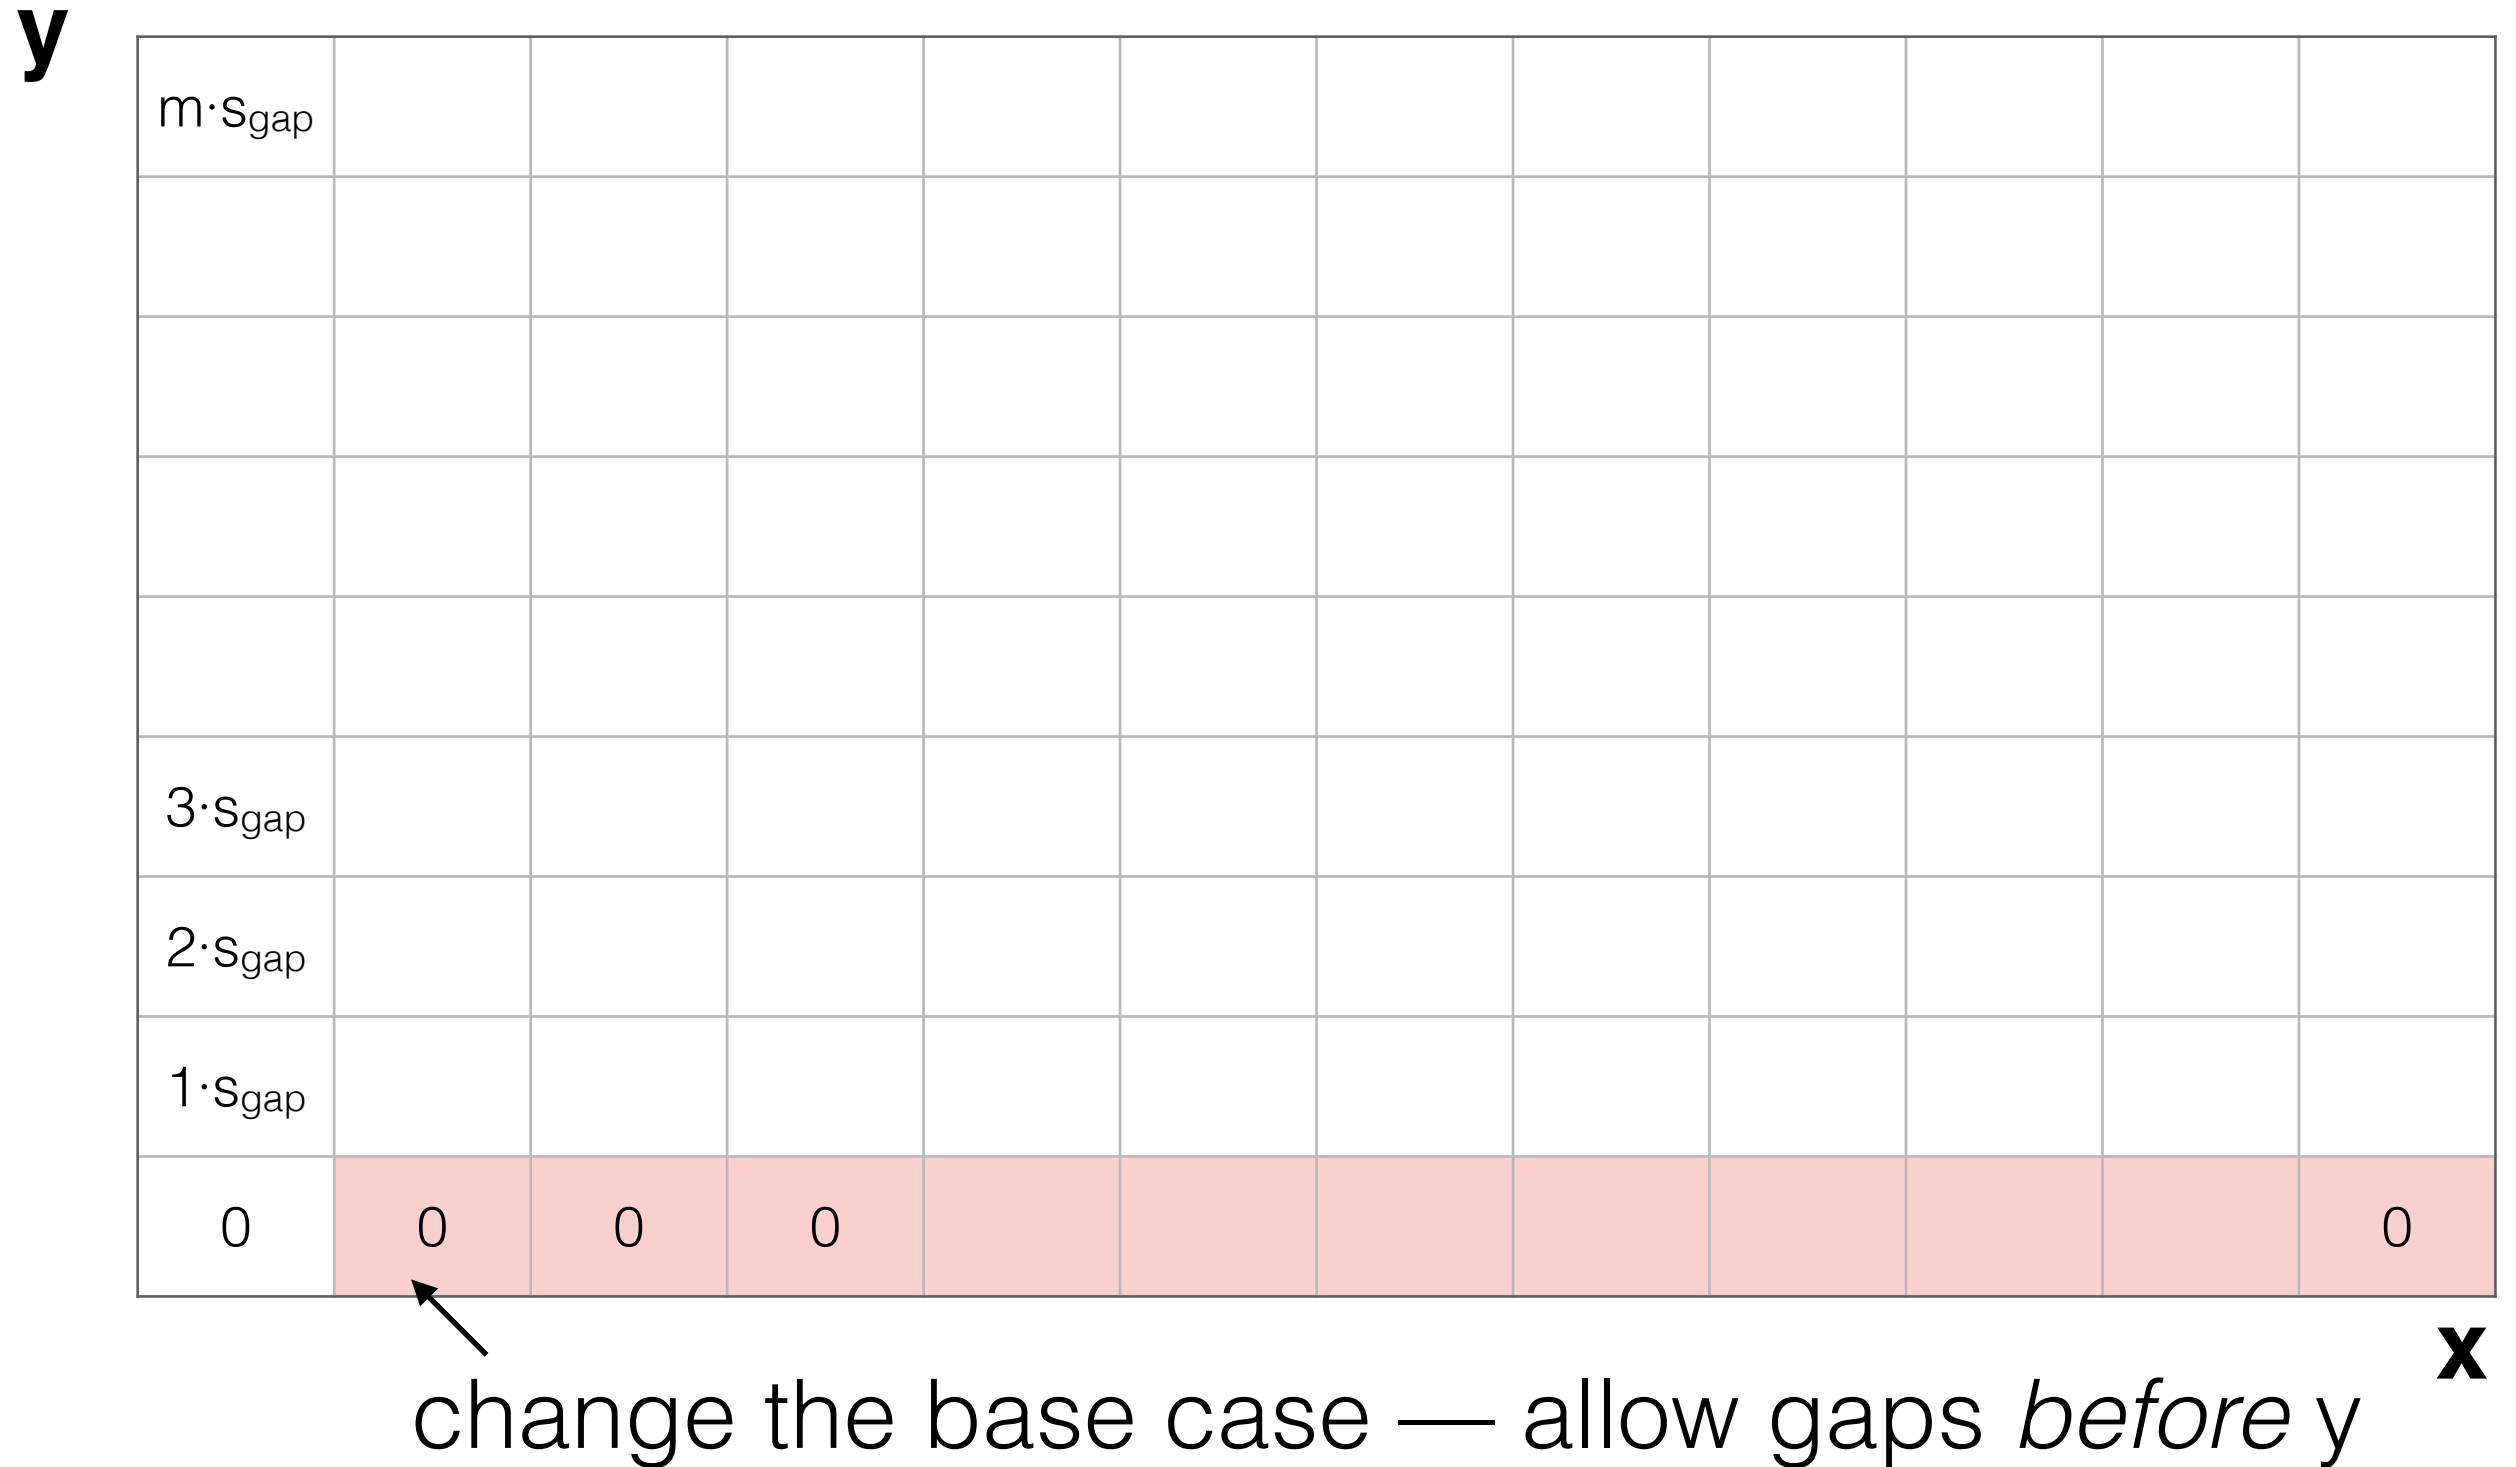
y

$m \cdot s_{\text{gap}}$											
$3 \cdot s_{\text{gap}}$											
$2 \cdot s_{\text{gap}}$											
$1 \cdot s_{\text{gap}}$											
0	$1 \cdot s_{\text{gap}}$	$2 \cdot s_{\text{gap}}$	$3 \cdot s_{\text{gap}}$								$n \cdot s_{\text{gap}}$

x

Start with the original global alignment matrix

How to do semi-global alignment?



How to do semi-global alignment?

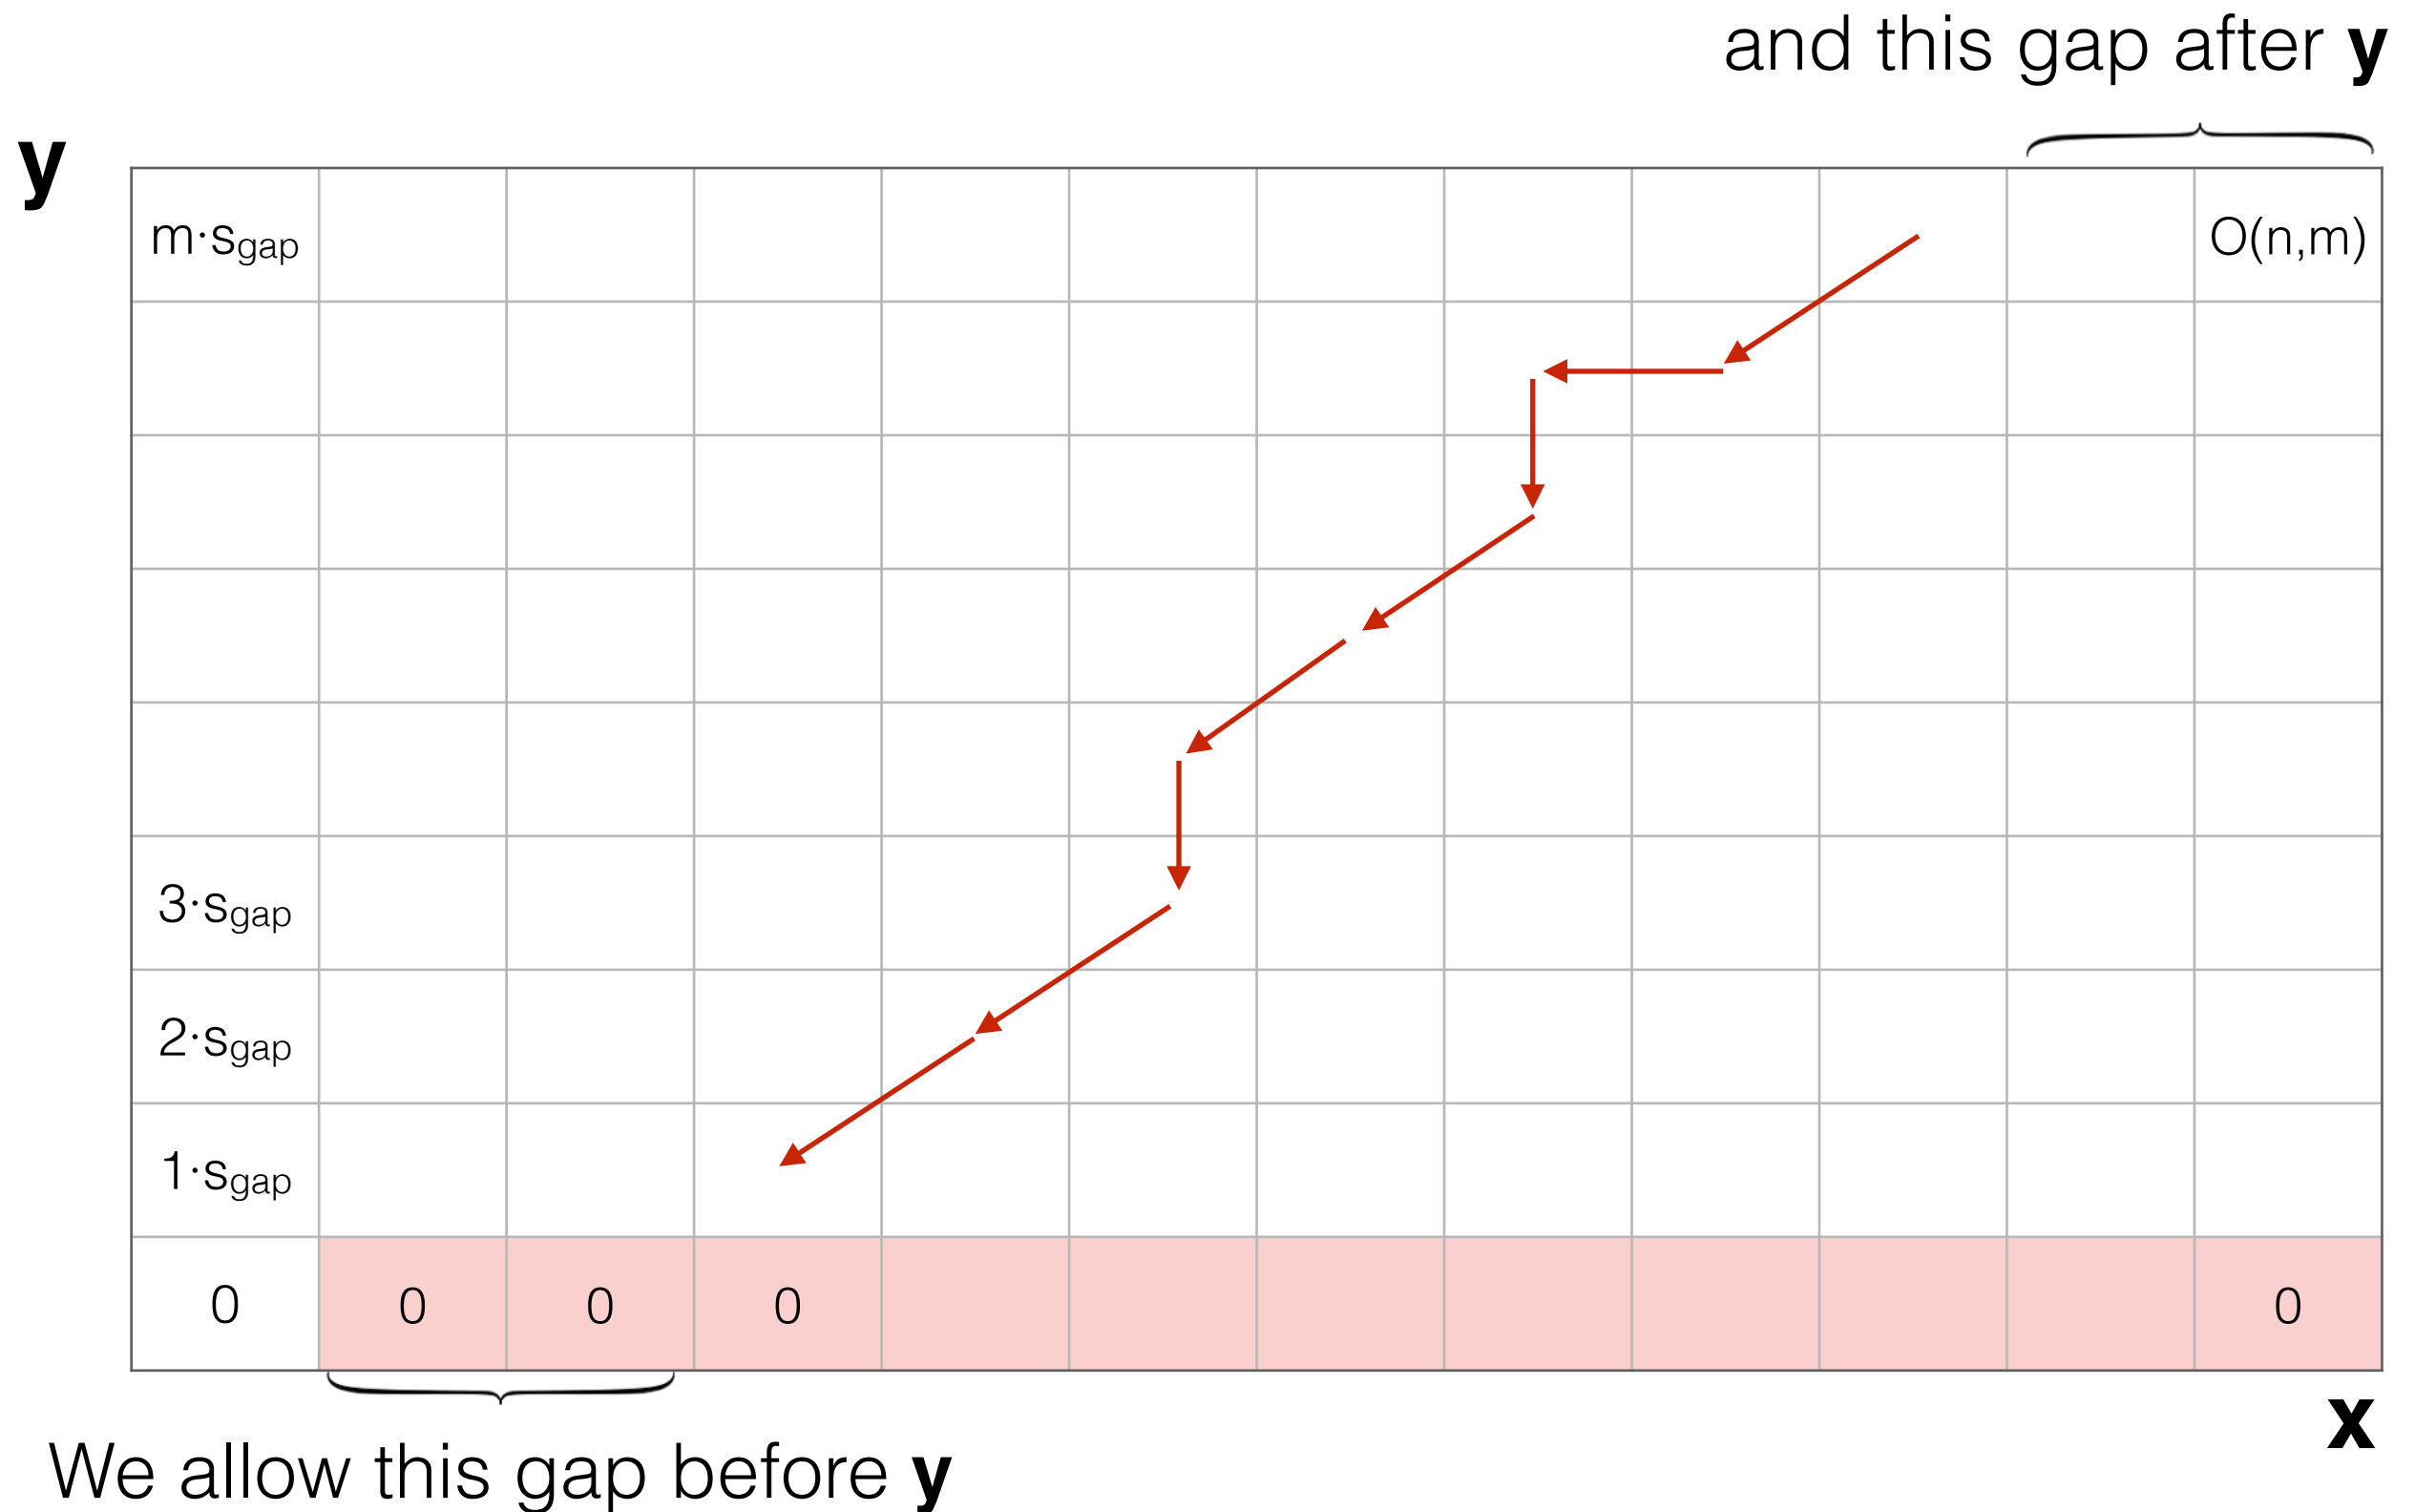
y

$m \cdot S_{\text{gap}}$											$O(n, m)$
$3 \cdot S_{\text{gap}}$											
$2 \cdot S_{\text{gap}}$											
$1 \cdot S_{\text{gap}}$											
0	0	0	0								0

x

start traceback at $\max_{0 \leq i \leq n} \text{OPT}(i, m)$ — this allows gaps after **y**; why?

Semi-global alignment example



Semi-global Alignment

What is the **same** and **different** between the “global” and semi-global (“fitting”) alignment problems?

*assuming $|y| < |x|$ and we are “fitting” y into x

Global

$$\text{OPT}(i, j) = \max \begin{cases} \text{score}(x_i, y_j) + \text{OPT}(i-1, j-1) \\ s_{\text{gap}} + \text{OPT}(i-1, j) \\ s_{\text{gap}} + \text{OPT}(i, j-1) \end{cases}$$

Base case: $\text{OPT}(i, 0) = i \times s_{\text{gap}}$

Traceback starts at $\text{OPT}(n, m)$

Semi-global (“fitting”)

$$\text{OPT}(i, j) = \max \begin{cases} \text{score}(x_i, y_j) + \text{OPT}(i-1, j-1) \\ s_{\text{gap}} + \text{OPT}(i-1, j) \\ s_{\text{gap}} + \text{OPT}(i, j-1) \end{cases}$$

Base case: $\text{OPT}(i, 0) = 0$

Traceback starts at **max** $\text{OPT}(j, m)$
 $0 < j \leq n$

Semi-global Alignment

The recurrence remains the *same*, we only change the base case of the recurrence and the origin of the backtrack

- | | | |
|-------------------------|---|---|
| 1) Ignore gaps before x | → | change base case;
$\text{OPT}(0,j) = 0$ |
| 2) Ignore gaps after x | → | change traceback;
start from $\max_{0 < j \leq m} \text{OPT}(n,j)$ |
| 3) Ignore gaps before y | → | change base case;
$\text{OPT}(i,0) = 0$ |
| 4) Ignore gaps after y | → | change traceback;
start from $\max_{0 < i \leq n} \text{OPT}(i,m)$ |

Semi-global Alignment

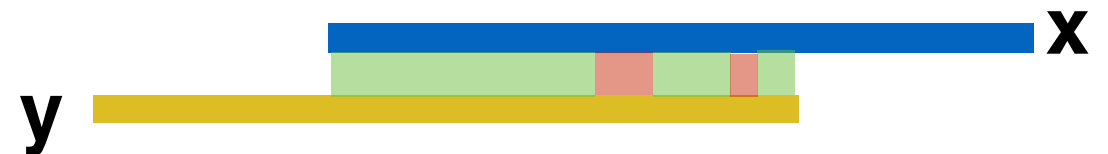
- 1) Ignore gaps before x
- 2) Ignore gaps after x
- 3) Ignore gaps before y
- 4) Ignore gaps after y

Types of semi-global alignments

use mods 3&4



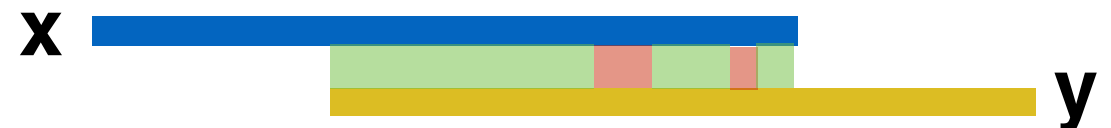
use mods 1&4



use mods 1&2



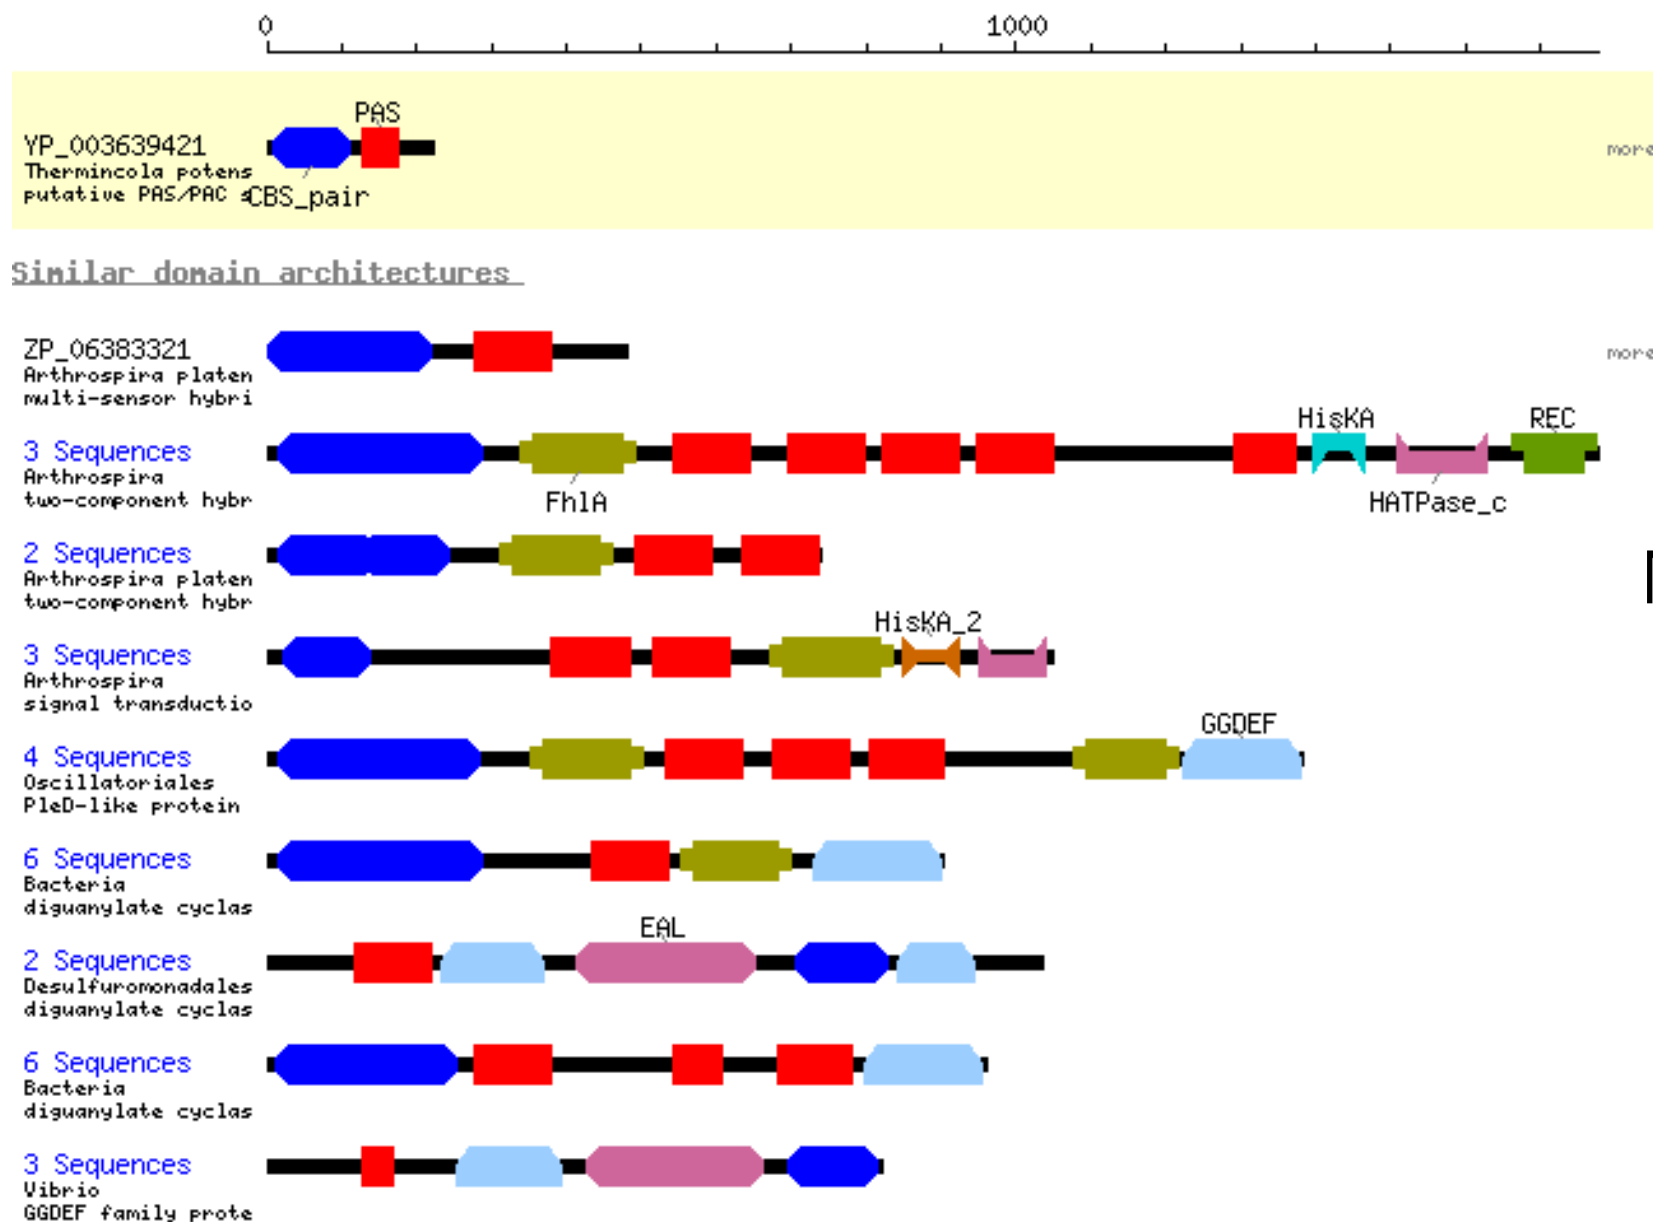
use mods 2&3



Local Alignment



Local alignment between a and b: Best alignment between a subsequence of a and a subsequence of b.



Motivation:

Many genes are composed of *domains*, which are subsequences that perform a particular function.

Local Alignment

New meaning of entry of matrix entry:

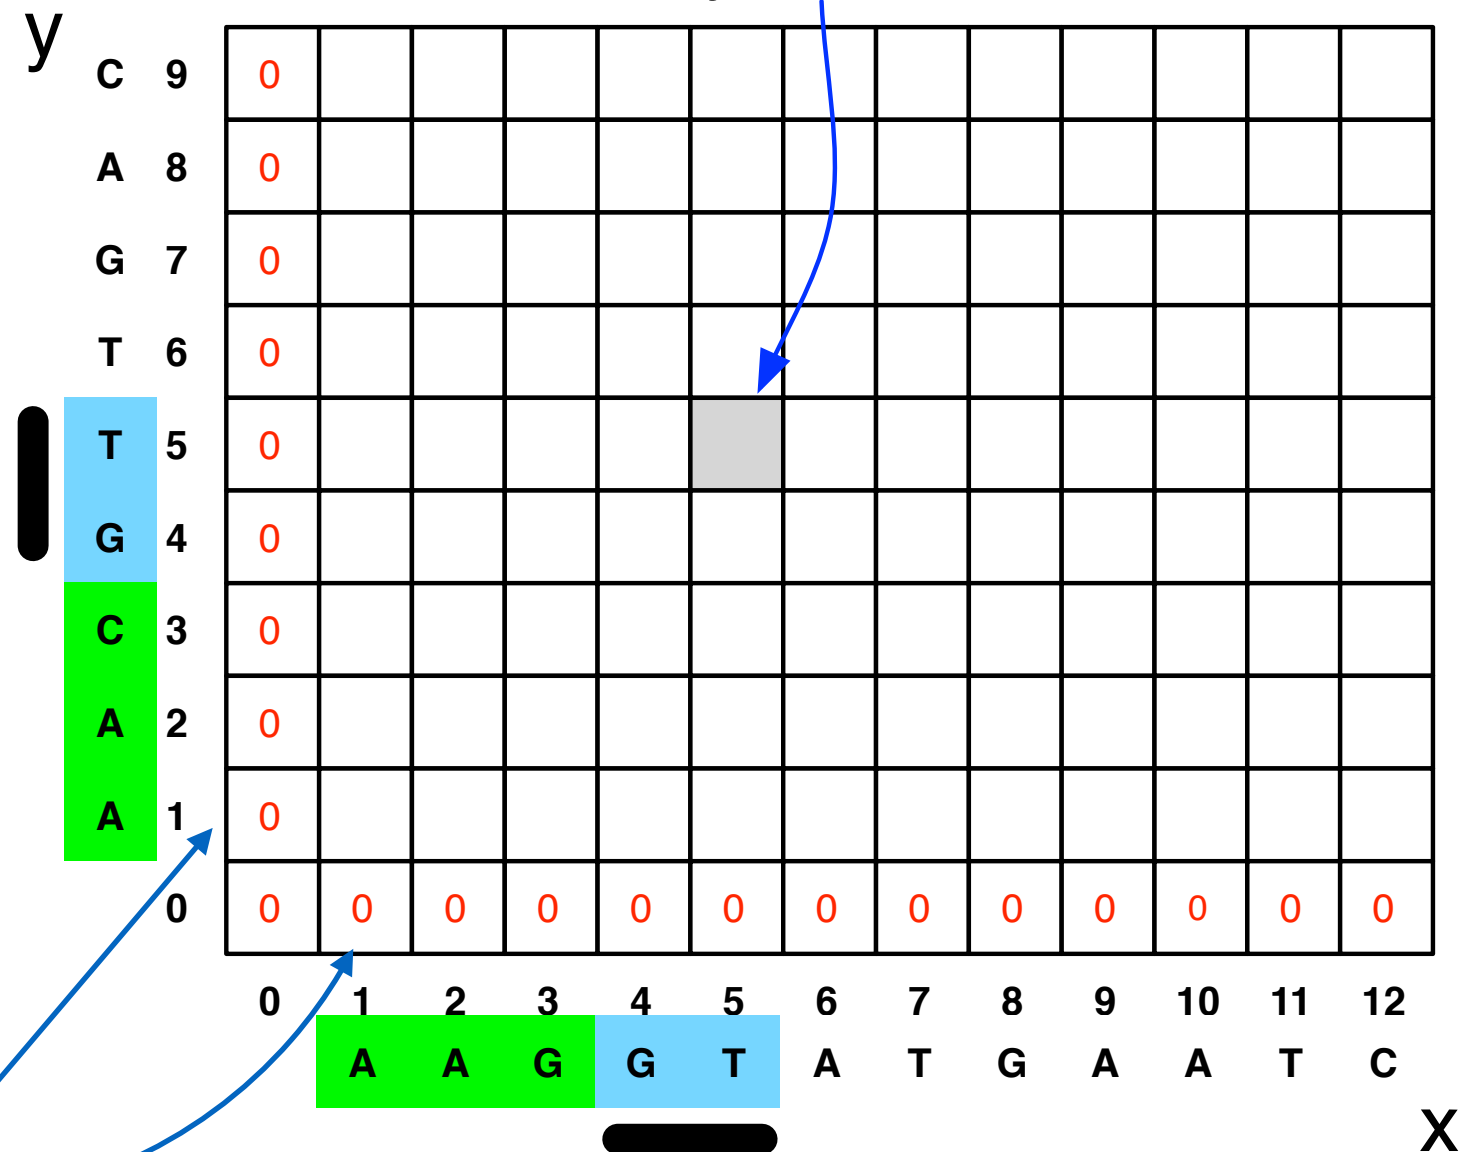
$\text{OPT}(i, j)$ = best score between:

some suffix of $x[1\dots i]$

some suffix of $y[1\dots j]$

Same base-case trick we used in semi-global alignment

Best alignment between a suffix of $x[1..5]$ and a suffix of $y[1..5]$



Local Alignment

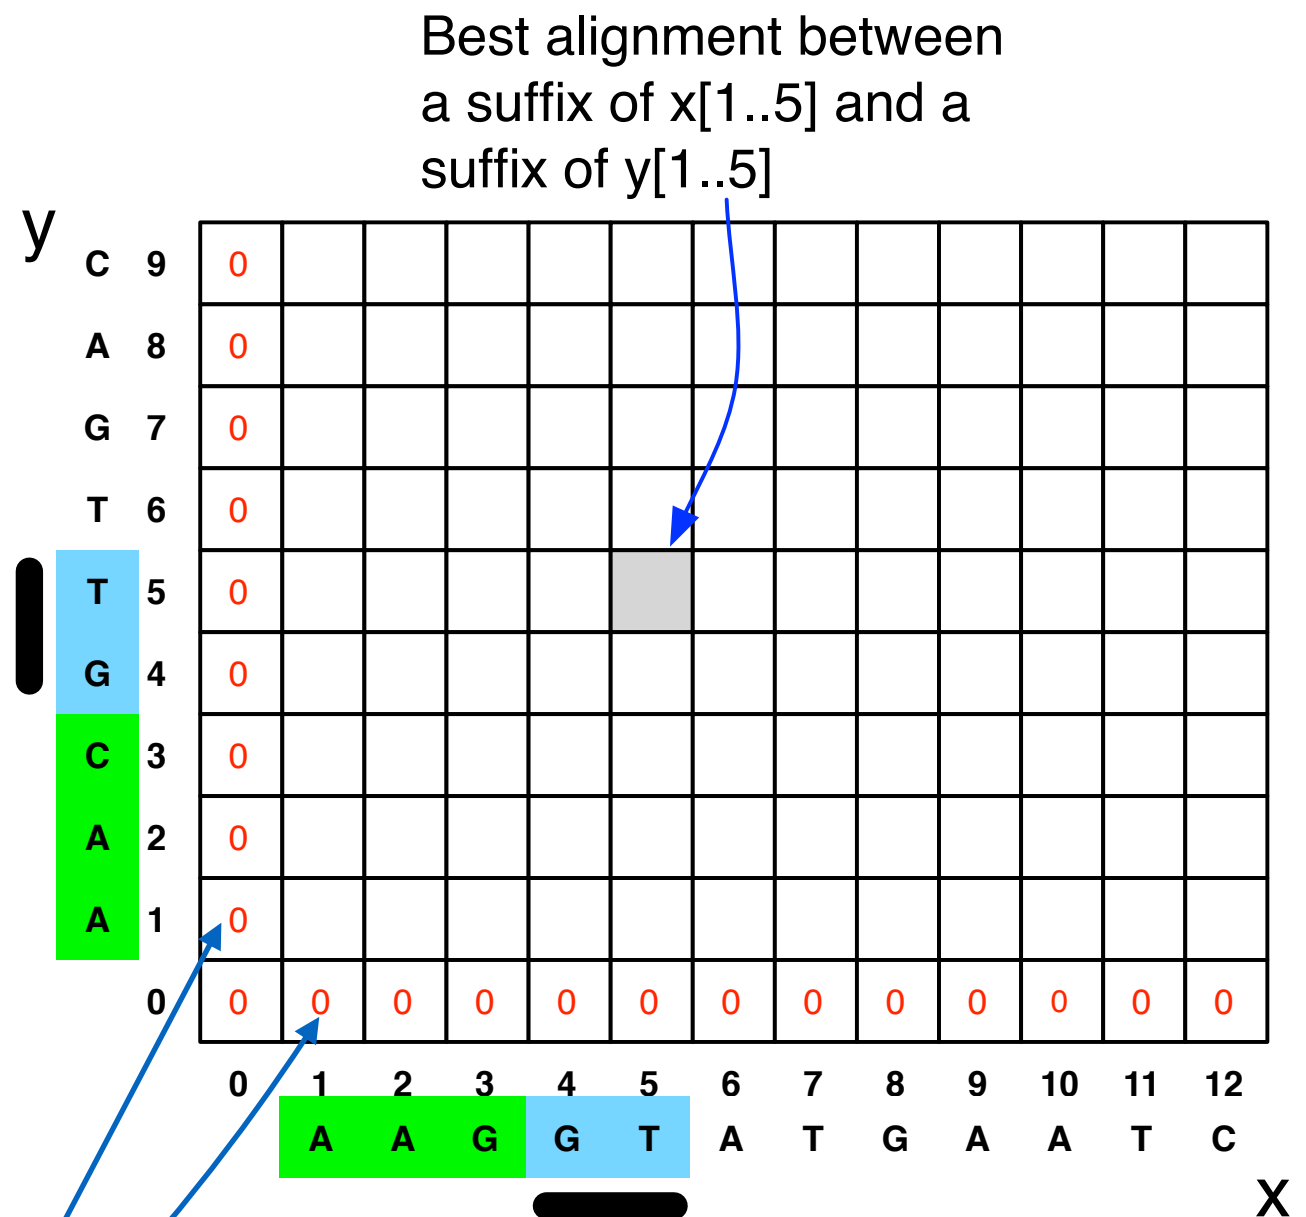
New meaning of entry of matrix entry:

$\text{OPT}(i, j)$ = best score between:
some suffix of $x[1..i]$
some suffix of $y[1..j]$

What else do we need to change to allow local alignments?

Hint: The empty alignment is always a valid local alignment!

Same base-case trick we used in semi-global alignment



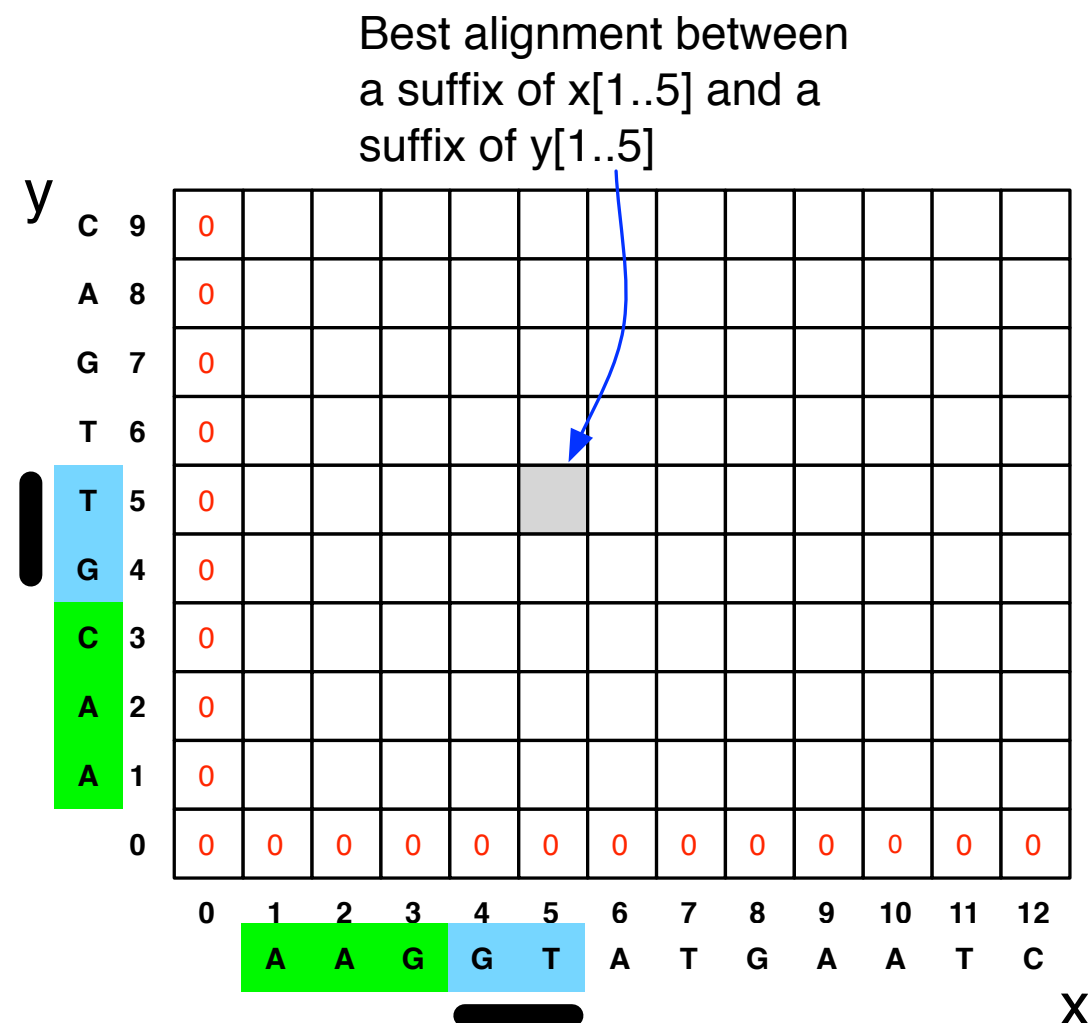
How do we fill in the local alignment matrix?

$$\text{OPT}(i, j) = \max \begin{cases} \text{score}(x_i, y_j) + \text{OPT}(i-1, j-1) & (1) \\ s_{\text{gap}} + \text{OPT}(i-1, j) & (2) \\ s_{\text{gap}} + \text{OPT}(i, j-1) & (3) \\ 0 & \end{cases}$$

(1), (2), and (3): same cases as before:
match x and y, gap in y, gap in x

New case: 0 allows you to say the
best alignment between a suffix of x
and a suffix of y is the empty
alignment.

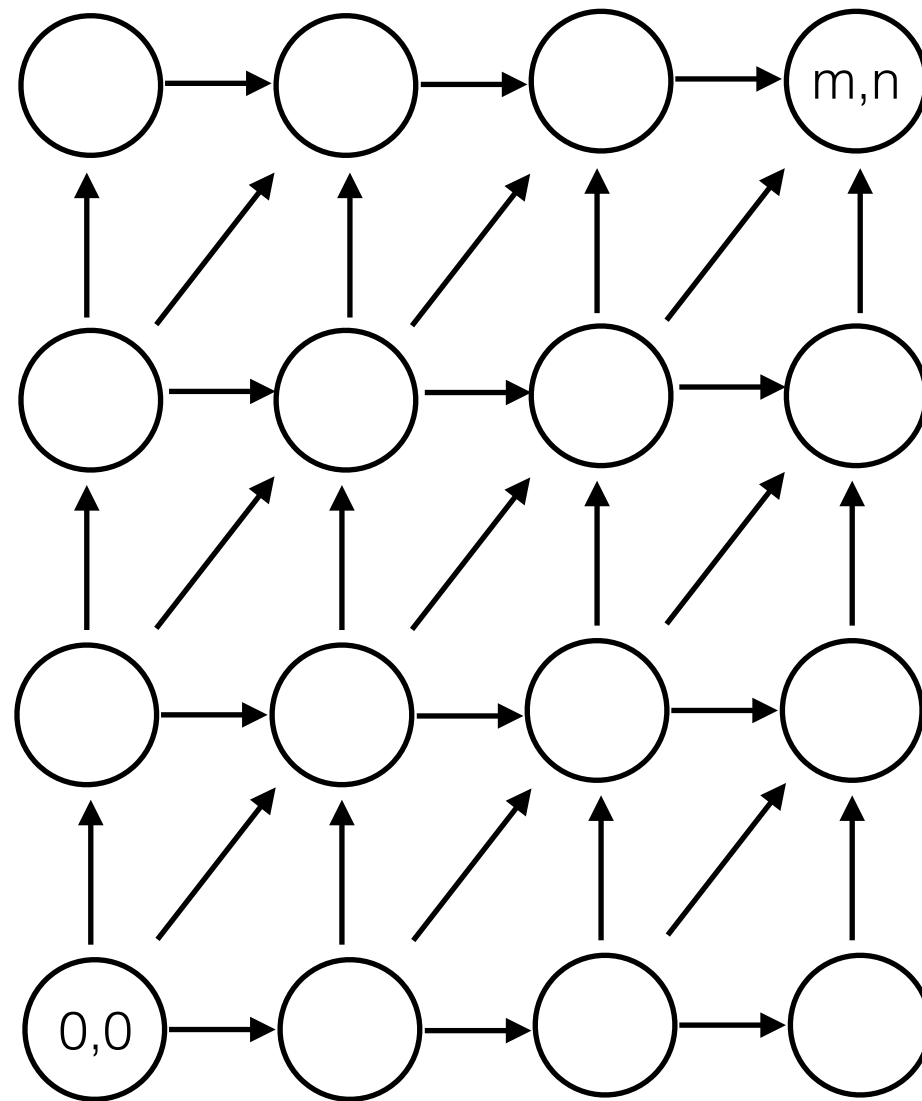
Lets us “start over”



Local Alignment

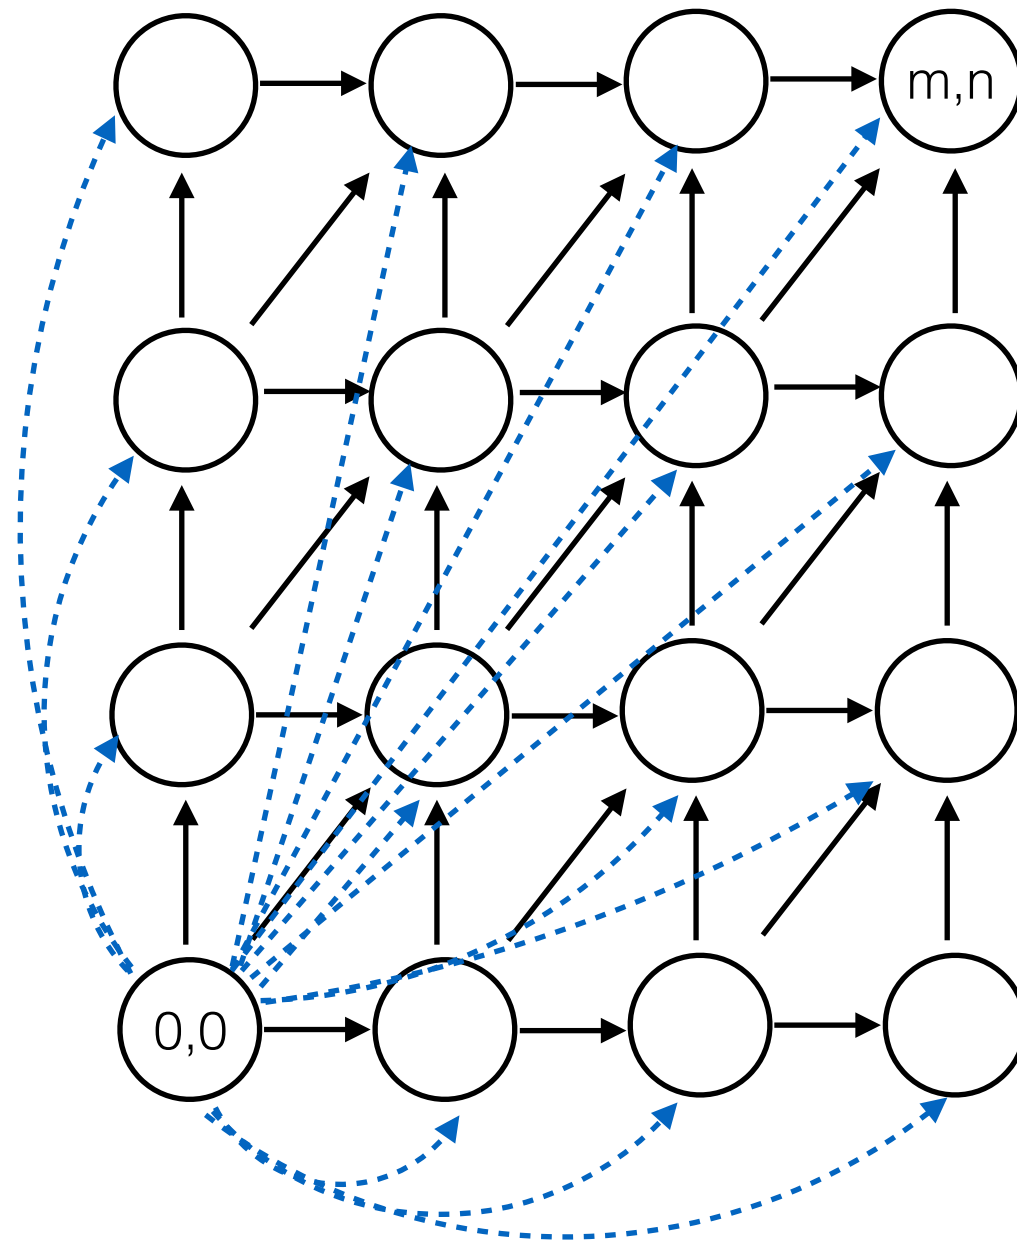
- Initialize first row and first column to be 0.
- The score of the best local alignment is the largest value in the entire array.
- To find the actual local alignment:
 - start at an entry with the maximum score
 - traceback as usual
 - stop when we reach an entry with a score of 0

Local Alignment in the DAG framework



Local Alignment in the DAG framework

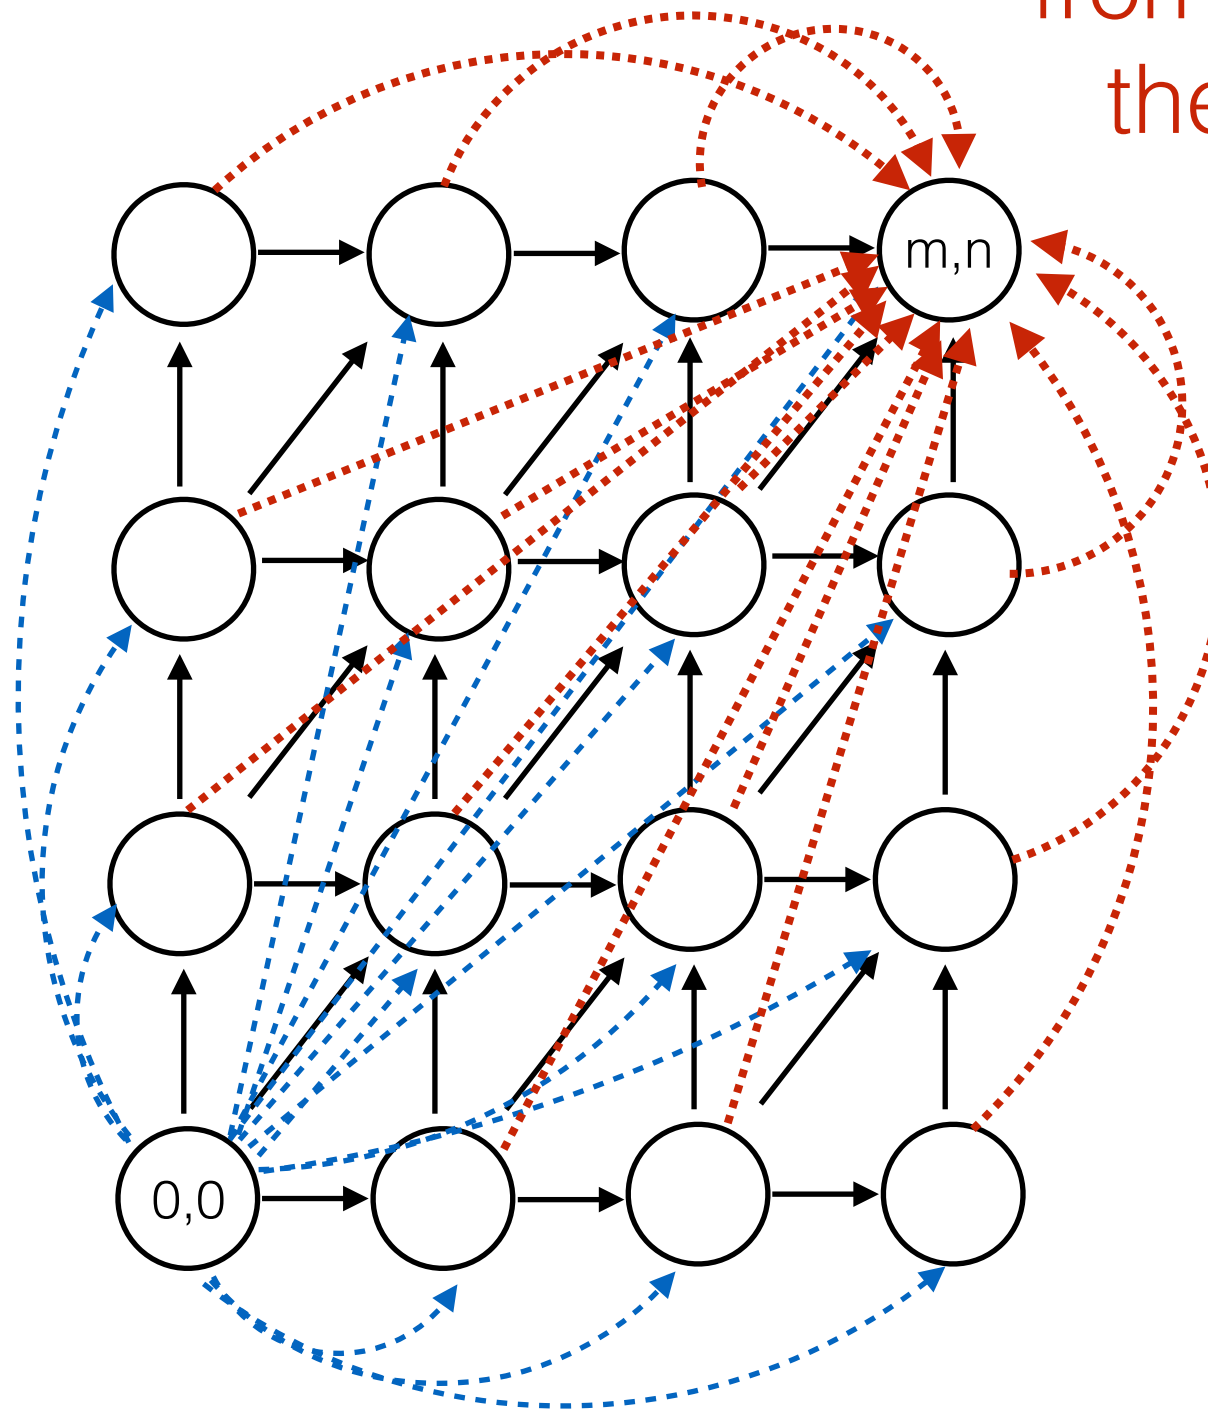
Add 0 score edge
from the source
to every node



Local Alignment in the DAG framework

Add 0 score edge
from the source
to every vertex

Add 0 score edge
from every vertex to
the target vertex



Side Note: Lower Bounds

- Suppose the lengths of x and y are n .
- Clearly, need at least $\Omega(n)$ time to find their global alignment (have to read the strings!)
- The DP algorithms show global alignment can be done in $O(n^2)$ time.
- A trick called the “Four Russians Speedup” can make a similar dynamic programming algorithm run in $O(n^2 / \log n)$ time.
 - We probably won’t talk about the Four Russians Speedup.
 - The important thing to remember is that only one of the four authors is Russian...
(Alazarov, Dinic, Kronrod, Faradzev, 1970)
- Open questions: Can we do better? Can we prove that we can’t do better? No#

#: Backurs, Arturs, and Piotr Indyk. "Edit distance cannot be computed in strongly subquadratic time (unless SETH is false)." *Proceedings of the forty-seventh annual ACM symposium on Theory of computing*. ACM, 2015.

Using semi-global alignment is not feasible for read mapping

The best algorithms we have (and like the best that could exist) to compute the optimal alignment of two strings are *quadratic*

If we have N reads, each of length ℓ , and the genome is of length L , then applying the optimal algorithm at each possible position (to test the edit distance) is **$O(N \cdot \ell \cdot L)$**

Consider a dataset with:

$N = 20 \times 10^6$ reads

$\ell = 100$

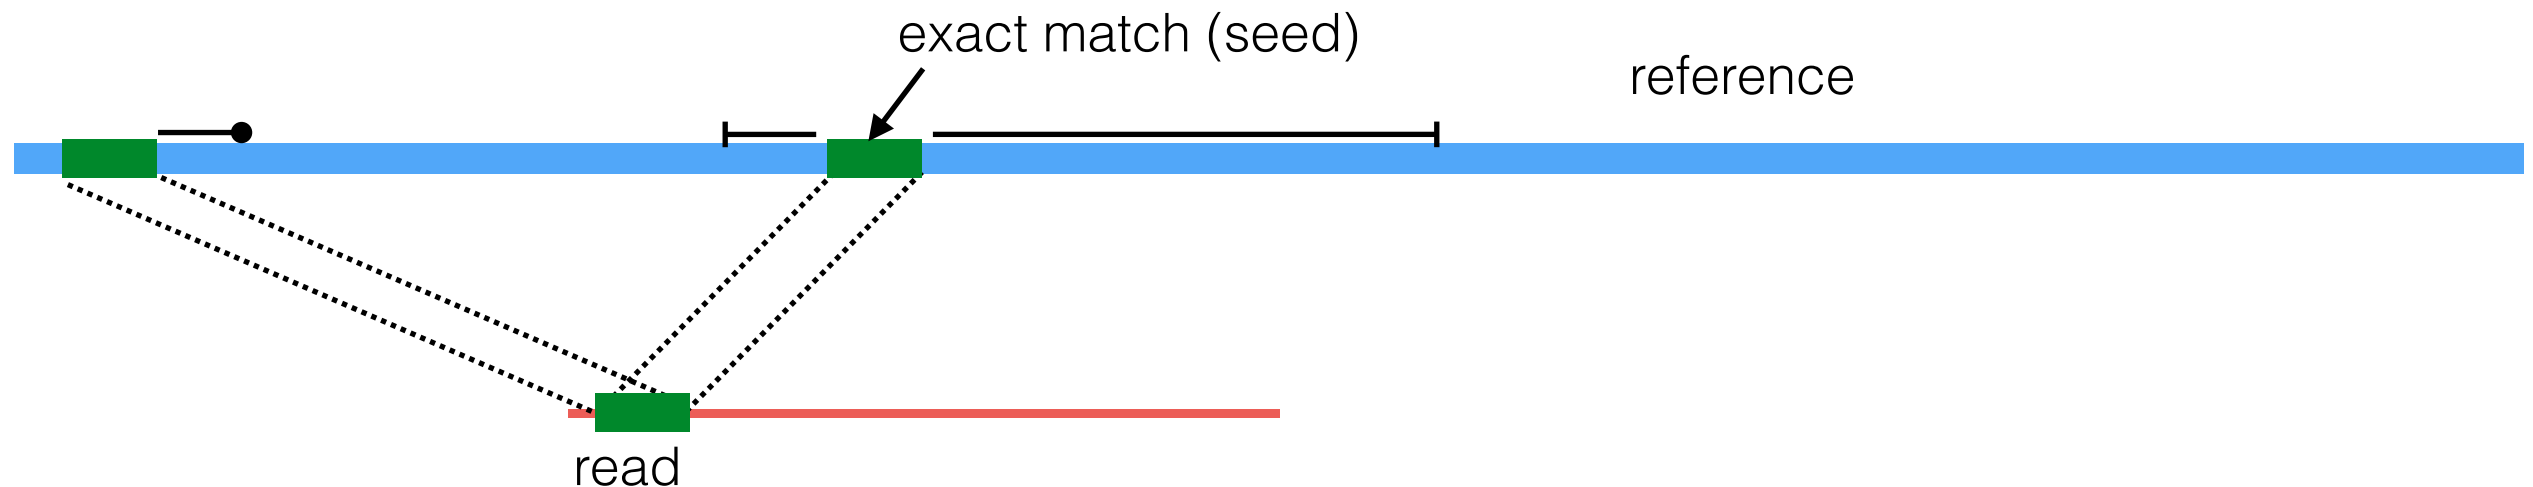
$L = 3 \times 10^9$ nucleotides

and a processor that can do $X = 3 \times 10^9$ operations / sec.

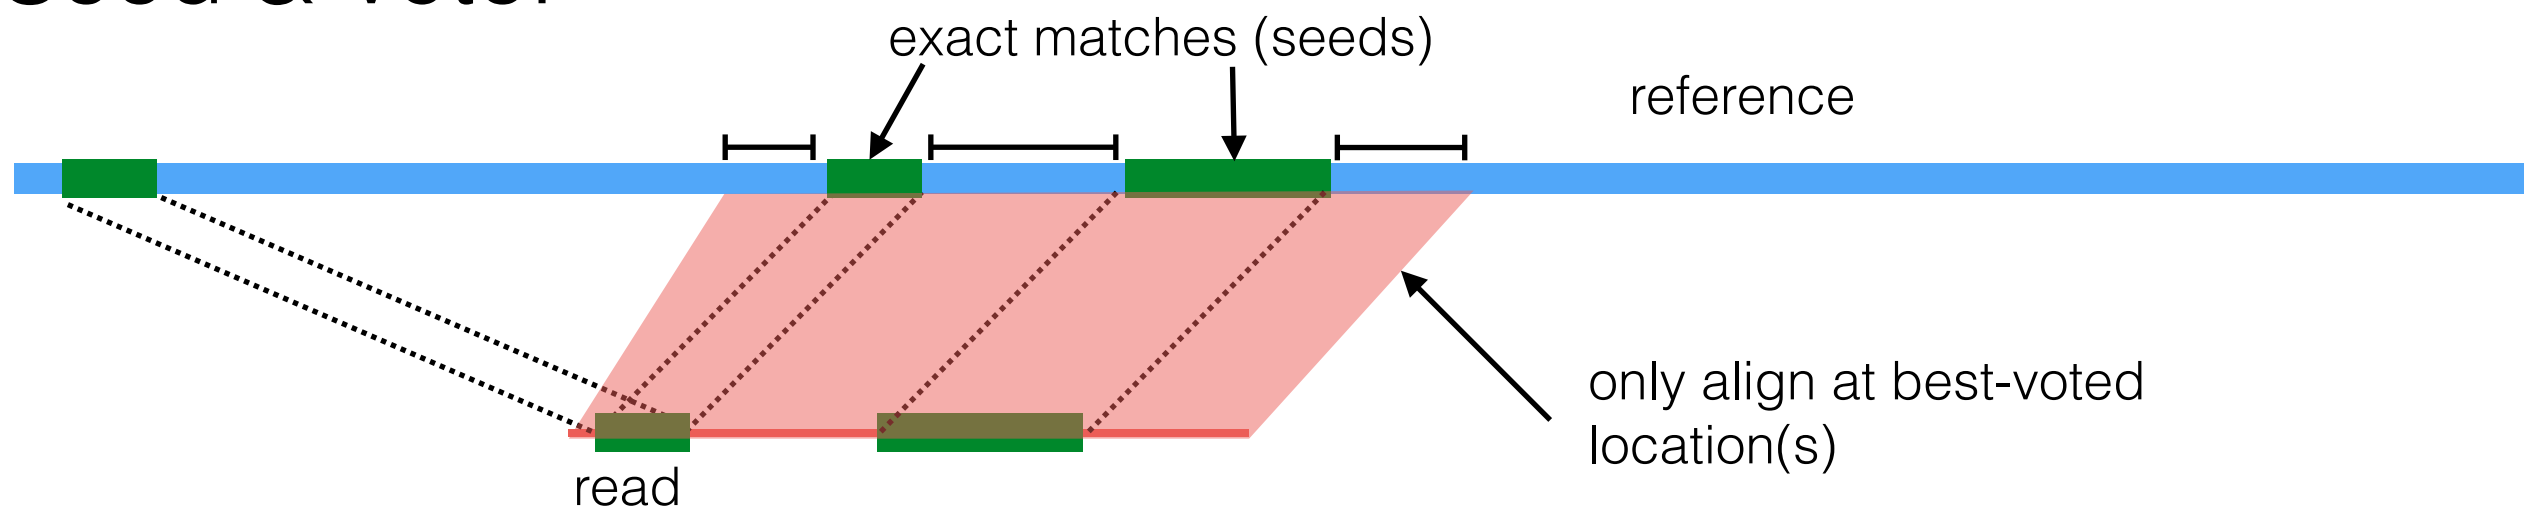
You'd wait about $(N \cdot \ell \cdot L) / X = 200,000,000$ sec = 6.34 **years** to align your reads.

How can exact matching help?

Seed & Extend:



Seed & Vote:



How can exact matching help?

GGTACCCCCAAATTCGAAAGAGCCGAATGAAGTGGAAAGAATGGCCATTCAAATGGCTTGGGA

CAGGGCGCGCCGCTGGCCGAGAAAACCGCATTACAAATACTCTCAAAGCAGGATTTCTTTCTATAGGTAGTCAGTCCACAATTTTCACAT
TTGCAGCCAGGTGGACATCTTTGGGCCAATGATGAGGAGAAGGTGGAGGAAGAAGGAAATGAGGAAGAGGAGAGAAAGAAGAGTATGGC
ATTTTACAAACTGTGACCGTTTCTGTGTGAAGATTTTATAGCTGTCTGTCTCGGACTTGGGGGTCTCAGGGGAAACTCACTTTGCCCGCCAGC
TGAGGTTTTTCAGGAAATCTGGAAACCTACAGTCTCCAAGCCTGCTCAGCCAAGAACGCGGGCGGGCGGGCGGGTGGCGACGGCGG
CGGCGGGCGCAGGACCTCCGCGCCTCCATTATGCTATTCTGCCCGCCGTGGGTGACAAACAGACGGATGCTACCAGCCCAGCCCAGT
CCCGGGGAGCCAGCTGGCCTGGGGTTCGGTCCCGCGCTCTTCCCTCATTCTGTGCCGCTGCCGAGCCTGTCTCAGCTCCACACACG
CTTGGGAGCTGCAGATGCCTCCGCCCCTCCTCTCTCCAGGCTCTTCCTGCCGTTGAACCCCGGGCGGGCGGGCTCTCGGCCAGCGG
CGCGCCCTGGTACCCCCAAATTCGAAAGAGCCGAATGAAGTGGAAAGAATGGCCACAATGGCTTGGGCCCCGAGTGACCATGGGA
TGGTTAGGTAGGATTTTAGAGGCGACTGCTCCTGGAATTAGAGAAAGAGTTTCATTACAACCGCTACCTGACCCGAAGGCACTCGGGC
CCATGCCTTCCTCTCCTTCGCTGTTTGATTTCTATTCTGTTTAGCCTGAGTTGGAGATGCCGGAGGTGCCGTGGTGGAAAGAAGTCTTGA
ACGAATTTGGAGGCGTCTCCGTGGCAGCTAAGCGAGCACGGGTCTGCTGGTGCAGGATGACACTGGCAGCCACTGCCGCGGACTT
GTATCTCTTGTCTTCCTGCTTTTATAGAGAATAGAATGACACTCACAACCTCTAACTACCTGTCAGAAGCAGGCAGGAGCTAGTAAGGATGA
ATTTGTAGCAAAATTAGCAAGTGGACTTCTTTCTCCTCTTCCTCATTTCCTTCTTCCTCCACCTTCTCCTCATCTTAAATCTTTAACATACTA
CCTAAAGGGAACCTGCAATAATCTTGAAAAAGGACTTCAATCCGACGTTTTTCGTGTCAAATAAGGATTAAAGAGAAACTCCTCCGCGAGC
CGTGCGCCGAGGGTGGCGGGCGGGGGCCTGAAGCGTGAGGAGCCTTCAATATGTATTTAACAGGGACCGTCGGTATGAGGTGGCCC
GGGTTCTTATTTGTTTGGGGGCTGGAGGGGGGAGACGGAGAAACAGTGAAAAGTTCCTGAGCCCCATAAAGGGACTGTCTGGGGAGC
GCCTCGTAGCCATAGAATTCCACCGCCGCGCCCGCCGCGTAGTCGTACTTGAAGCCGAGCGCAGGCGGGGTGGTTCATTAACCTCTGA
CTTTGCCTTTGATTTTGCTCGACCTCTGCTTCGTCAAATCTGGTTTCAGAATCGAAGGATGAAGATGAAAAAGATGAATAAAGGAGGAA
AAGGAAGAAAACAAGGACTAAGCAAAAAAGAAAGACCCCCCTATAGCAGGATTTTAAAATTTTCTCTTTTTCTTTTTCAAGATTATTGCA
AGGCGAGCGTGGTGCAATATCCCGACTGTAAATCCTCCGCCAACACTAACTTTTAAAAAAAACACCCAGCAGGTACCATGCTAAGACA
ACATCACATGCATTATTATGACTCACGTATACAATACAAAGTACTTGGACCAGGAACAGGGTCTTTAATCCTTATTTGACACGAAAACGTCG
GATTGAAGTCTCTCATGCCCAACTAGTGGGGTTTCCTGGCACTGGACCCCAAGTGGTCTAGAGGCGAAAAGGAAGAAAACAAG
GACTAAGCAAAAAAGAAAGACCCCCCCCCCTGAAAGGATATCATGGATTCCTTAAAGAATGAGAACTTCGACATGGTGATATCACATGG
ACTTCTGGGGCCGAGTGAAGAATTTTCTGATGTTCTTTAGTTAGCAGTCCTAAACCCTACCCAGCCTGCTGCCTCAGCACAGCCAAGG
GAAAATCAGCAAGGGCTATGCTGTGATTCTGTCTCTGAGTGACTTGGACCACTGTTGATTTTTATTTTTTCTCTTTTTTCTCCTATAGCA
GGATTTTAAAATCGGGCCCACCTTAACTCGGGAGGGCCGCGCTGAGGCTGGGAGCCGGAGATTCGGGGCGAGGGGCAGTGTCTGCGG
GGCGCGGTGCGCGCAGCTCCCCGGGCGAGCCAGGTGCAGCCTTGGCGGGGTCTGTTTGGTGGGCGATGTCACCATTTCGCGCCGCC
GCCGTGCGCACCGCCGCGCCCGCCGCGTAGTCGTACTTGAATAGCTGGACATAAAGACAAATGACAAAAAATTATTATTATAGATATATT
TGGTCTGTGTGTTATGTCCTAAGGTGTTTTGTCTGCAGTTTGAGAGCATGTTGCTGGTAGCCTGAGTTGGAGAT

How can exact matching help?

List of tuples of 13-mers;

(position in the query, position in the reference)

-1 indicates not found (Python behavior)

```
(0, 615); (1, 616); (2, 617); (3, 618);  
(4, -1); (5, -1); (6, -1); (7, -1);  
(8, -1); (9, -1); (10, -1); (11, -1);  
(12, -1); (13, -1); (14, -1); (15, -1);  
(16, 631); (17, 632); (18, 633); (19, 634);  
(20, 635); (21, 636); (22, 637); (23, 638);  
(24, 639); (25, 640); (26, 641); (27, 642);  
(28, 643); (29, 644); (30, 645); (31, 646);  
(32, 647); (33, 648); (34, 649); (35, 650);  
(36, -1); (37, -1); (38, -1); (39, -1);  
(40, -1); (41, -1); (42, -1); (43, -1);  
(44, -1); (45, -1); (46, -1); (47, -1);  
(48, -1); (49, 662); (50, 663); (51, -1);
```

So the basic strategy is :

Find **exact** matches shared between the read and reference

Group exact matches into regions likely to support a high-quality alignment

Score / validate each hit location, and filter the ones that fail to yield a high-quality alignment.

There are **many** variations on this theme. What is a good type of seed? How should we search for seeds? How should we group seeds? How aggressively should we filter potential loci?

Common types of seeds

K-mers: fixed substrings of length k.

string

AGATTACGACATAGAGCCAATATTTAGACAGATAGC

In some sense, k-mers are the “simplest” type of seed.

K-mers in the text are independent of the query.

How might we “index” k-mers in the text?

all 13-mers



AGATTACGACAT
GATTACGACATA
ATTACGACATAG
TTACGACATAGA
TACGACATAGAG
ACGACATAGAGC
CGACATAGAGCC
GACATAGAGCCA
ACATAGAGCCAA
CATAGAGCCAAT
ATAGAGCCAATA
TAGAGCCAATAT
AGAGCCAATATT
GAGCCAATATTT
AGCCAATATTTA
GCCAATATTTAG
CCAATATTTAGA
CAATATTTAGAC
AATATTTAGACA
ATATTTAGACAG
TATTTAGACAGA
ATTTAGACAGAT
TTTAGACAGATA
TTAGACAGATAG
TAGACAGATAGC

Common types of seeds

MEMs: Maximal Exact Matches

AGATTACGACATAGAGCCAATATTTGGACAGATAGC	Query
GCCAGATTACGACATAGAGCCAATATTTAGACAGATAGCTT	Ref

An exact match shared between the query and the reference that *cannot be extended, in either direction, without introducing a mismatch.*

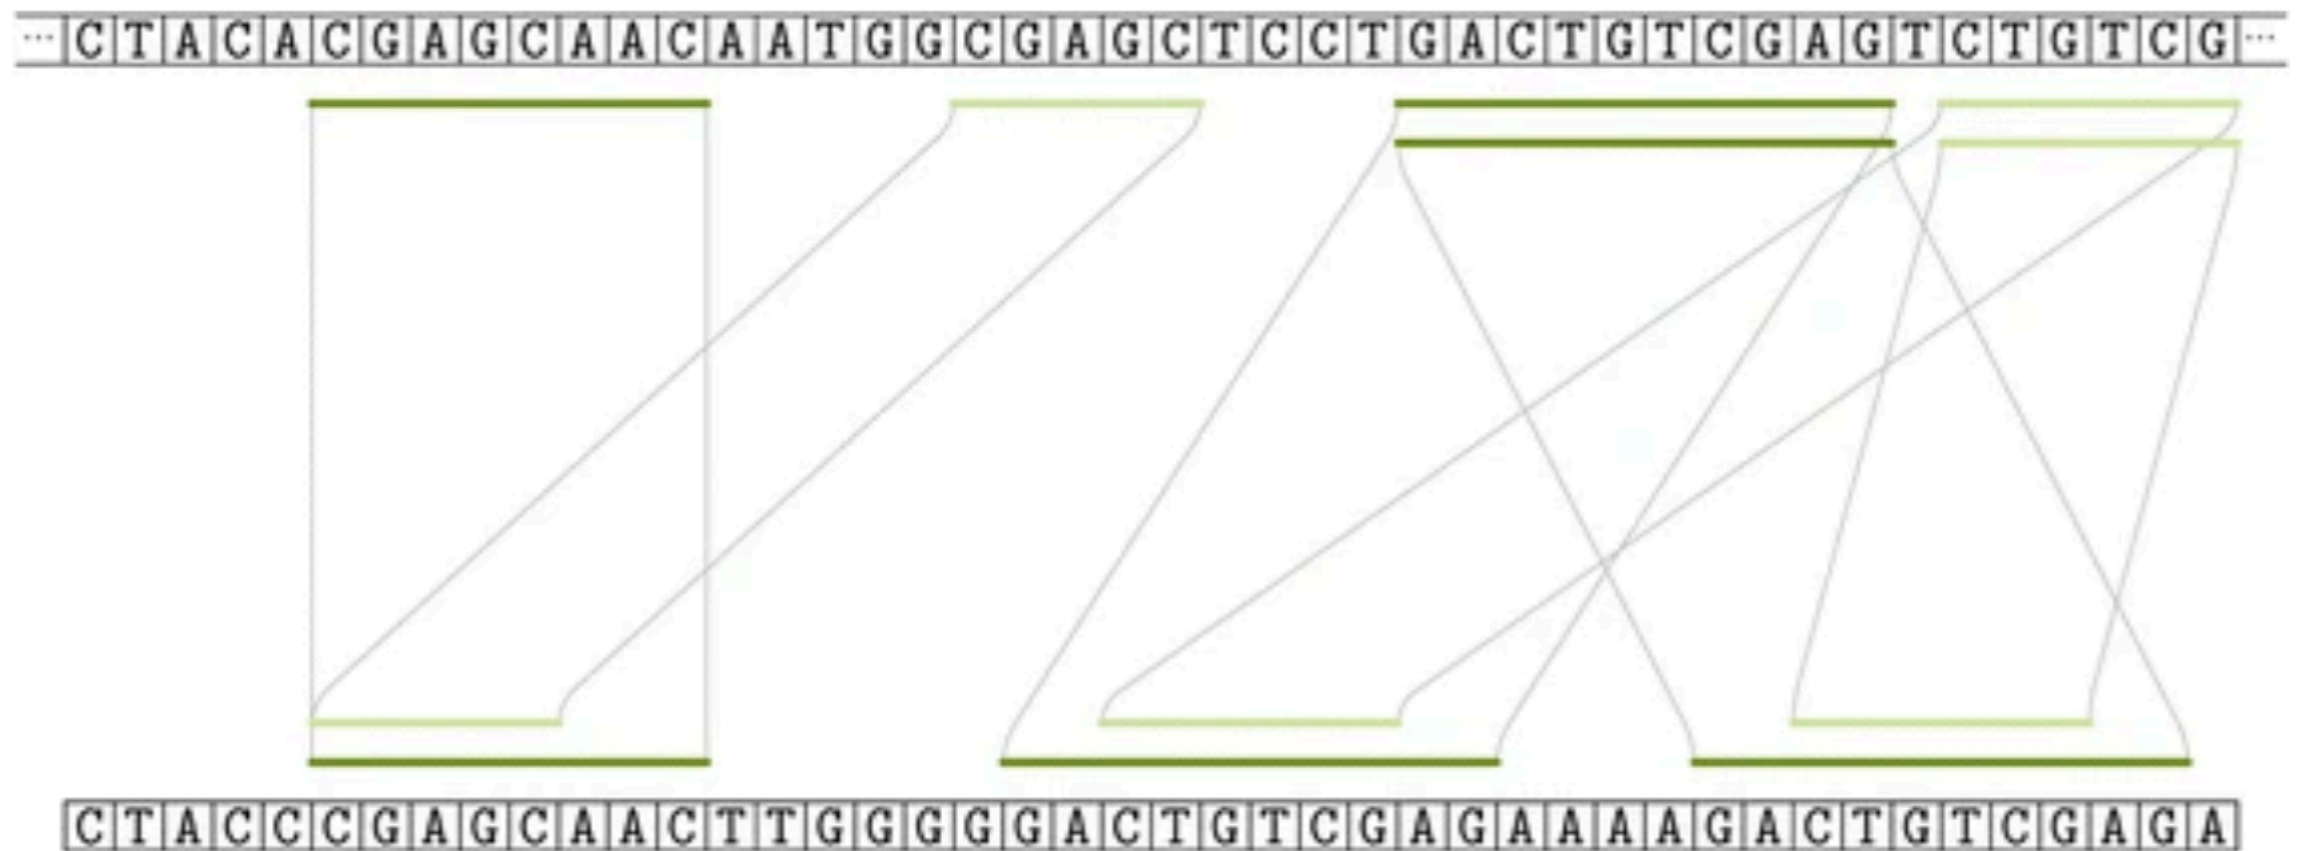
Unlike k-mers, the MEMs depend on both the reference and the specific query.

To find MEMs efficiently, we'll need a “full-text” index, not just a token / k-mer index.

Because of their “context dependence”, MEMs can be more specific than k-mers, though we don't often deal with individual seeds in isolation.

Common types of seeds

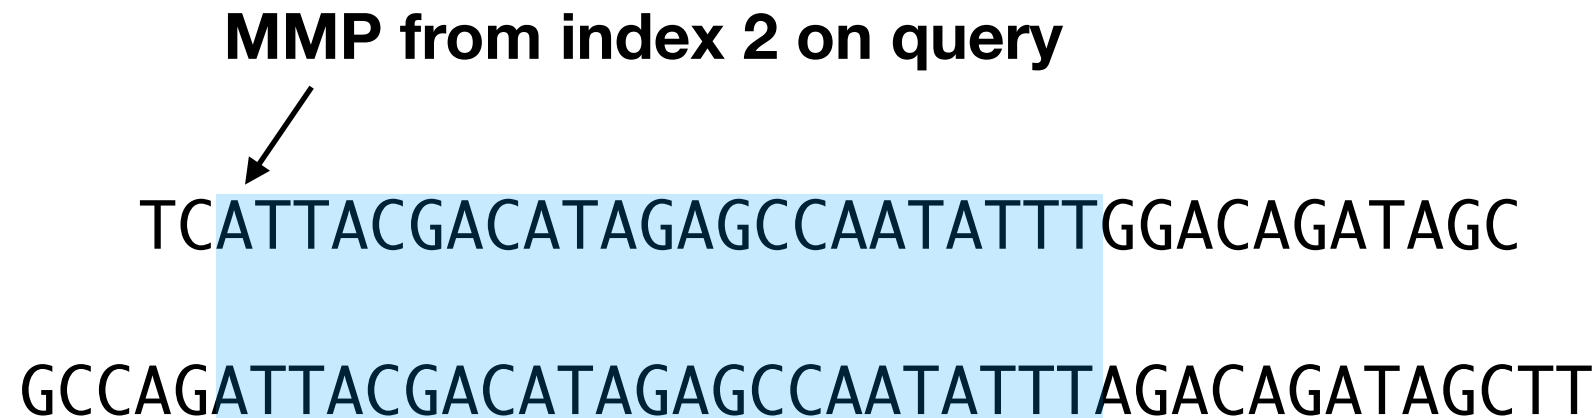
SMEMs: Super-Maximal Exact Matches



An exact match shared between the query and the reference that *cannot be extended, in either direction, without introducing a mismatch*. Also, an SMEM is not contained within any other exact match on either the query or reference.

Common types of seeds

MMPs: Maximum Mappable Prefixes



A prefix of (some suffix of) a query that is an exact match with a substring of the reference, and which cannot be extended further without introducing a mismatch.

Similar to a MEM, but extension only works in one direction — the MMP depends on the query, reference and start position. Originally introduced in STAR aligner. Useful for mapping read “pieces” across exons.

Indexing

Hopefully, I've convinced you of the importance of being able to quickly find different types of exact matches (seeds).

The next few lectures will be about data structures, and the corresponding search algorithms, that will enable this on *genome-scale* data.

We will consider both “full-text” indexing and inverted “token-based” indexing (k-mers).

There are many amazing results in the literature about indexing large text corpora; but there are likely still improvements to be made!